

Semantic Segmentation of Orthorectified Tiles: A Modern pipeline and workflow

Subtitle

DEFENCE AND SPACE

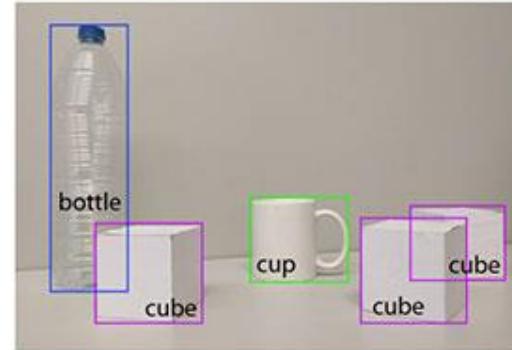
Maher SEBAI – TIN3D
December 2019

AIRBUS

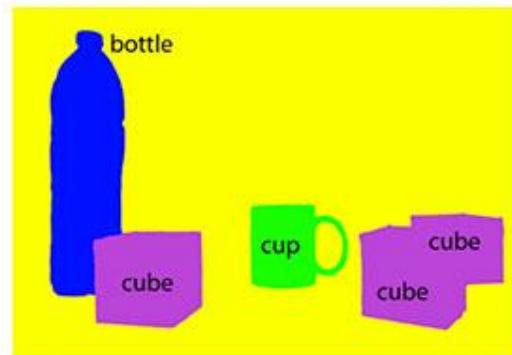
What is Semantic Segmentation in Deep Learning Context



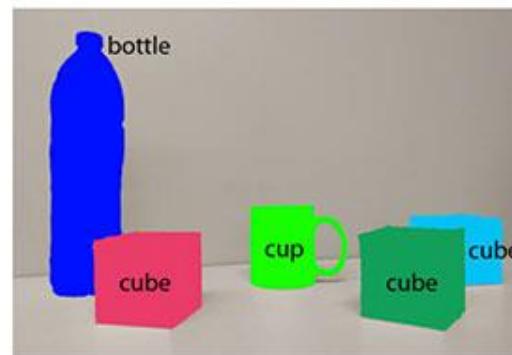
(a) Image classification



(b) Object localization



(c) Semantic segmentation



(d) Instance segmentation

Smart 3D City: Potential Benefits

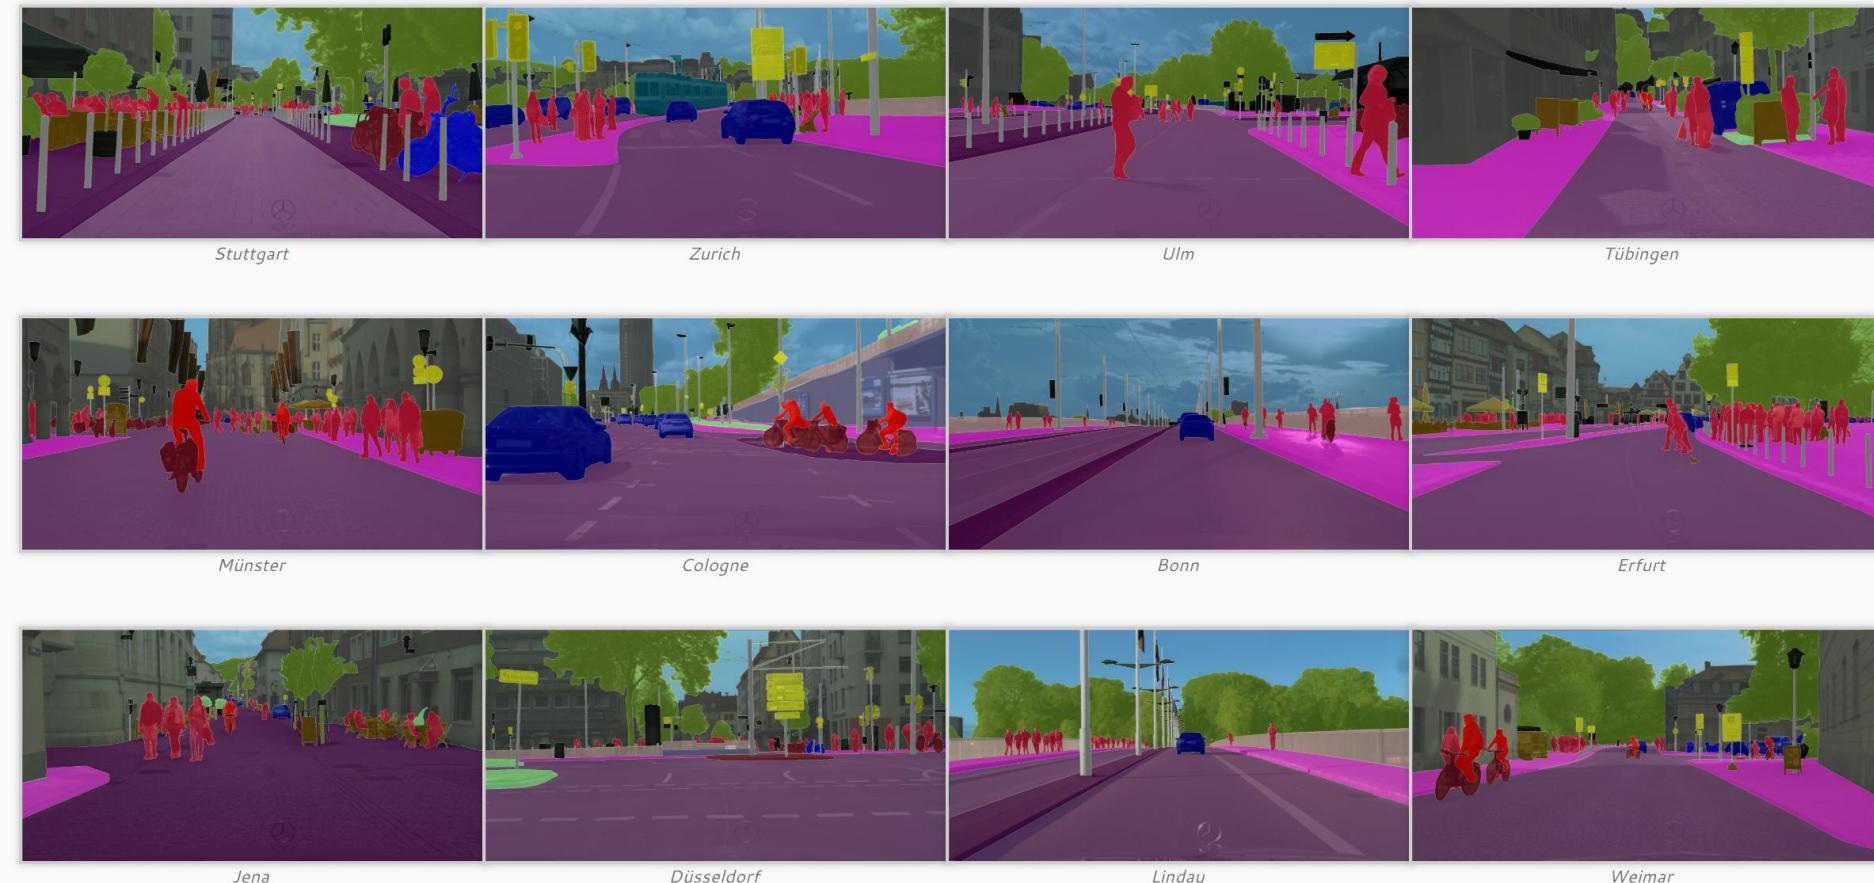
- Product added Value
 - Point Cloud Semantic
 - Urban planning
 - Statistics ...



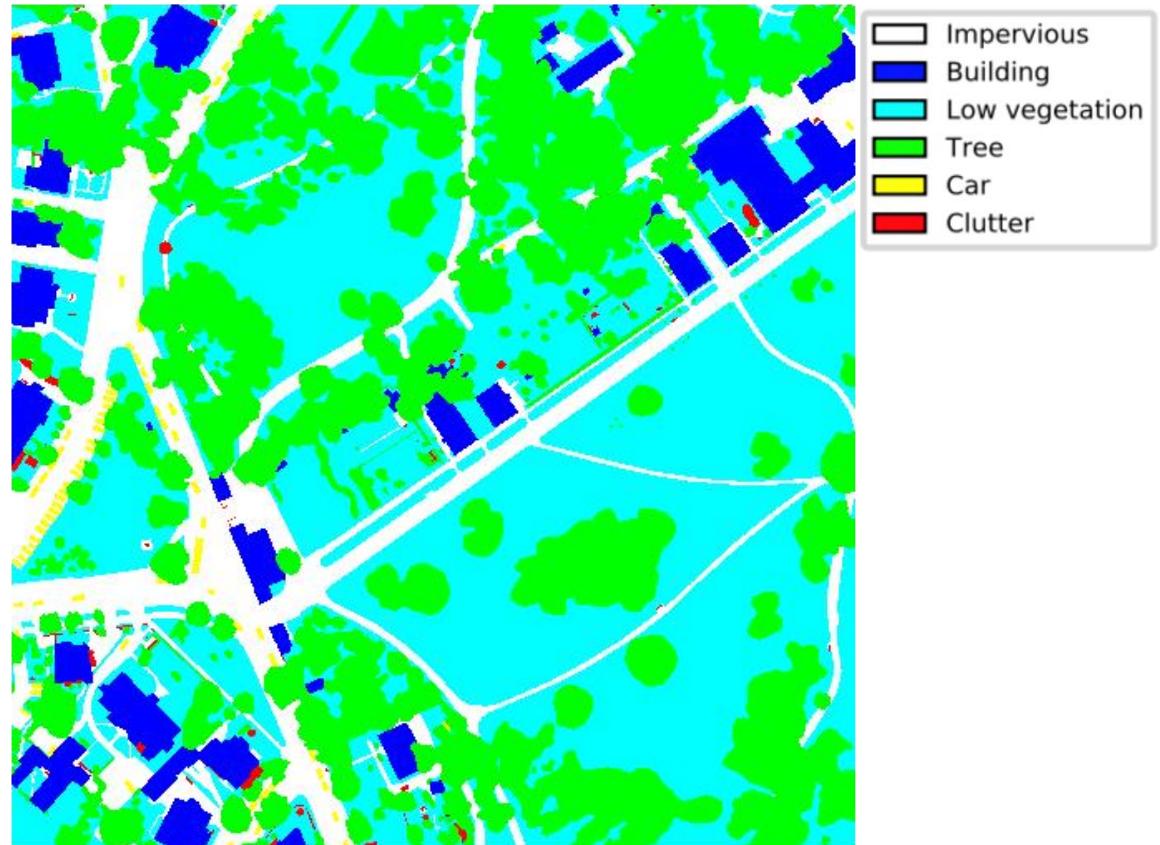
Supervised Semantic Segmentation Task Examples: MS COCO Dataset



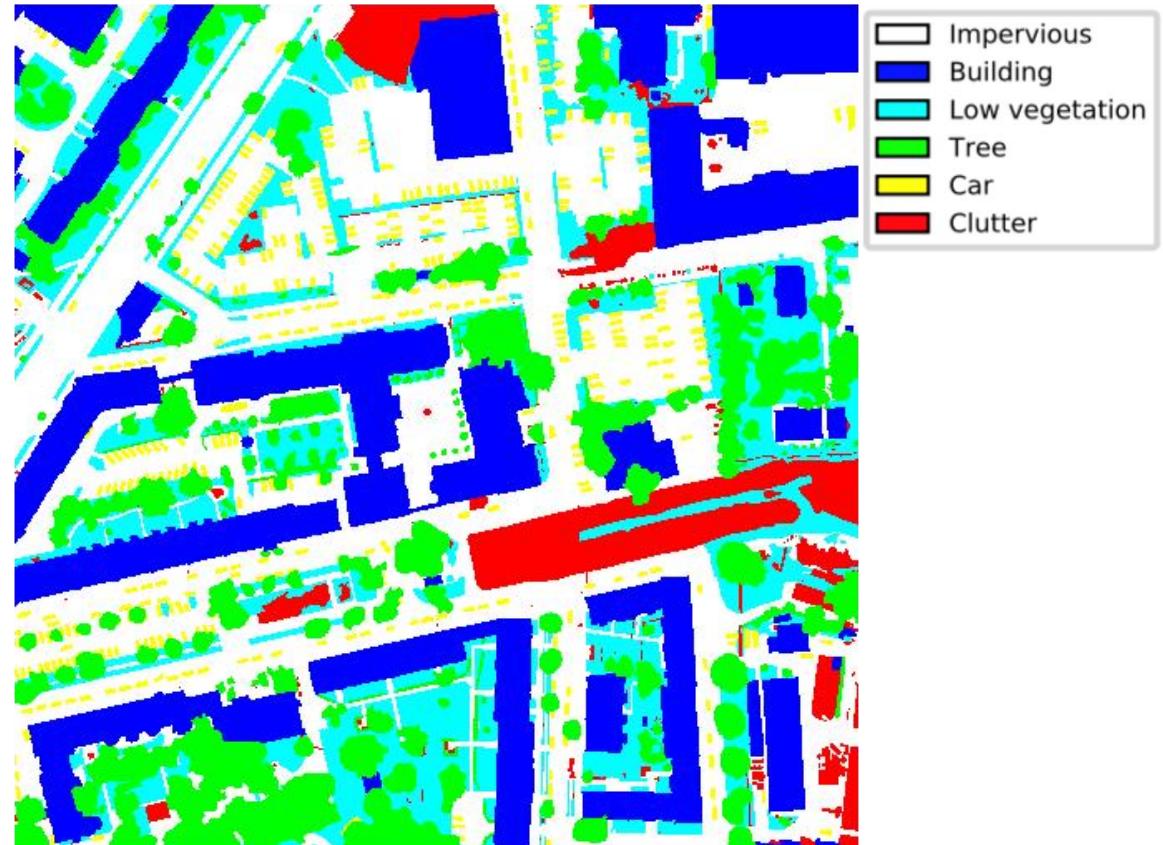
Supervised Semantic Segmentation Task Examples: Cityscapes Dataset



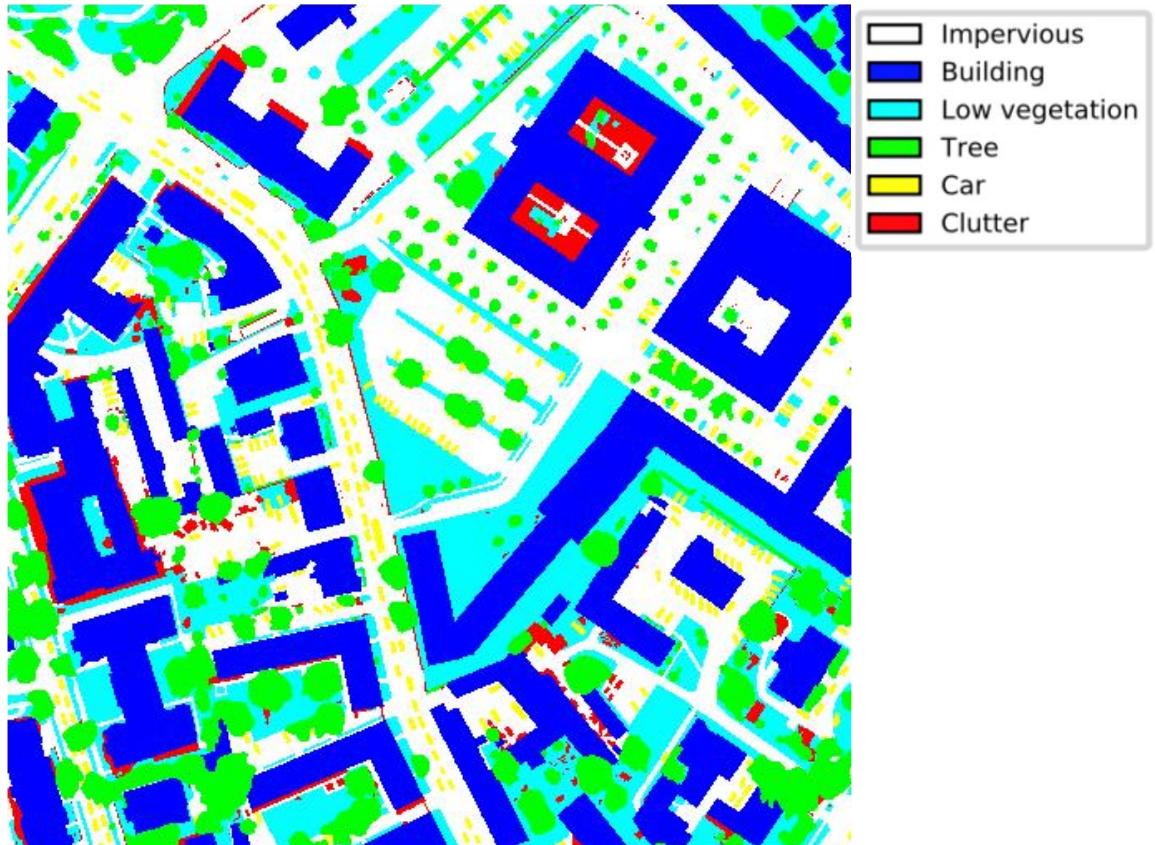
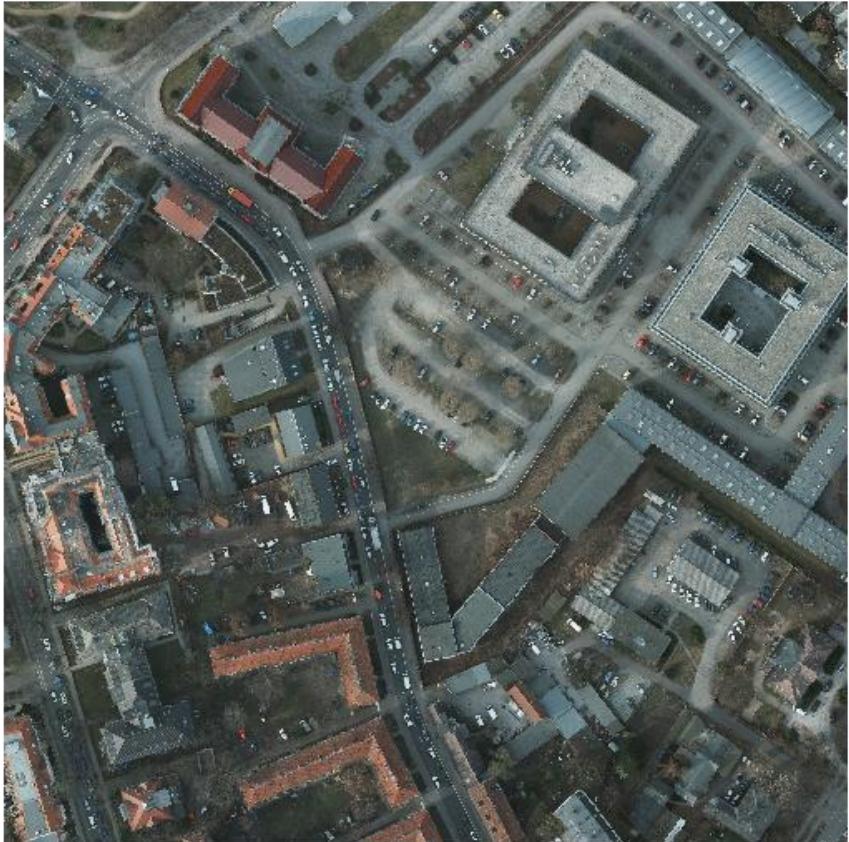
Supervised Semantic Segmentation Task Examples: ISPRS dataset (Potsdam)



Supervised Semantic Segmentation Task Examples: ISPRS dataset (Potsdam)



Supervised Semantic Segmentation Task Examples: ISPRS dataset (Potsdam)



ISPRS dataset details

Potsdam Benchmark dataset for ISPRS
2D Semantic labeling challenge
in Potsdam:

- 38 tiles (6000x6000) at \approx 5cm resolution
 - 24 Tiles dedicated for training
 - 14 Tiles for validation
- 4-bands IRRGB + DSM + NDSM
- Ground truth mask for 6 classes:
 - Building
 - Impervious Surfaces
 - Low Vegetation Tree
 - Tree
 - Car
 - Background



Development environment

- Python 3.5
- Imageio 2.6
- Deep Learning Libraries:
 - Low-Level API: Tensorflow 2.0
(with eager_execution enabled)
 - High-Level API: Keras 2.2
 - Input pipeline API: Tensorflow.data
 - Monitoring API: TensorBoard
- Infrastructure:
 - 16-Core, 64GB RAM
 - Nvidia 16GB GPU (Tesla P100)
 - VM instance on GCP



TensorFlow



Keras



NVIDIA®

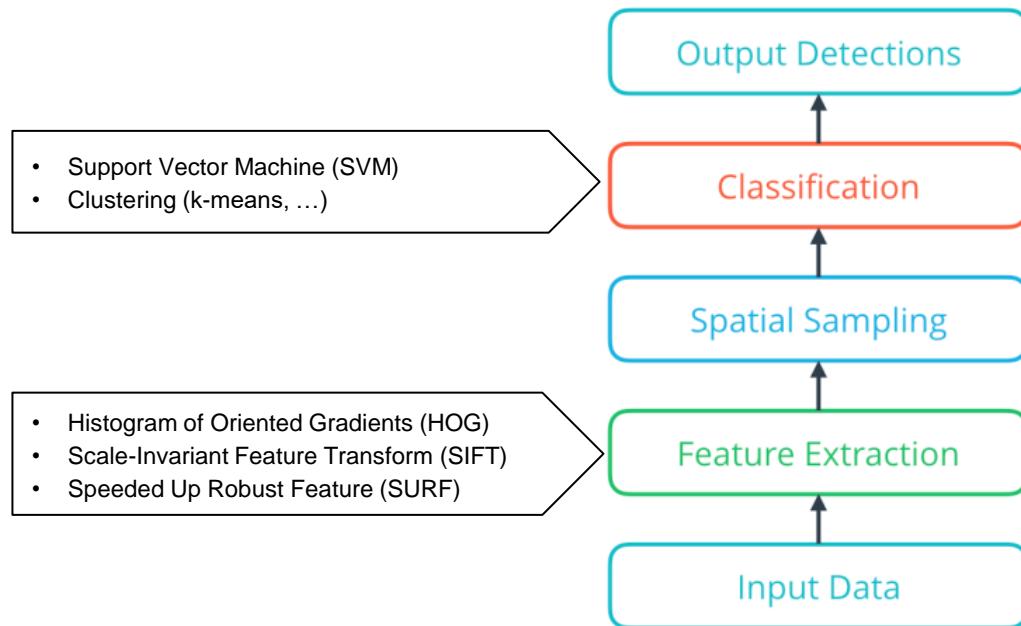
TESLA®



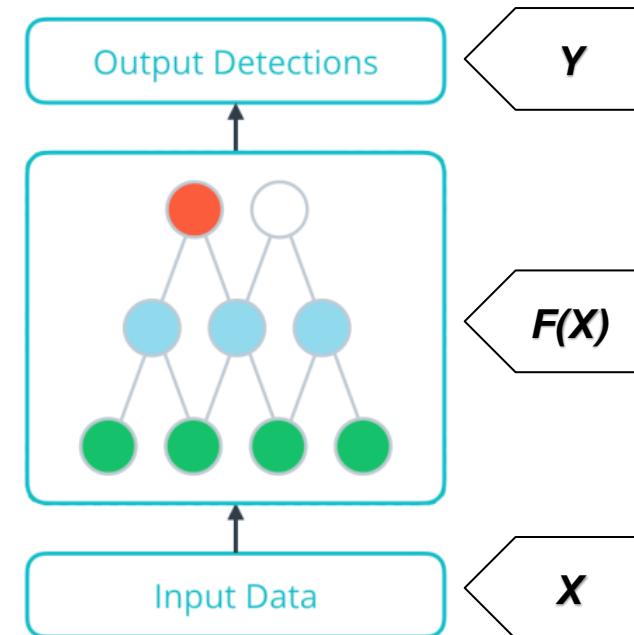
Google Cloud Platform

AIRBUS

Computer Vision Segmentation Approaches:

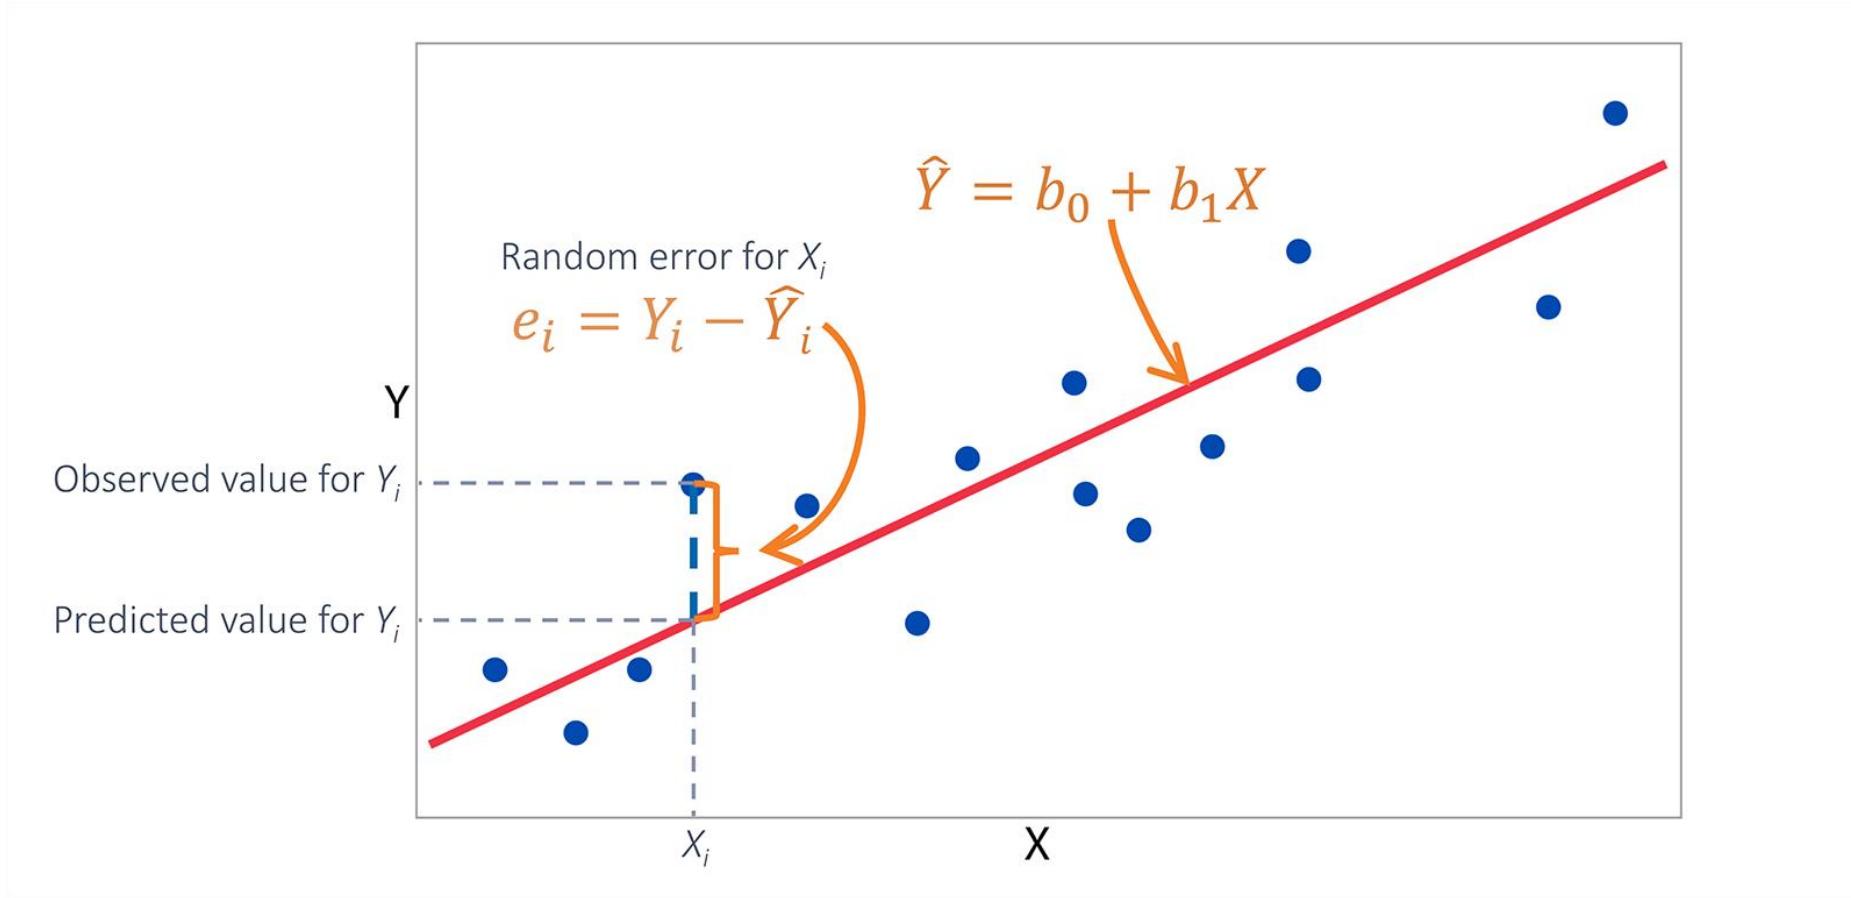


Traditional Computer Vision



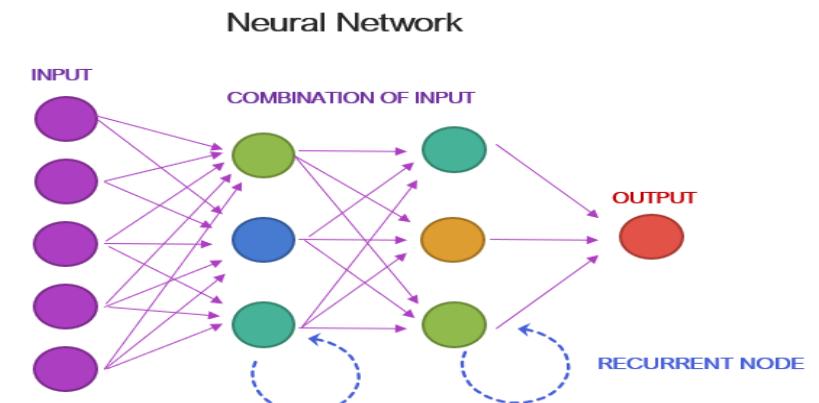
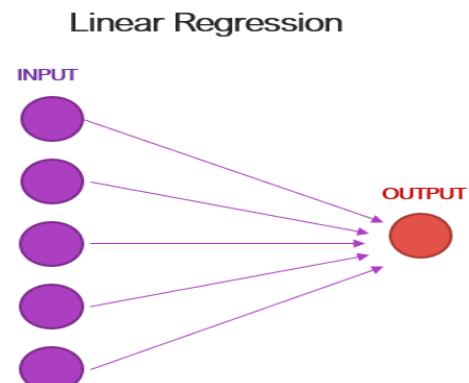
Deep Neural Network

Parametric Modeling: Linear Regression

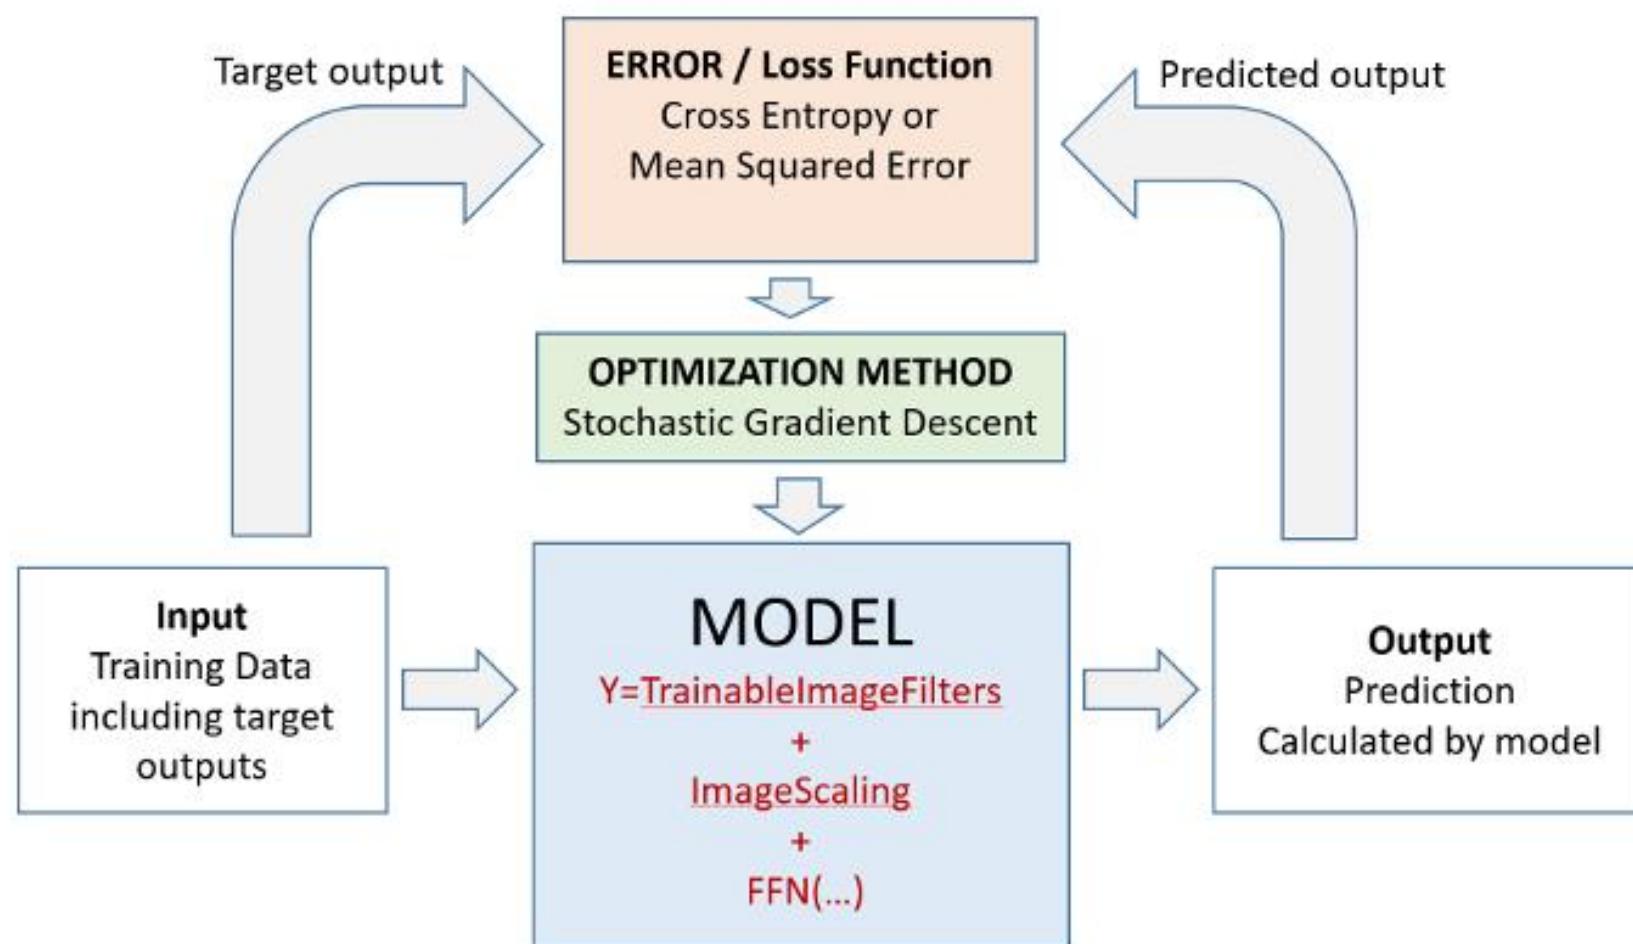


Linear Regression Vs Neural Networks:

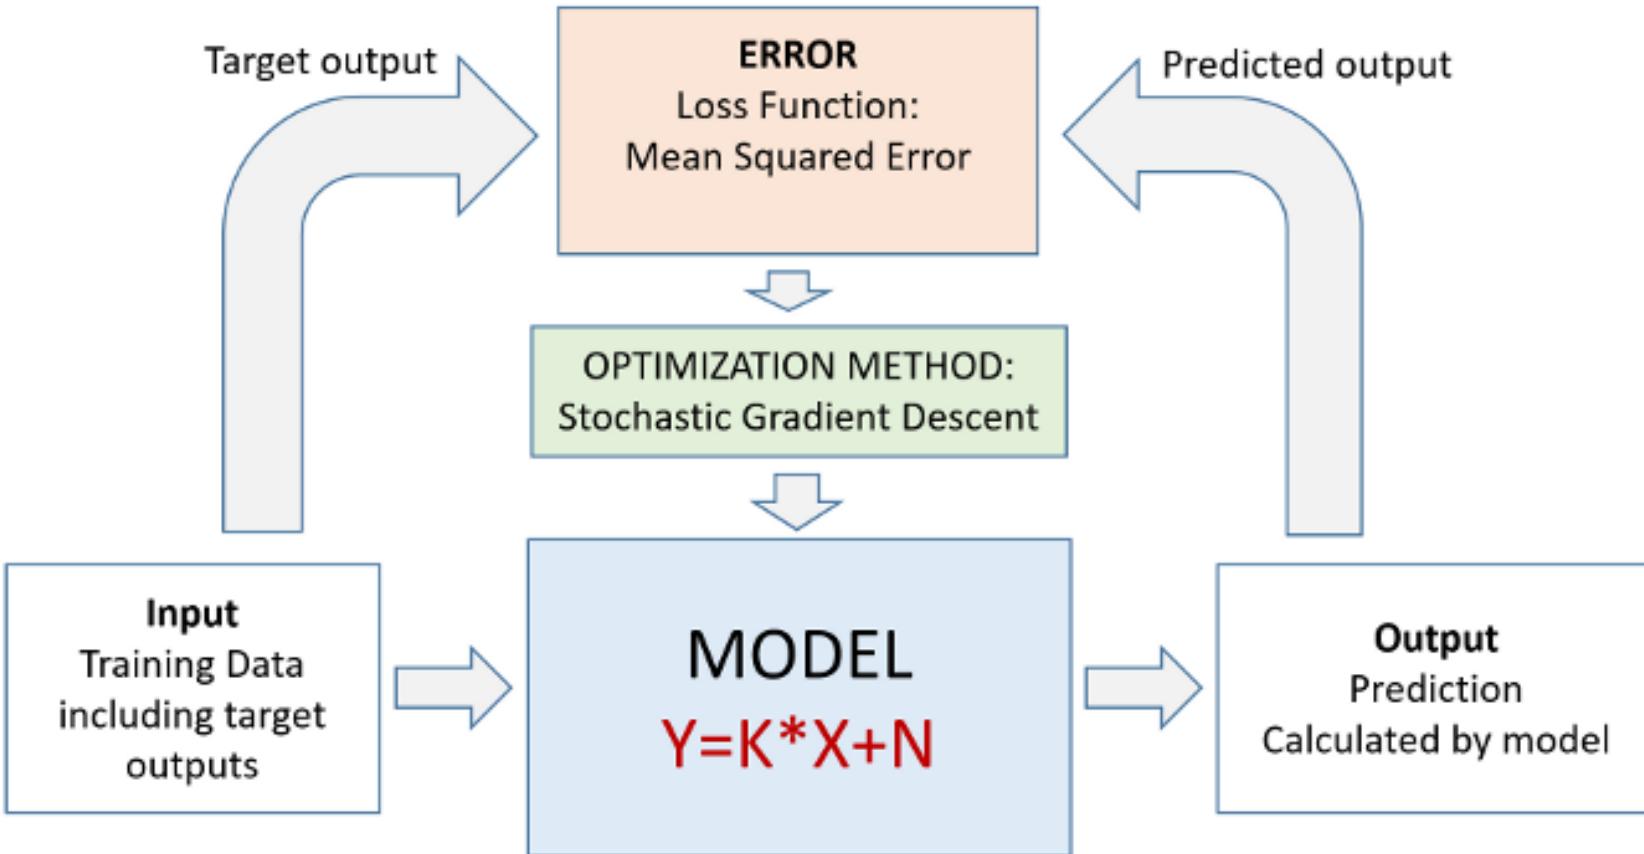
	Linear Regression	Neural Networks
Supervised Learning	Explanatory variables X Response Variables Y	Input Aerial image Pixels X Reference Mask Pixels Y
Parametric Model	$Y = b_0 + b_1 X$	Pixel Label = $F(X W_1 W_2 W_3 W_4 \dots)$
Minimizing Loss Function (Error)	Mean Squared Error	Cross-entropy
Optimization Algorithm	Analytical Solution	Gradient Descent
Goodness of fit Metrics	R ² , F-test, p-value ...	Accuracy, Jaccard Index, ...



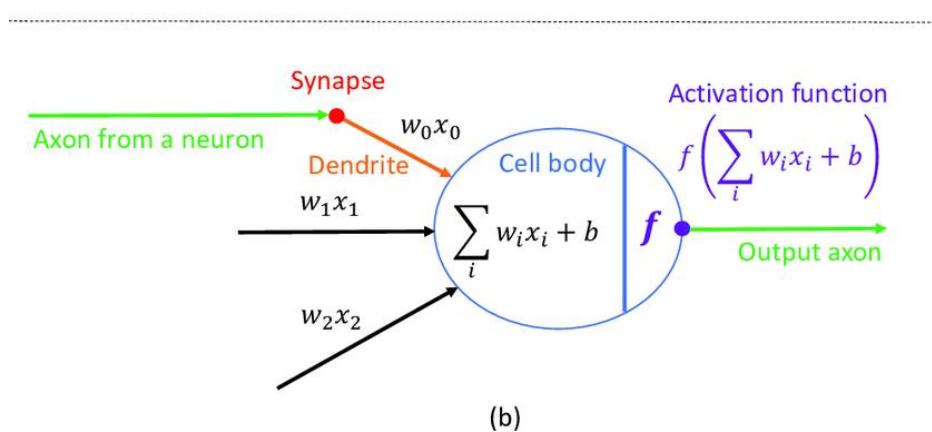
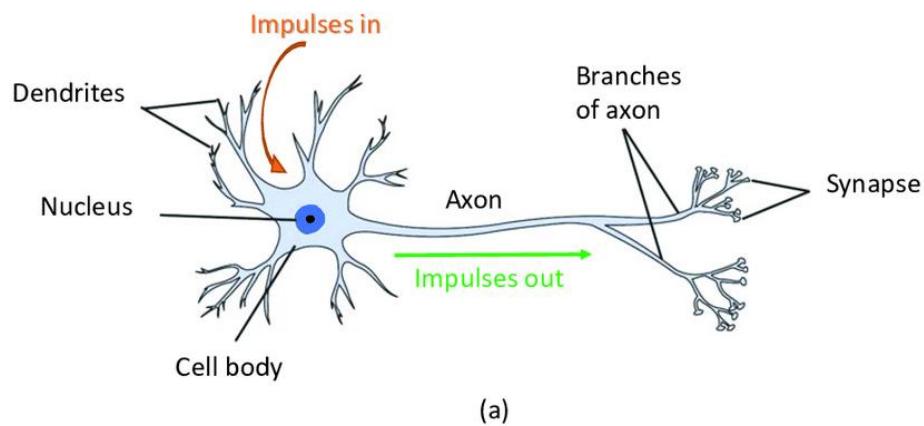
Neural Network Training Flowchart



Linear Regression Training Flowchart:



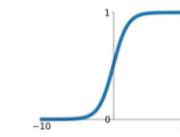
Neural Networks Inspiration:



Activation Functions

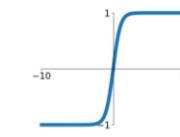
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



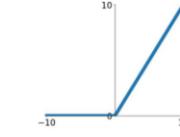
tanh

$$\tanh(x)$$

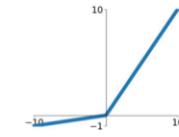


ReLU

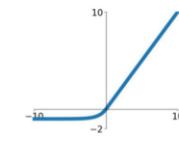
$$\max(0, x)$$



Leaky ReLU
 $\max(0.1x, x)$



Maxout
 $\max(w_1^T x + b_1, w_2^T x + b_2)$

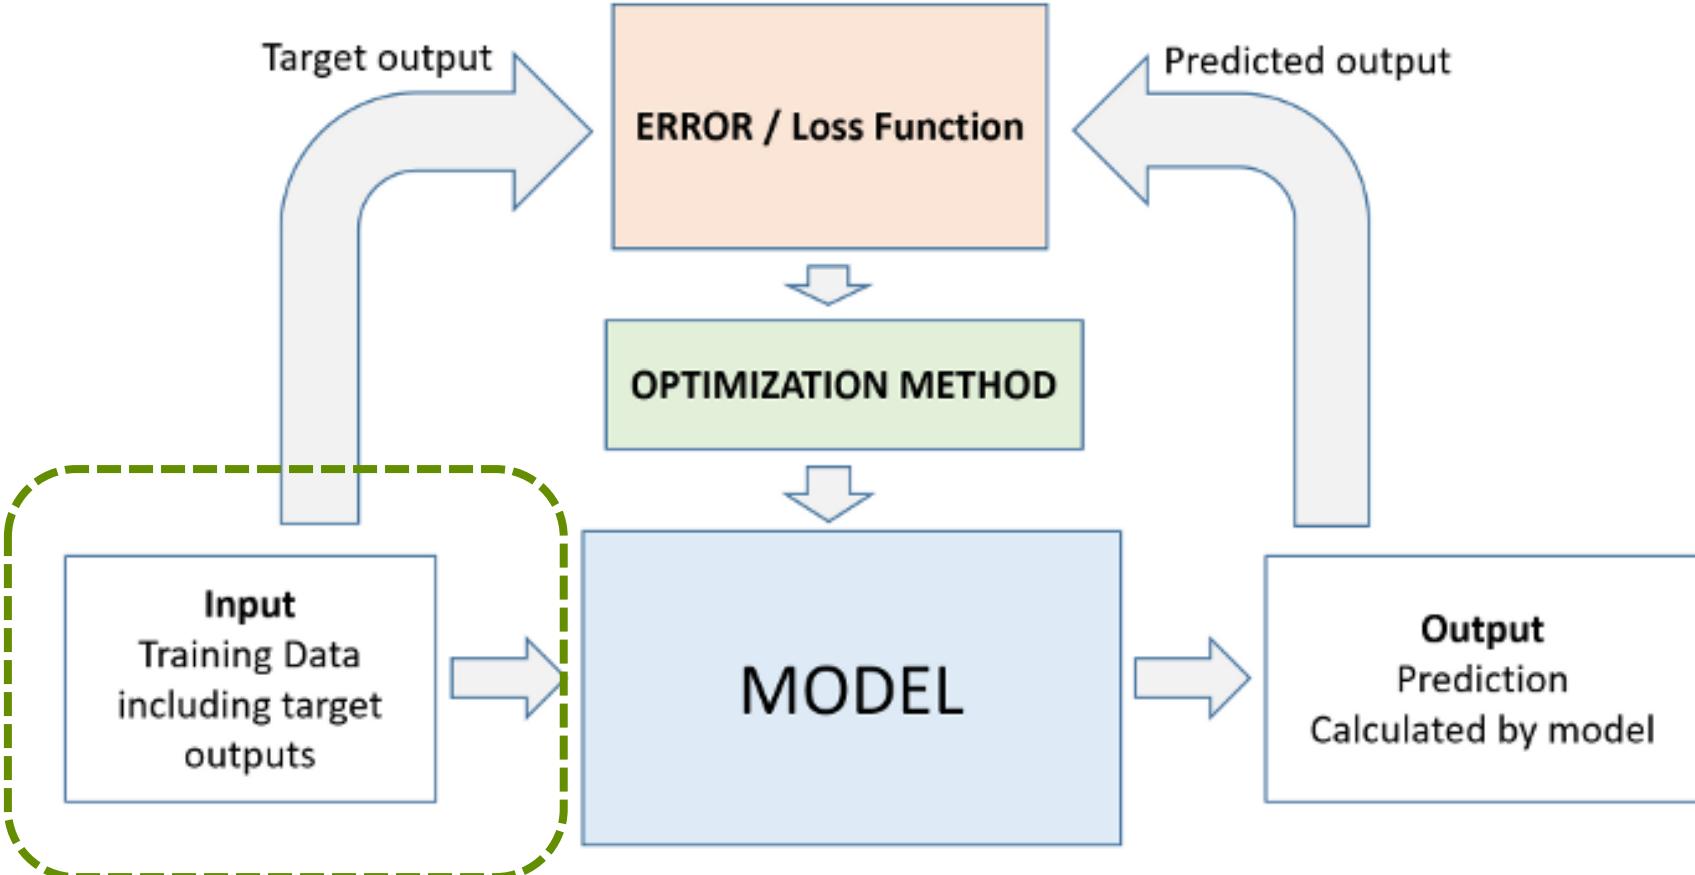


ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$

Input Pipeline: Tensorflow.data

Input Pipeline

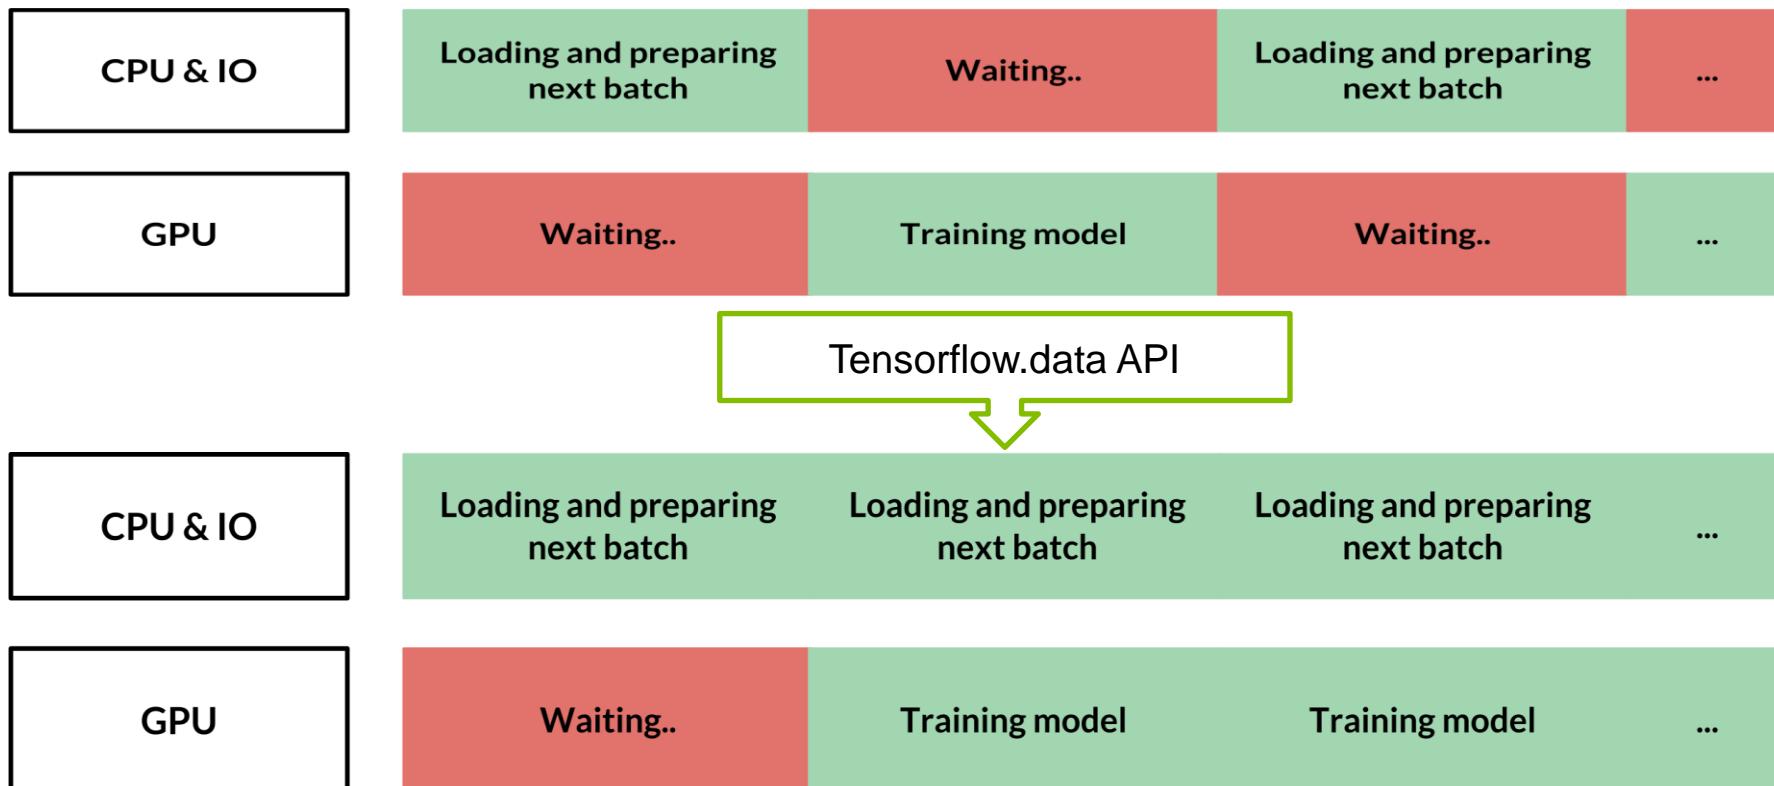


Terminology:

- **Training Dataset:** The dataset that is used to compute the Loss function gradient to update the Models weights
- **Validation Dataset:** The dataset that is used to compute the Generalization Performance of the Model (Metrics)
- **Epoch:** One Epoch is when an ENTIRE dataset is passed forward and backward through the neural network only ONCE.
- **Batch size:** Total number of training examples present in a single batch.
- **Example:** a dataset made of 100 images with Batch size of 20 images => 1 epoch consists of 5 Batch iterations

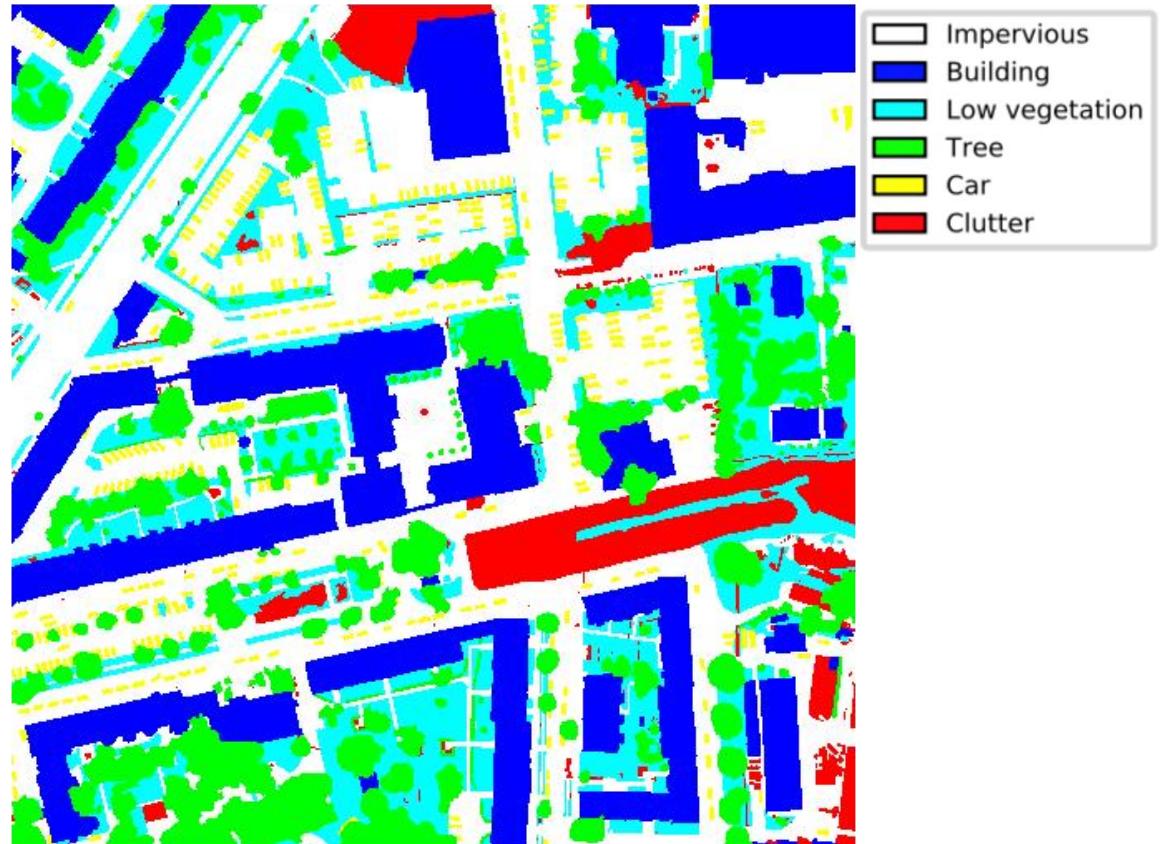
Input Pipeline challenges: avoid CPU-GPU bottleneck

- Goal: Decoupling data preprocessing from Accelerator Training

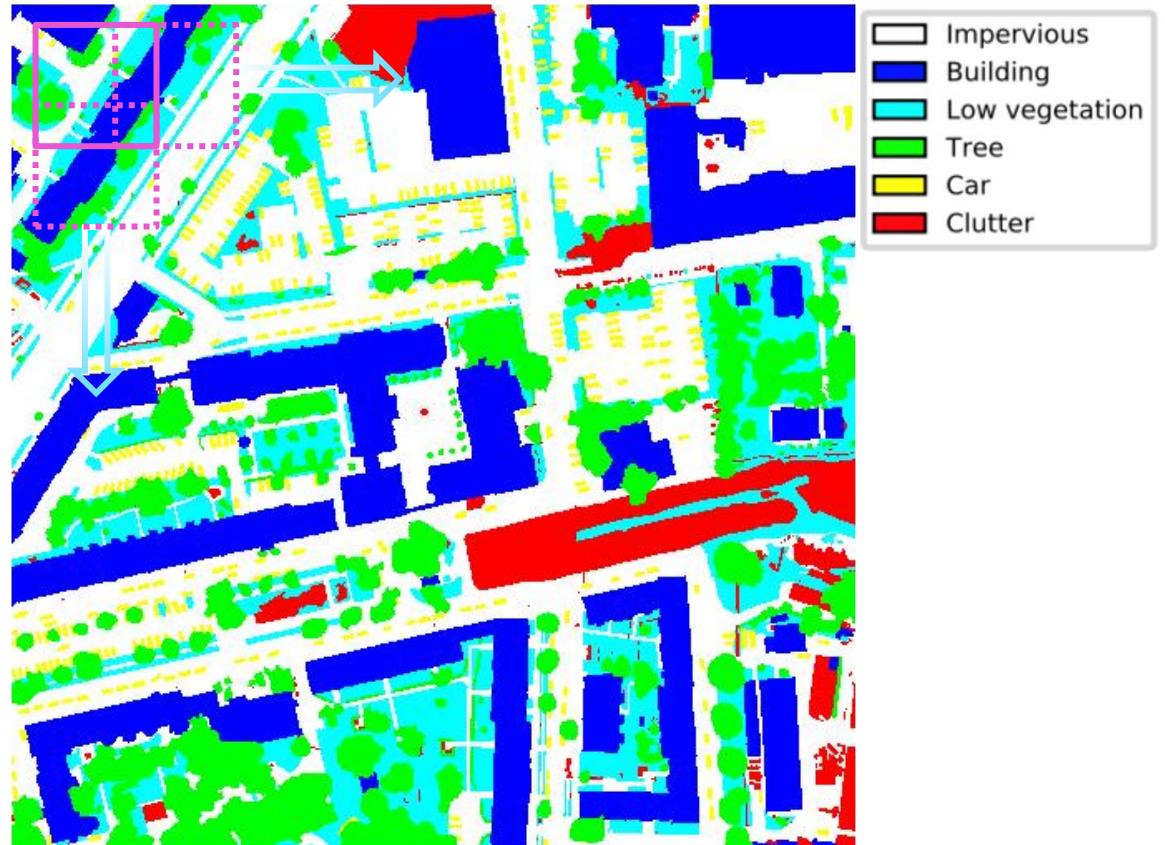


(Simplified, in reality there is always some idle time on both devices)

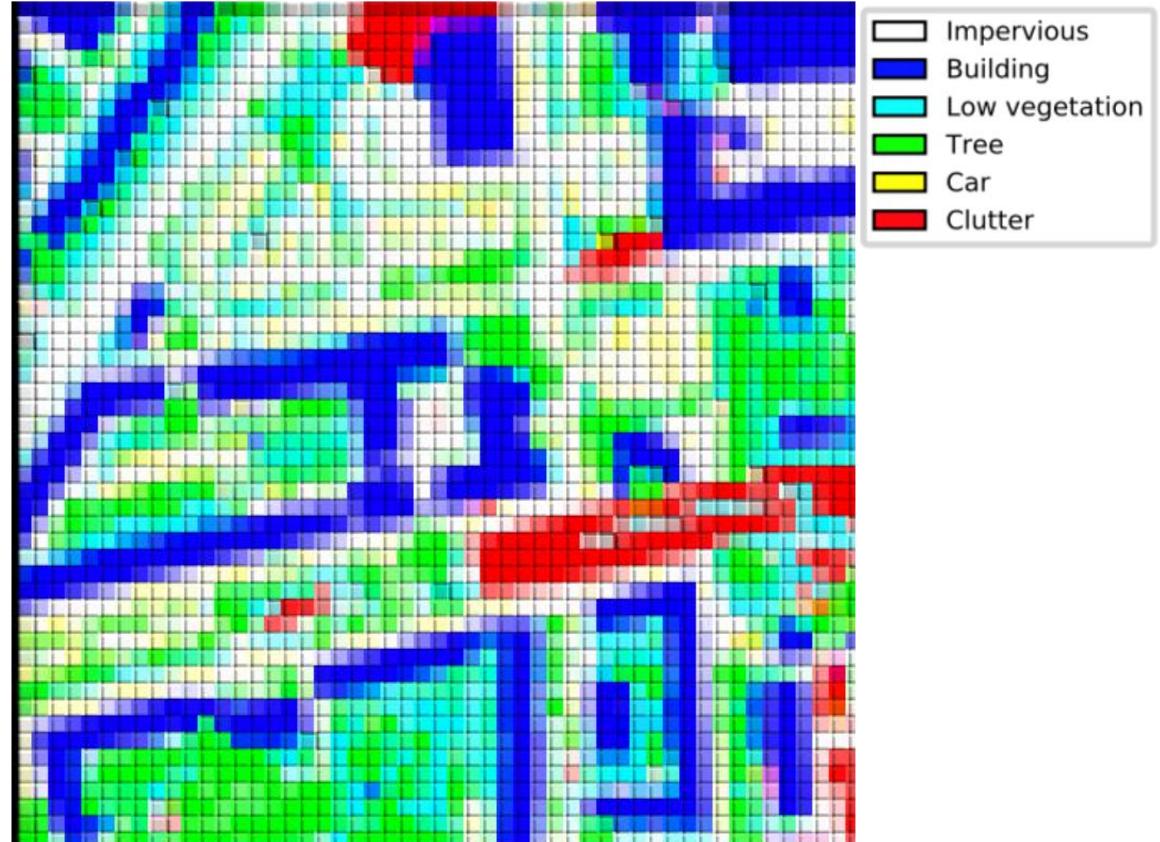
Patch Extraction



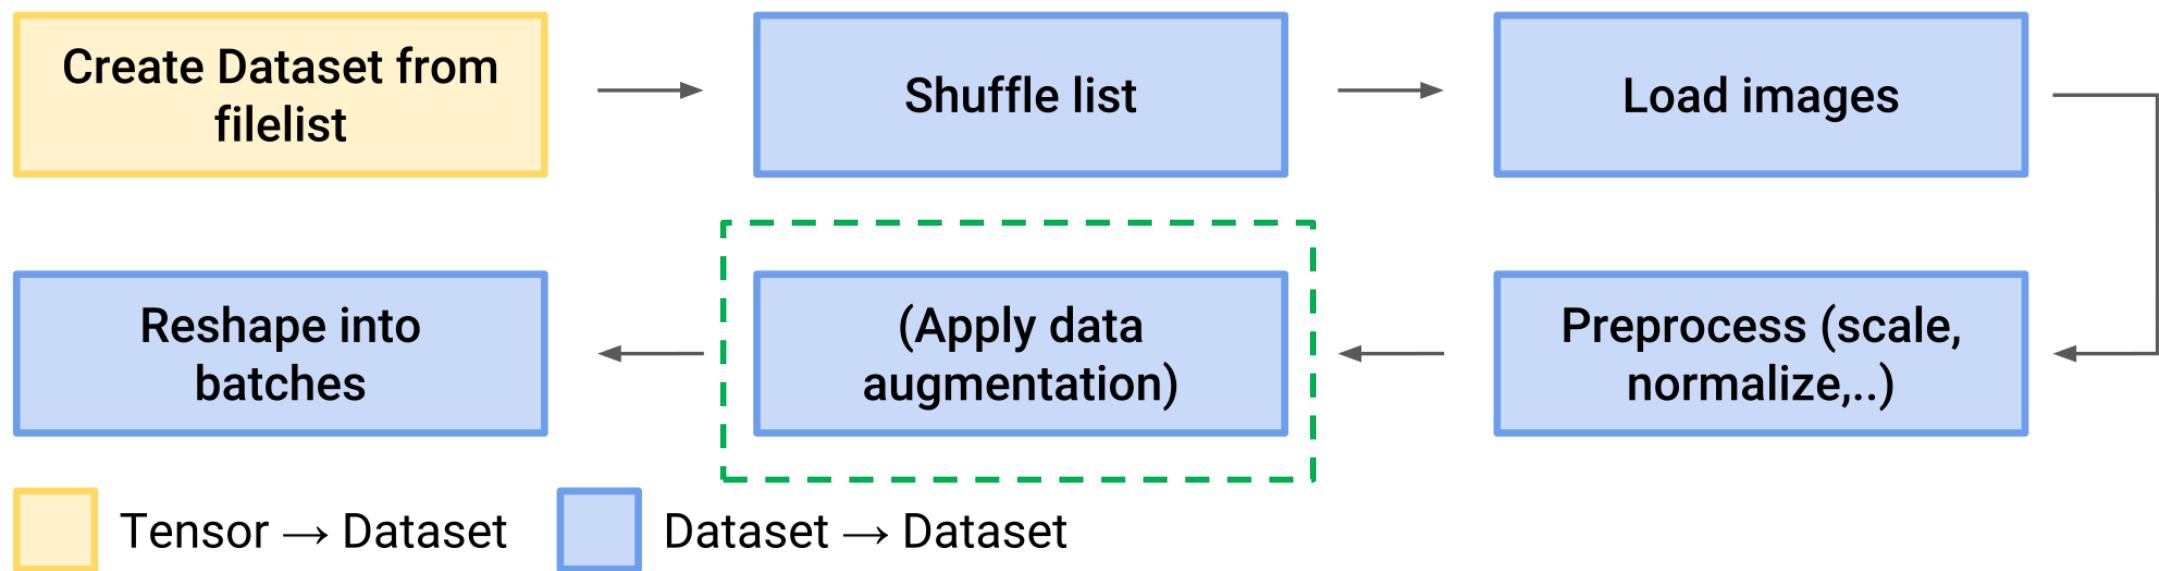
Patch Extraction



Supervised Semantic Segmentation Task Examples: ISPRS dataset (Potsdam)



- Practical advantage: Easy ETL (Extract-Load-Transform) chaining for complex data augmentation



Input Pipeline: Data Augmentation

- More data helps Model Performances
- Data Augmentation is a way to artificially produce more Annotated data at low cost to train the Model
- Apply same spatial transformation to both aerial and mask patch pair, color transformation only to aerial patch



Input Pipeline: implementation details

```
with tf.device('/CPU:0'):
    def get_dataset(img_dir, mask_dir, batch_size, shuffle=True, repeat=False, augment=True, scaled=True, cachefile=None):
        images_path = Path(img_dir)
        masks_path = Path(mask_dir)
        images_files = sorted([str(x) for x in images_path.iterdir() if x.is_file() and x.suffix == '.tif'])
        masks_files = sorted([str(x) for x in masks_path.iterdir() if x.is_file() and x.suffix == '.tif'])

        # checking the equality of the two list lengths
        assert len(images_files) == len(masks_files), " (Error) The Aerial image count does not match Mask counts!"
        len_images = len(images_files)

        if cachefile:
            cache_file_name = '/tmp/' + cachefile
        else:
            cache_file_name = ""

        # images file names dataset
        images_ds = tf.data.Dataset.from_tensor_slices(images_files)
        masks_ds = tf.data.Dataset.from_tensor_slices(masks_files)

        .
```

Input Pipeline: implementation details

```
with tf.device('/CPU:0'):
    def get_dataset(img_dir, mask_dir, batch_size, shuffle=True, repeat=False, augment=True, scaled=True, cachefile=None):
        .
        .

        # zipping the aerial and mask dataset
        patch_pair_ds = tf.data.Dataset.zip((images_ds, masks_ds))
        if shuffle:
            patch_pair_ds = patch_pair_ds.shuffle(len_images)
        patch_pair_ds = patch_pair_ds.map(tf_get_image, num_parallel_calls=AUTOTUNE).prefetch(AUTOTUNE)

        # start making patches
        if scaled:
            patch_pair_ds = patch_pair_ds.map(make_patches_scaled, num_parallel_calls=AUTOTUNE)
        else:
            patch_pair_ds = patch_pair_ds.map(make_patches, num_parallel_calls=AUTOTUNE)
        patch_pair_ds = patch_pair_ds.apply(tf.data.Dataset.unbatch).cache(filename=cache_file_name)

        if augment:
            patch_pair_ds = patch_pair_ds.map(make_augmentation, num_parallel_calls=AUTOTUNE).prefetch(AUTOTUNE)
        else:
            patch_pair_ds = patch_pair_ds.map(no_augmentation, num_parallel_calls=AUTOTUNE).prefetch(AUTOTUNE)
        # repeat is not necessary as keras fit method will use epochs parameter for repetition
        if repeat:
            patch_pair_ds = patch_pair_ds.repeat()
        if shuffle:
            patch_pair_ds = patch_pair_ds.shuffle(600)
        patch_pair_ds = patch_pair_ds.batch(batch_size)
        patch_pair_ds = patch_pair_ds.prefetch(AUTOTUNE)

    return patch_pair_ds
```

Input Pipeline Performance

Without Caching:

```
[24]: timepipeline(train_ds, batch_size=BATCH_SIZE)
```



100% 1000/1000 [03:40<00:00, 4.53it/s]

batche size: 50
1000 batches: 220.8209400177002 s
4.52856 Batches/s
226.42780 Images/s

With Disk Caching:

```
[35]: timepipeline(train_ds, batch_size=BATCH_SIZE)
```



100% 1000/1000 [01:14<00:00, 13.48it/s]

batche size: 50
1000 batches: 74.26196789741516 s
13.46584 Batches/s
673.29215 Images/s

With Memory Caching:

```
[38]: timepipeline(train_ds, batch_size=BATCH_SIZE)
```



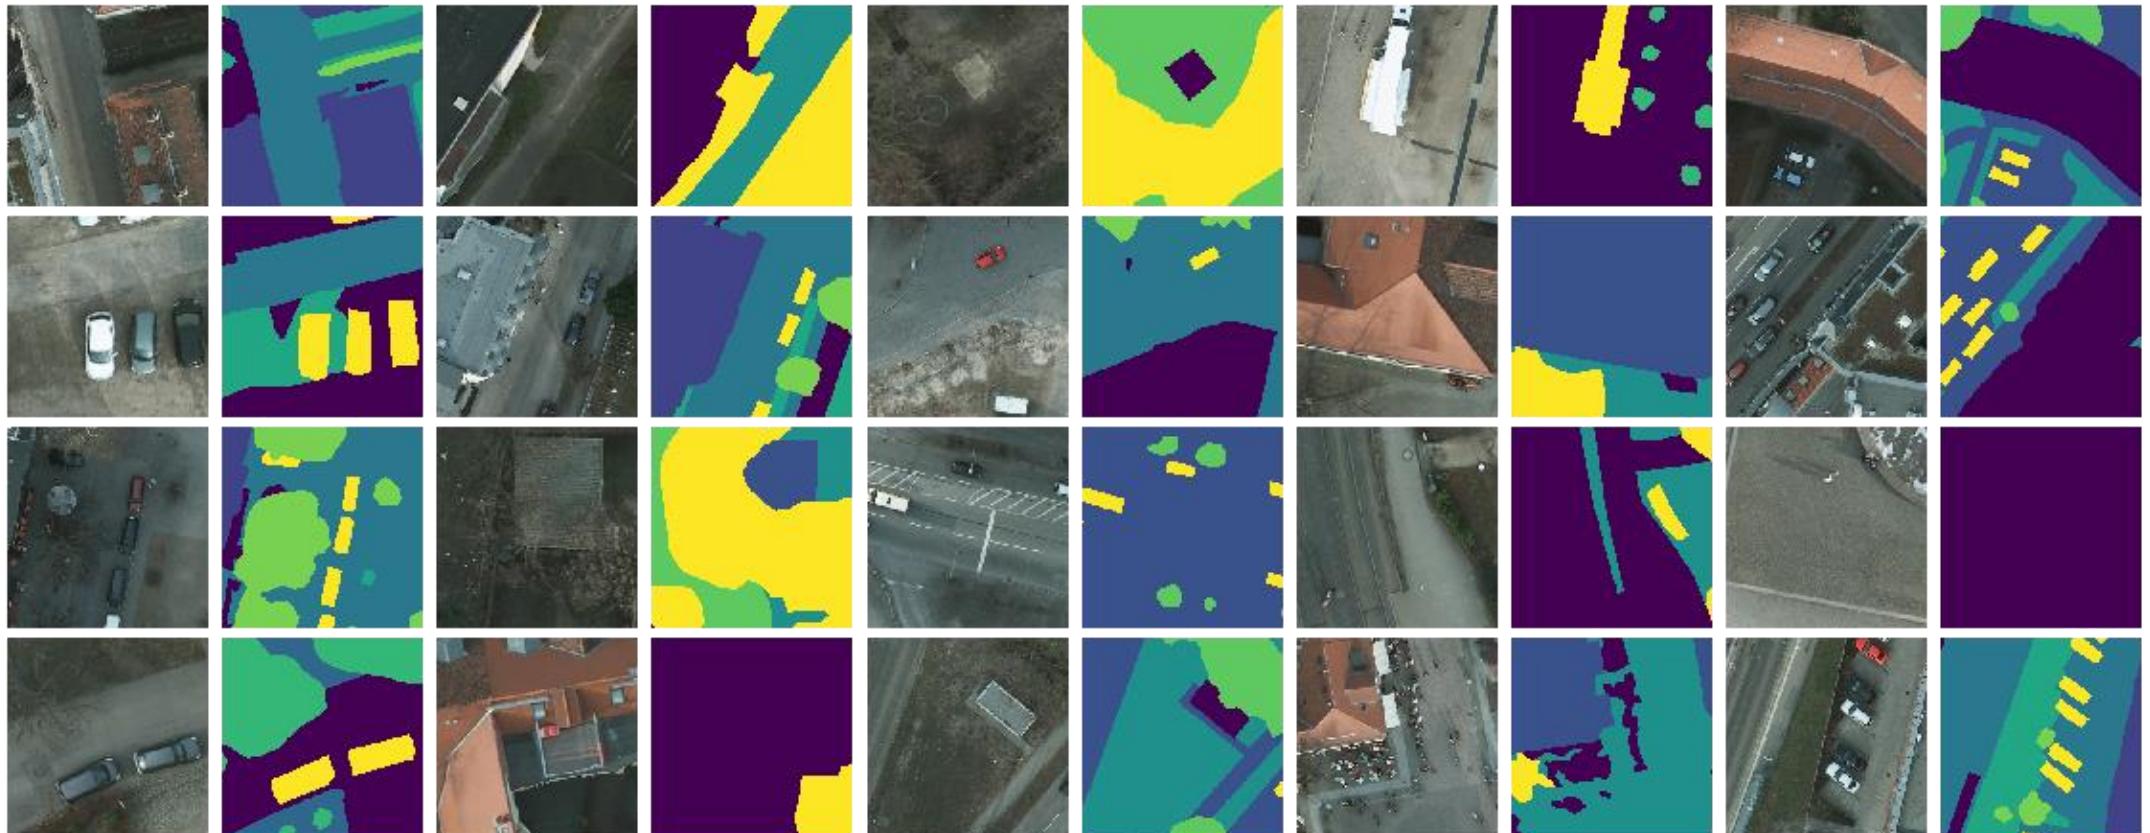
100% 1000/1000 [01:13<00:00, 13.56it/s]

batche size: 50
1000 batches: 73.79464364051819 s
13.55112 Batches/s
677.55595 Images/s

Input Pipeline Patch Pairs Example

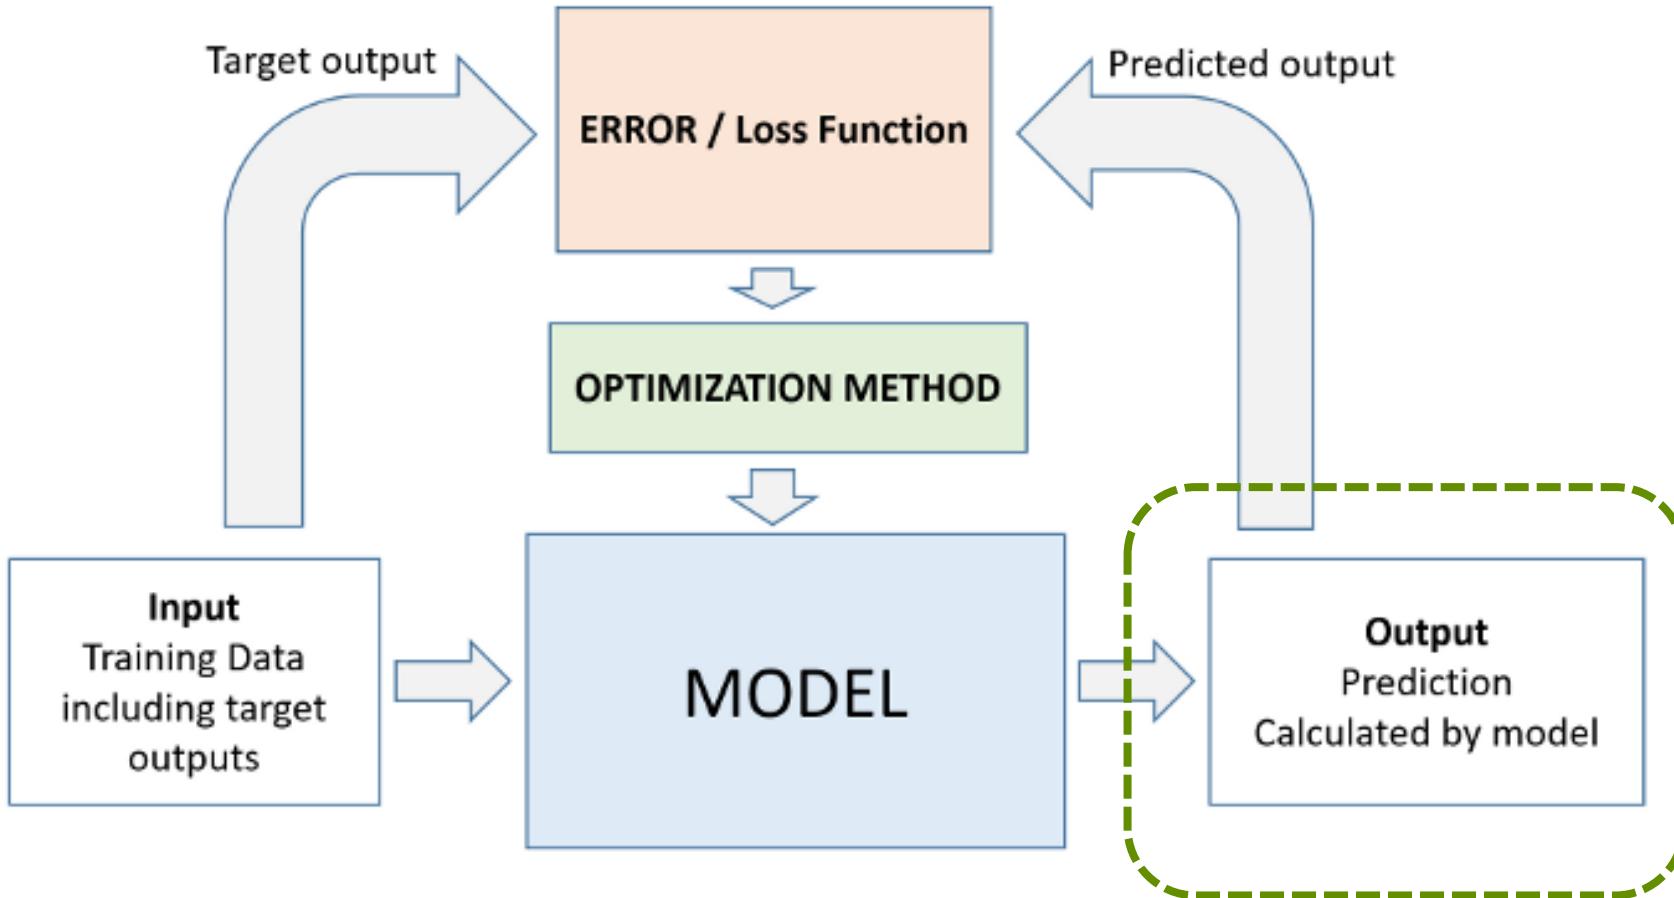
```
[47]: train_ds = get_dataset(img_dir=train_img_dir, mask_dir=train_mask_dir, batch_size=BATCH_SIZE, shuffle=True, repeat=True, augment=True, scaled=False, cachefile=None)
```

```
[ ]: val_ds = get_dataset(img_dir=val_img_dir, mask_dir=val_mask_dir, batch_size=BATCH_SIZE, shuffle=False, repeat=True, augment=False, scaled=False, cachefile=None)
```

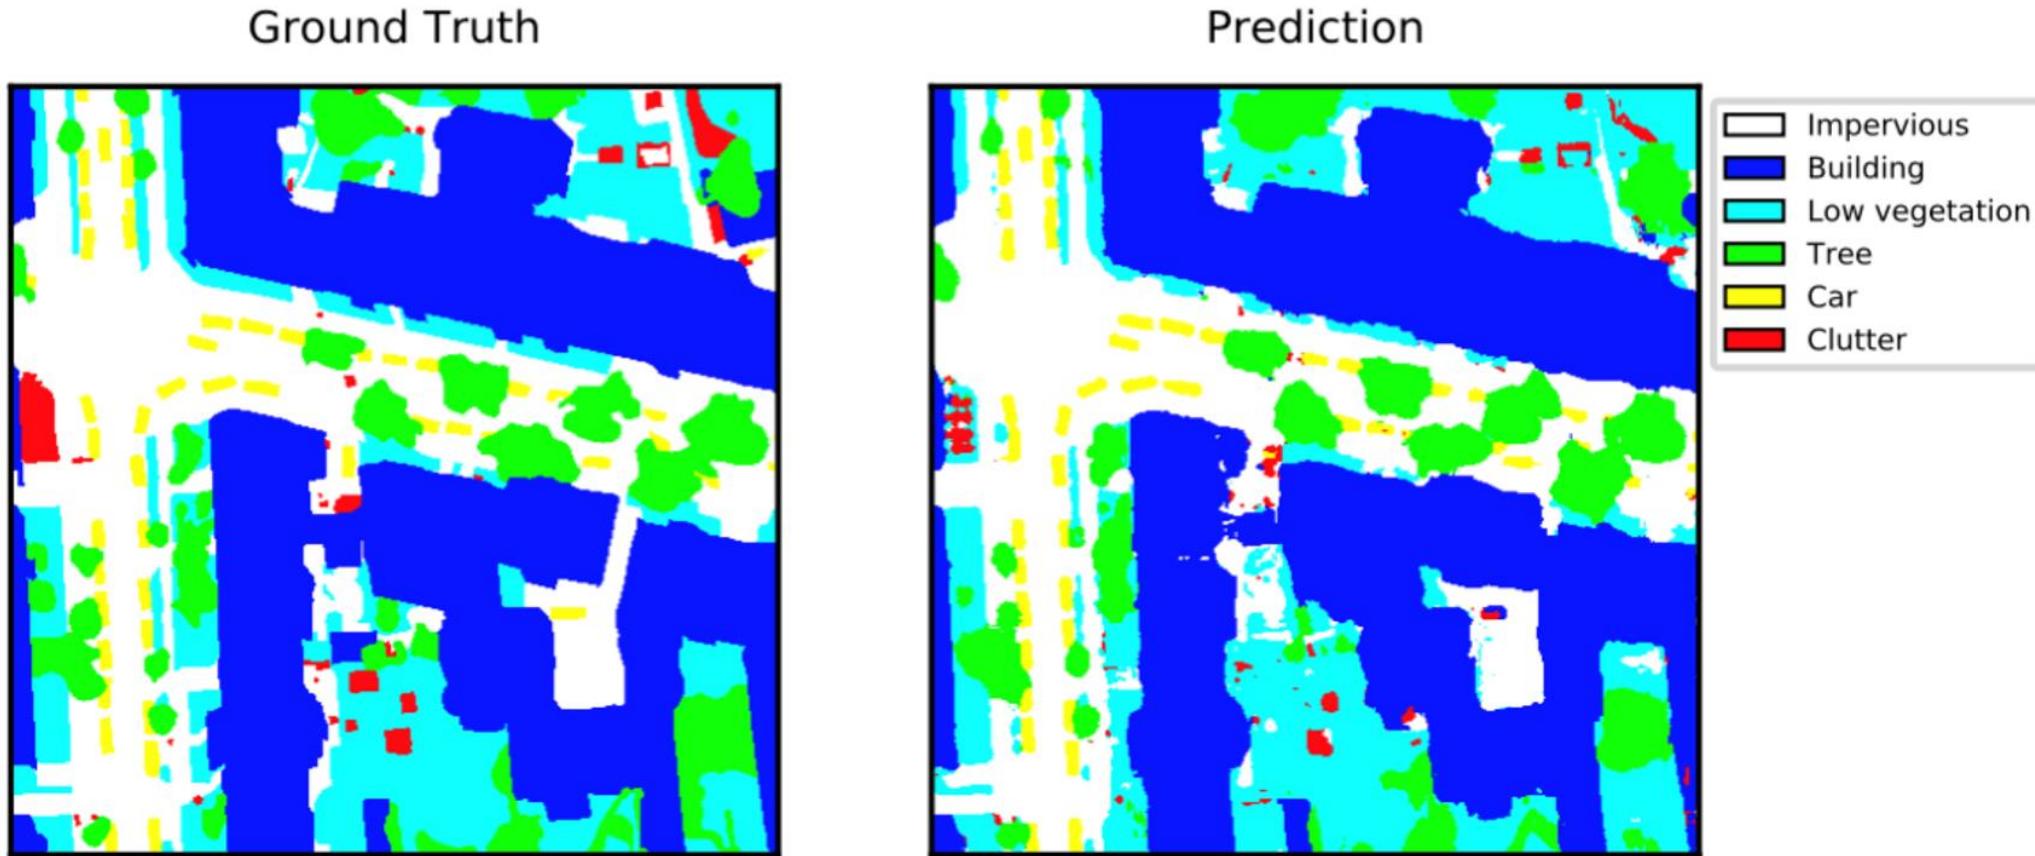


Model Performance Metrics

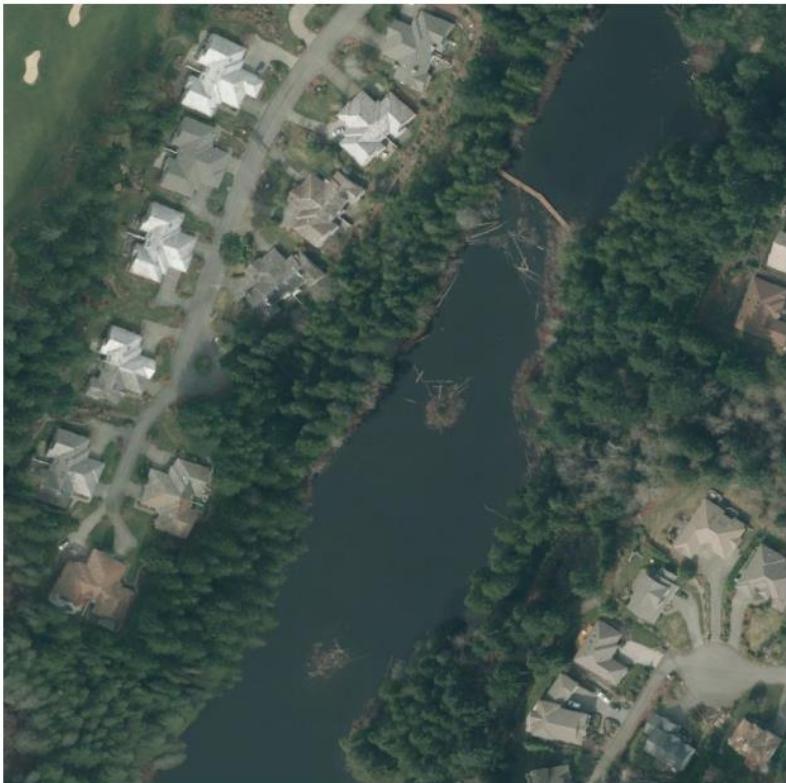
Performance Metrics



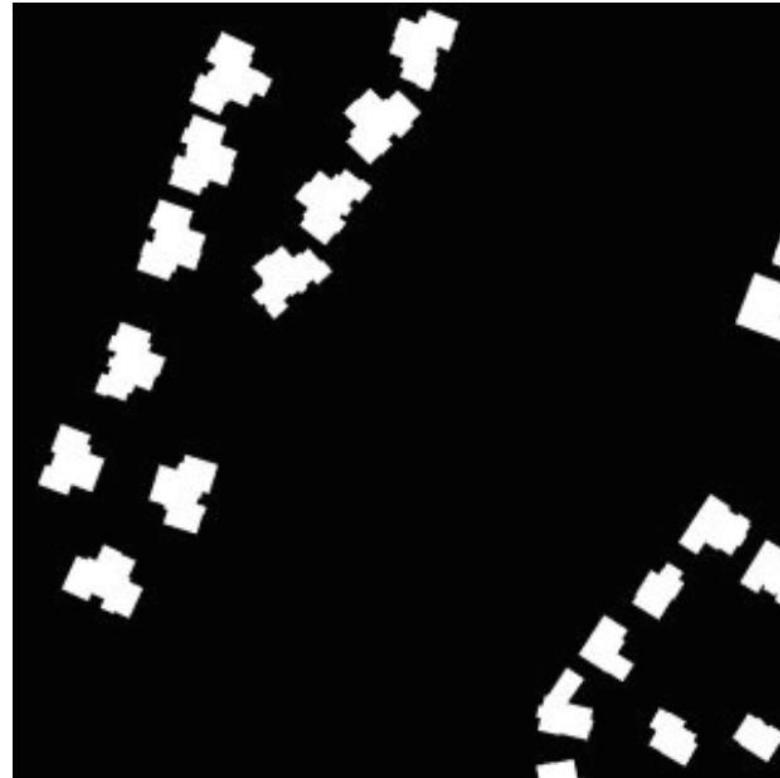
Natural Metric for Image Comparison: Accuracy



Performance Metrics: Accuracy is not enough For Segmentation



Kitsap County, WA



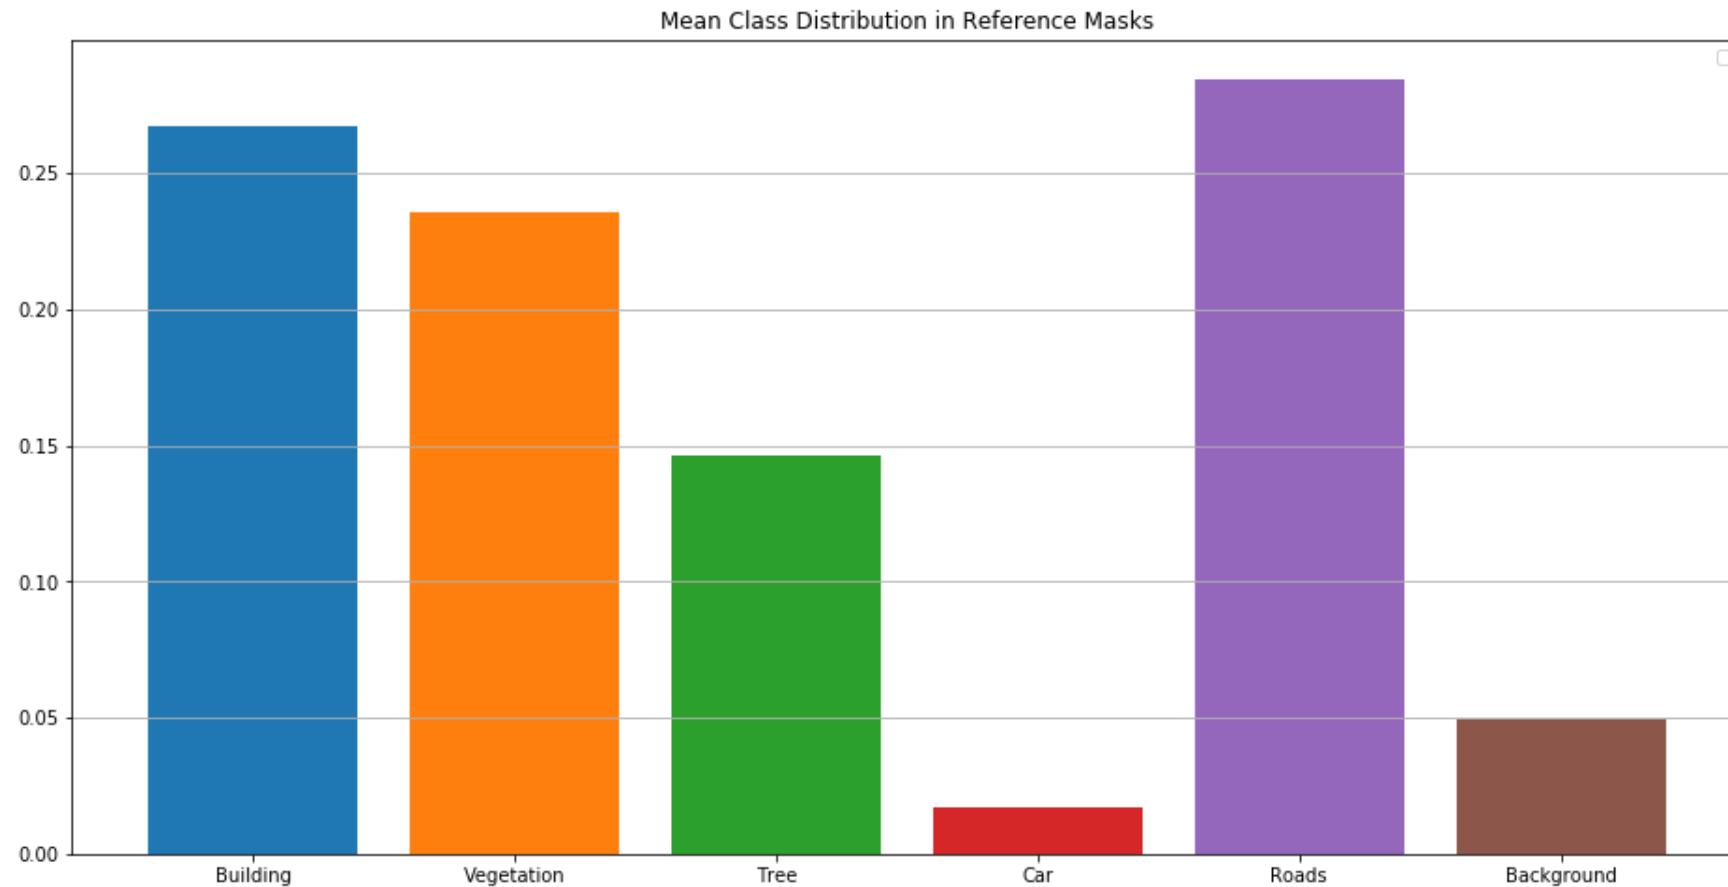
Kitsap County, WA – Reference

≈ 10% Building

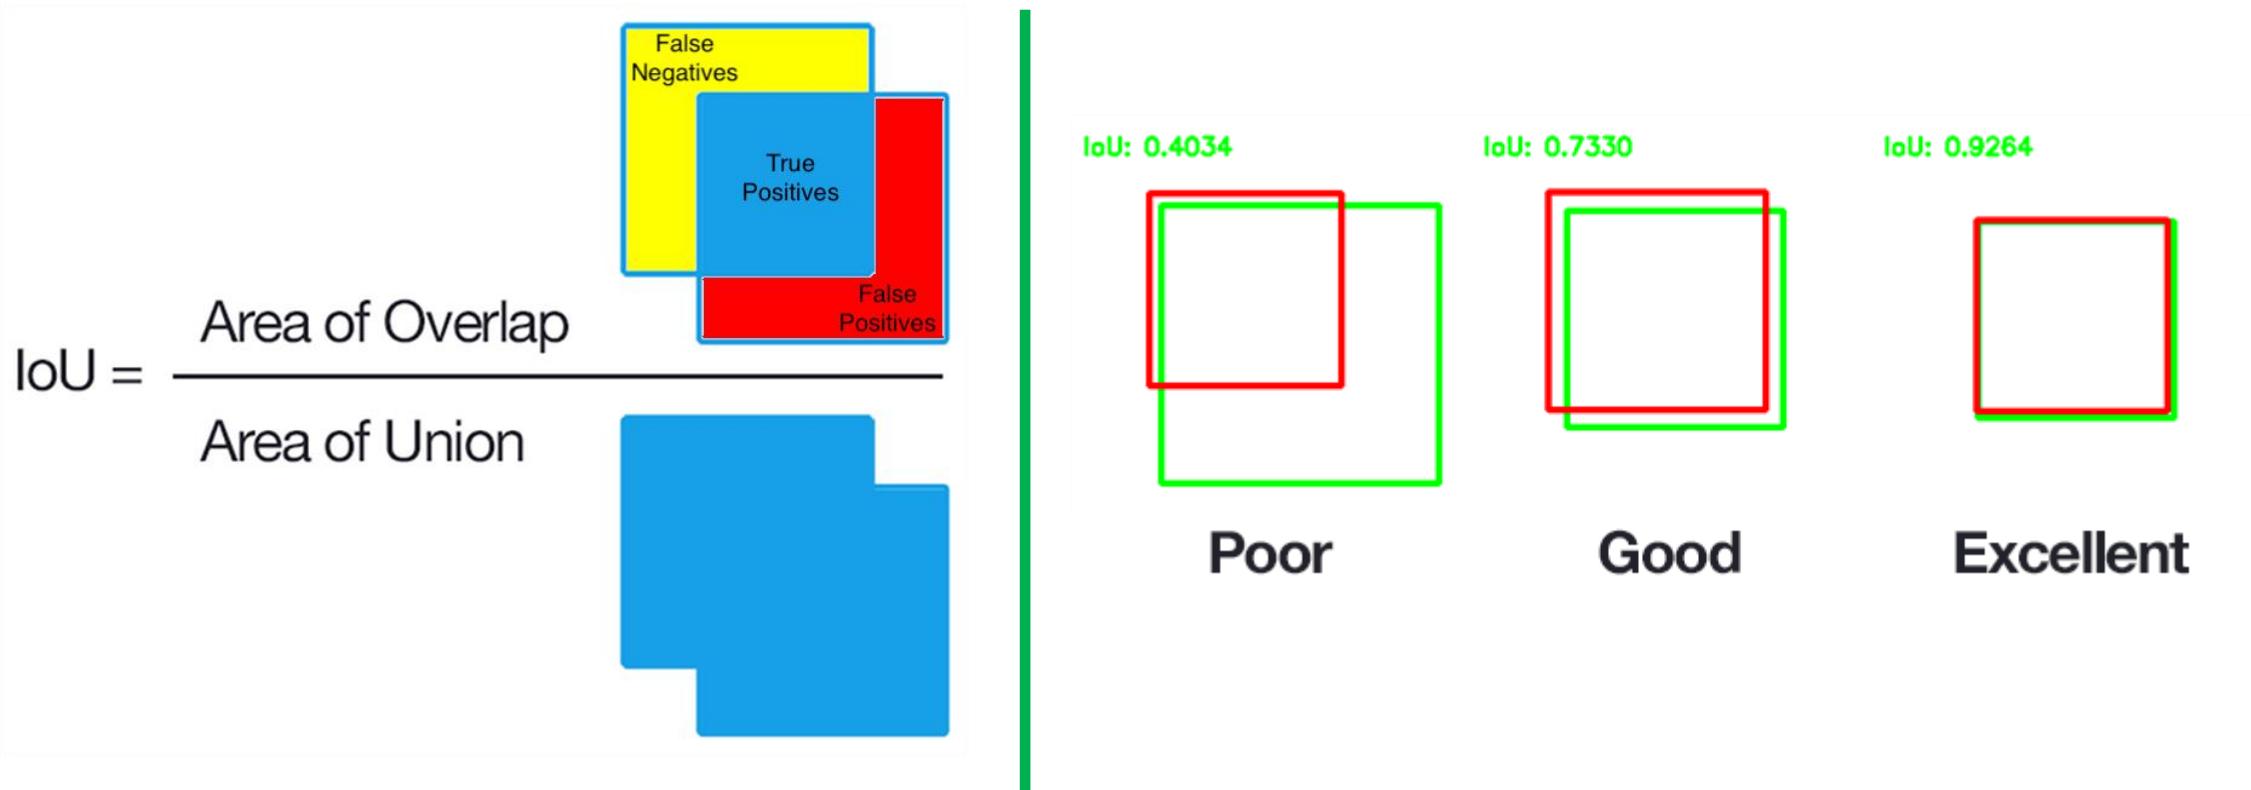
Model = all Background

≈ 90% Accuracy

ISPRS Class Distribution



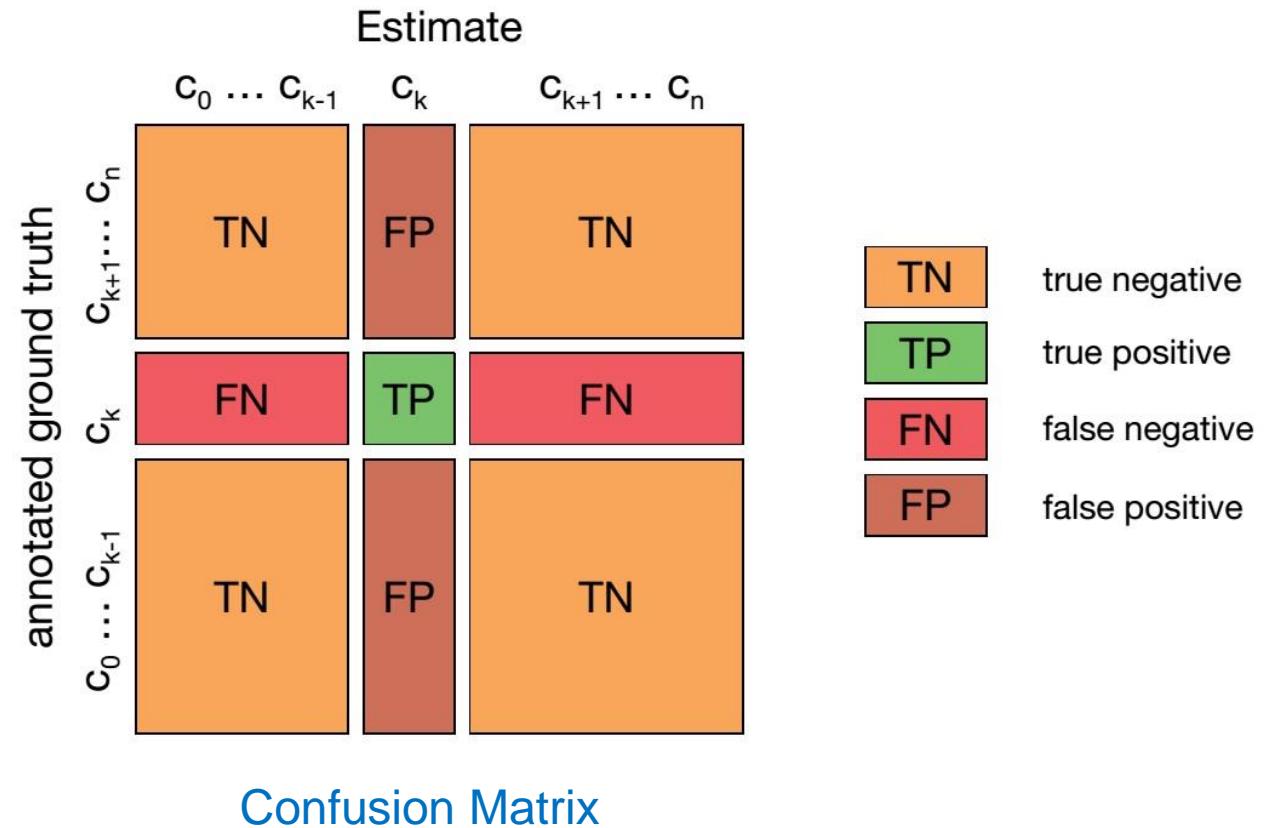
A Better Segmentation Metric: Jaccard Index (Intersection Over Union)



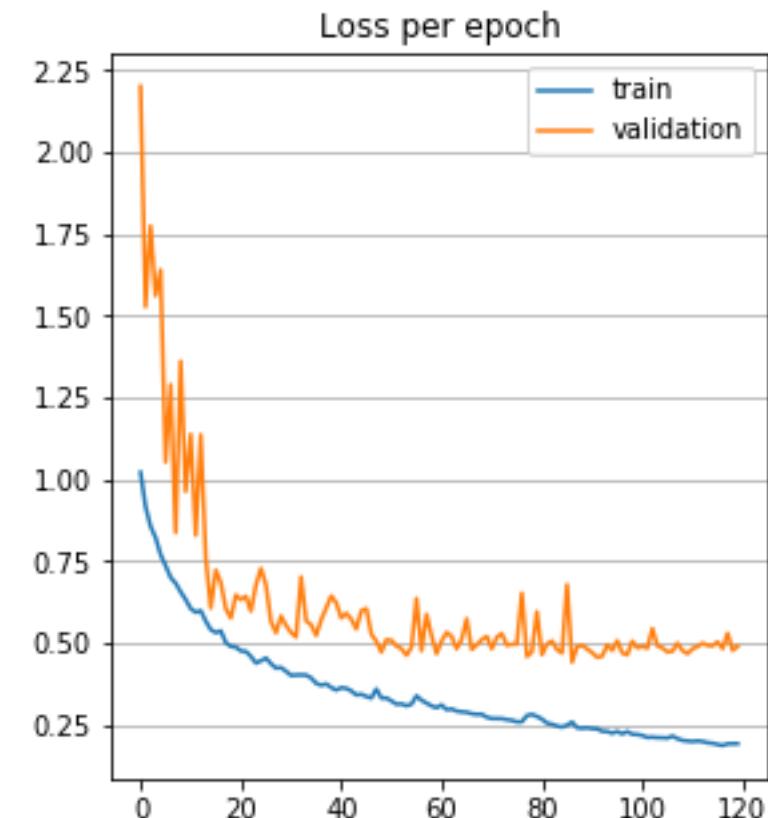
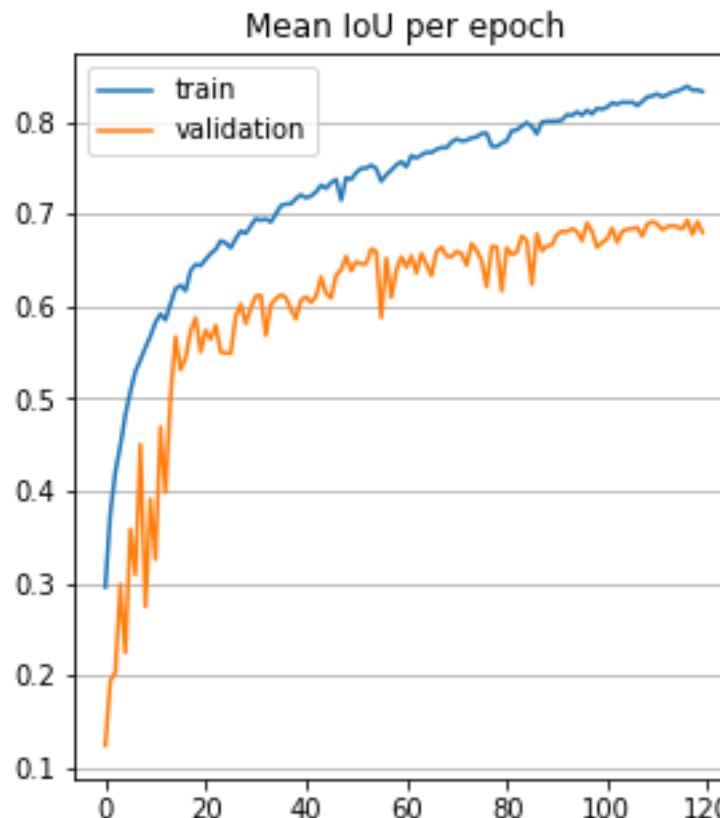
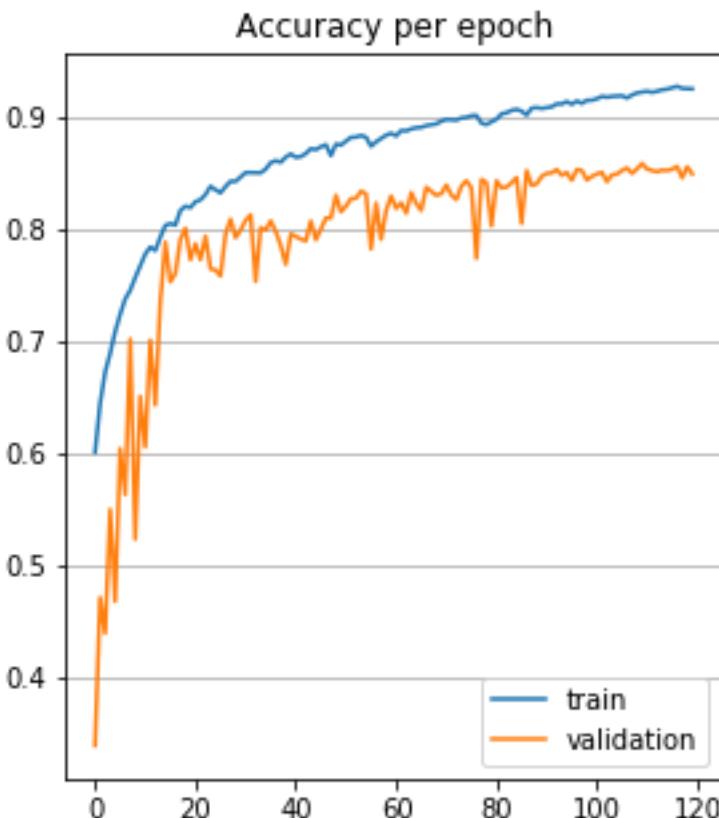
Metrics Formulas:

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + TN + FN}$$

$$\text{Jaccard Index (IoU)} = \frac{TP}{TP + FP + FN}$$



Monitoring Metrics During Training

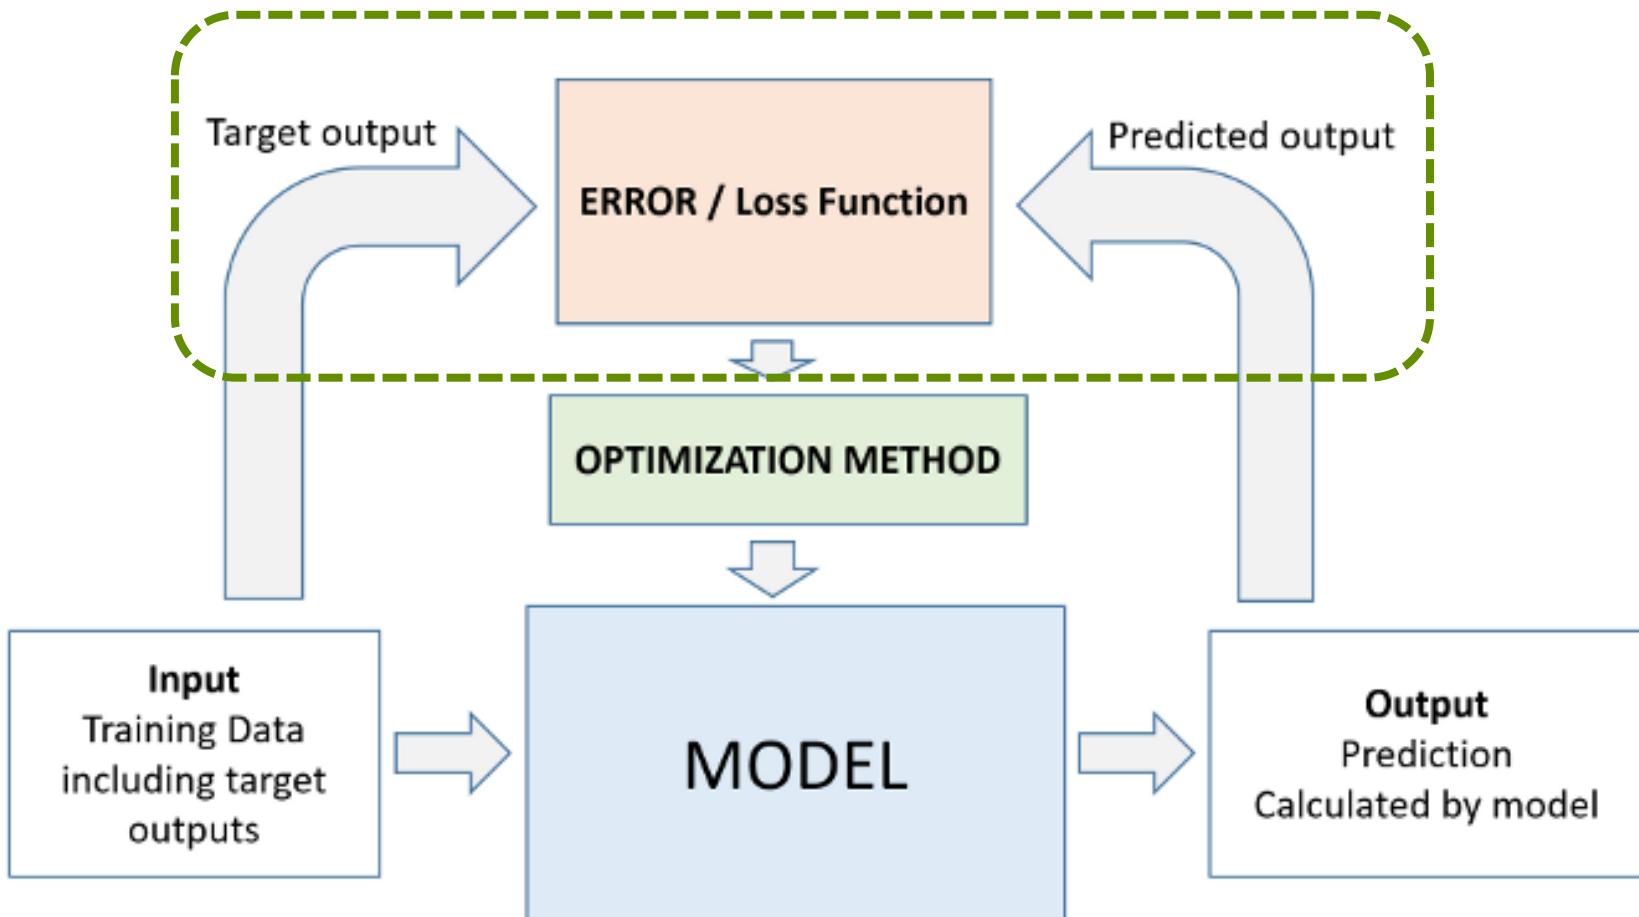


Training Callbacks

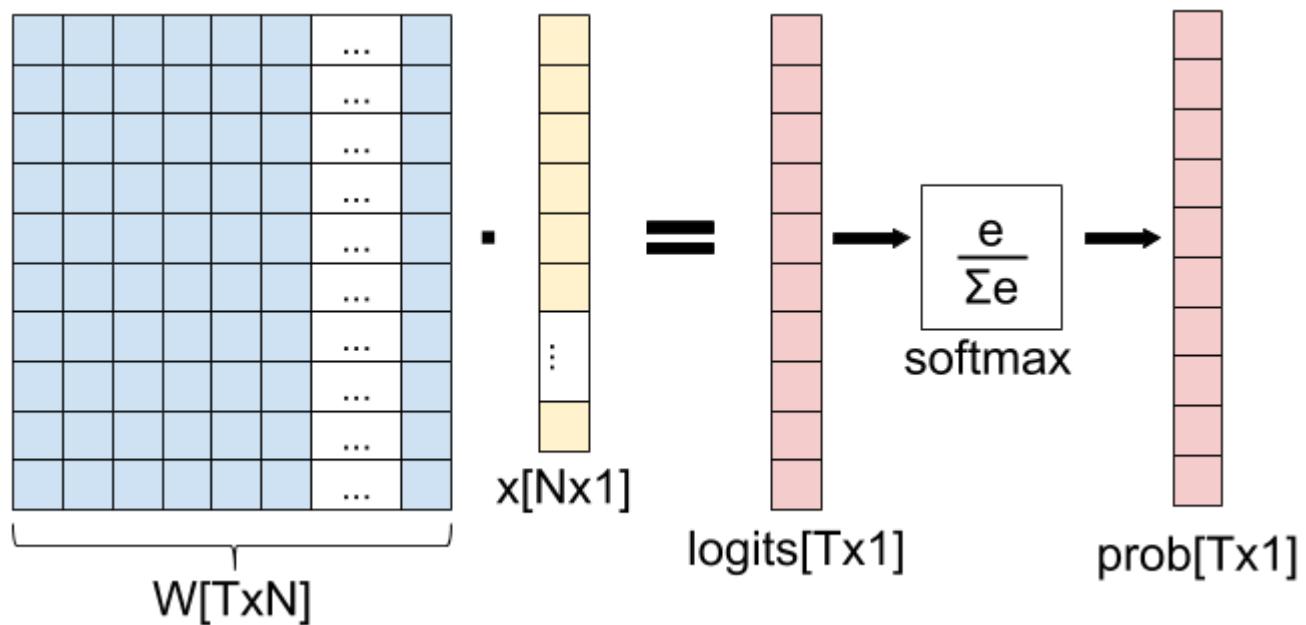
- Keras Callbacks are Predefined Functions automatically executed at the end of every epoch:
 - **Checkpoint Callback:** monitor a metric on validation dataset (IOU for instance) and stores the model weights only if the Performance improved
 - **Reduce on Plateau Callback:** reduce the learning Rate by a factor if the Performance Metric did not improve after a number of epochs
 - **Early stopping callback:** automatically stops training if Performance Metric did not improve after a number of epochs
 - **Tensorboard callback:** stores training and Validation Metrics for live Monitoring

Model Loss Function

Loss Function



The Softmax Layer



The Softmax Layer

Input pixels, x



Feedforward output, y_i

cat	dog	horse
5	4	2
4	2	8
4	4	1

Forward
propagation

Softmax output, $S(y_i)$

cat	dog	horse
0.71	0.26	0.04
0.02	0.00	0.98
0.49	0.49	0.02

Softmax
function

Shape: (3, 32, 32)

Shape: (3,)

Shape: (3,)

Used Loss Functions:

- Multiclass Cross Entropy:

$$L(\hat{y}, y) = - \sum_k^K y^{(k)} \log \hat{y}^{(k)}$$

- Focal Loss:

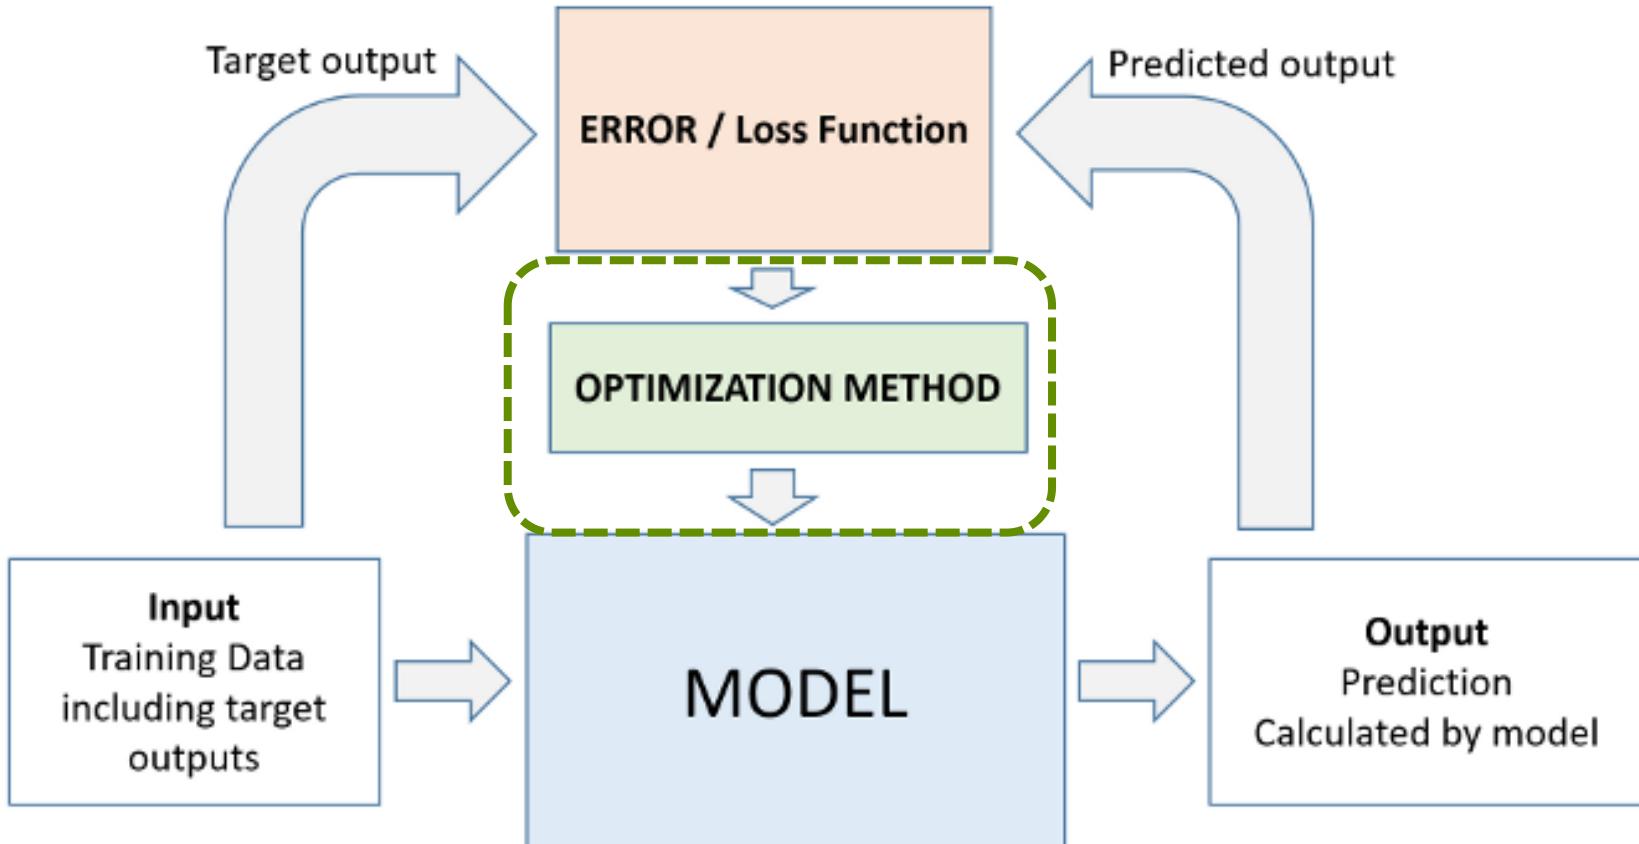
$$\text{FL}(p_t) = -\alpha_t (1 - p_t)^\gamma \log(p_t).$$

- Jaccard Index Loss:

$$L_{IoU} = 1 - IoU = 1 - \frac{I(X)}{U(X)}$$

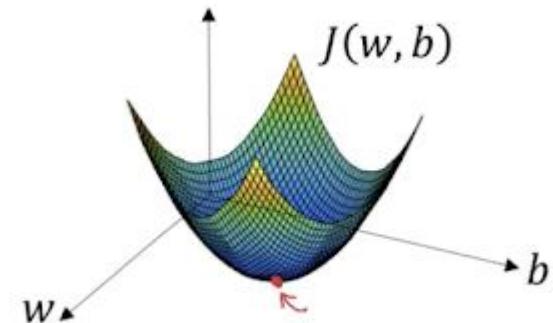
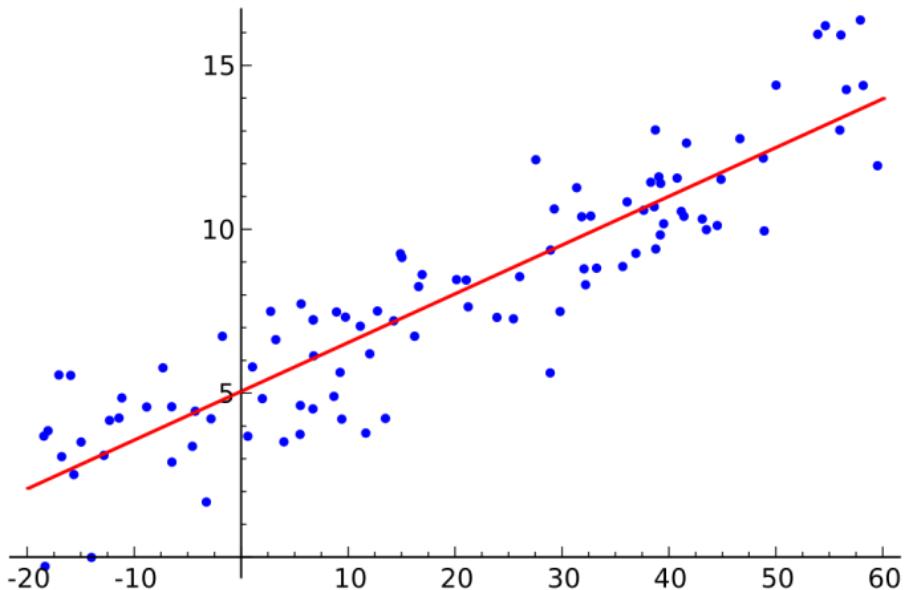
Model Optimization

Model Optimization



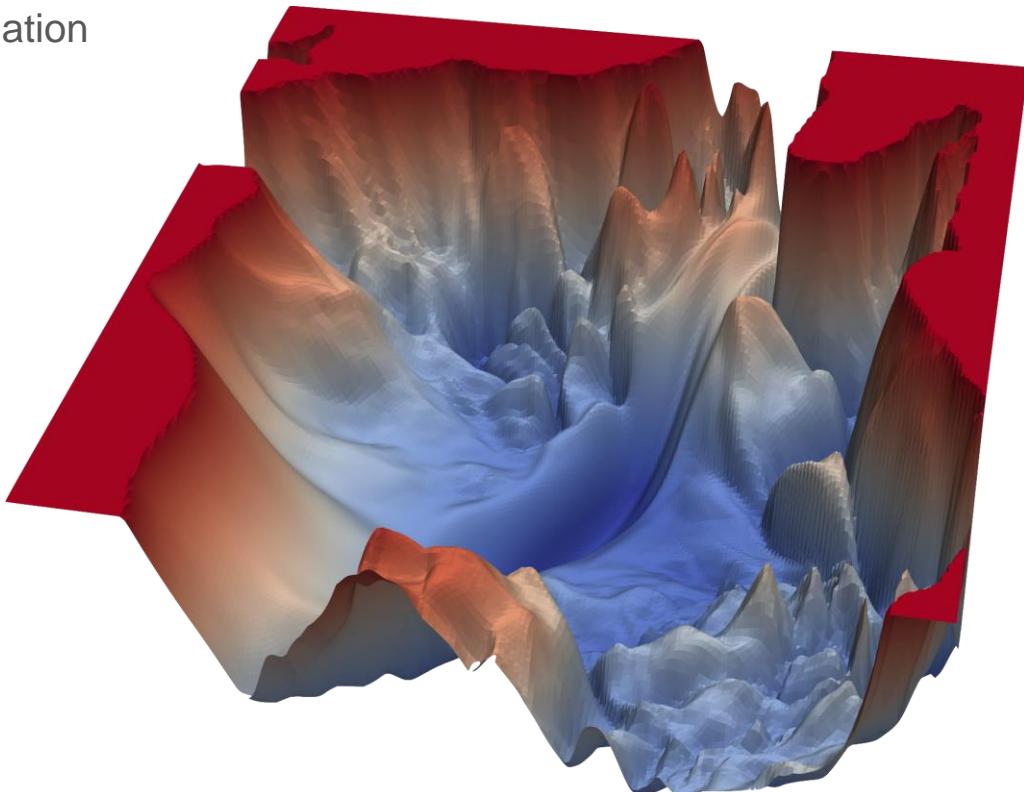
Optimization Algorithms

- The goal is to find the set of weights that minimizes the Loss function
- For Linear Regression:
 - Model: $Y = w X + b$
 - Loss Function (Squared Mean Error): $J(w,b) = \|Y - w X - b\|^2$

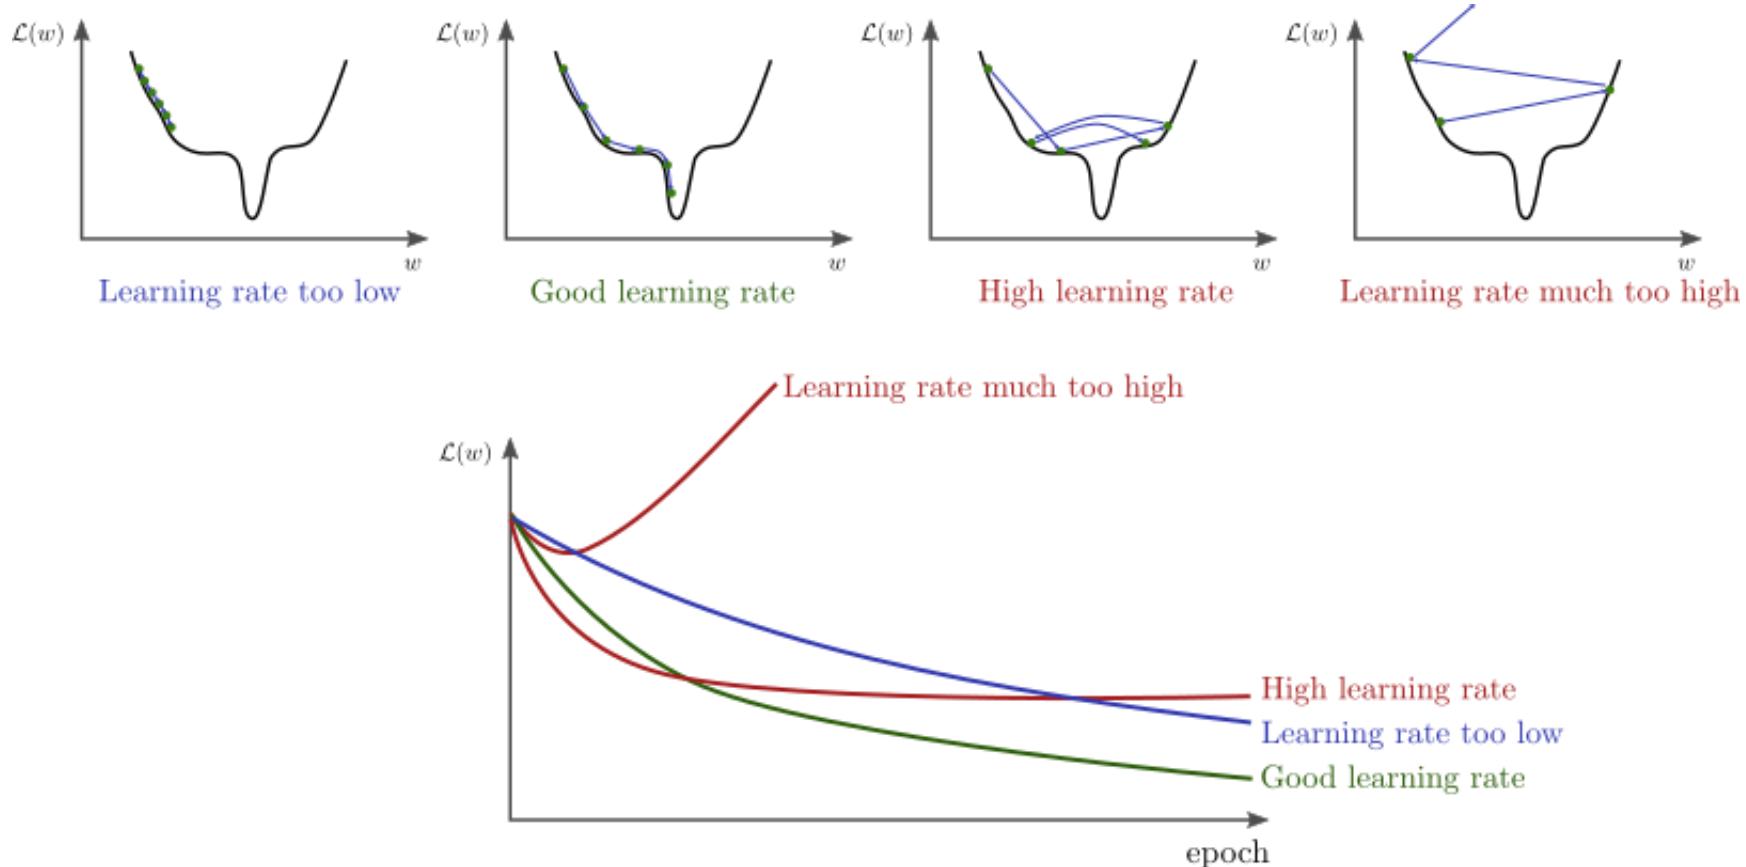


Loss Function Landscape for Neural Network

- All Optimization Algorithm are variation of Gradient Descent and are provided by Tensorflow library
- Gradient of loss function is computed using Automatic-differentiation
- Some of the best known are: SGD, RMSPROP, ADAM
 - In our work SGD with momentum is used
- The common algorithm Parameter: Learning Rate

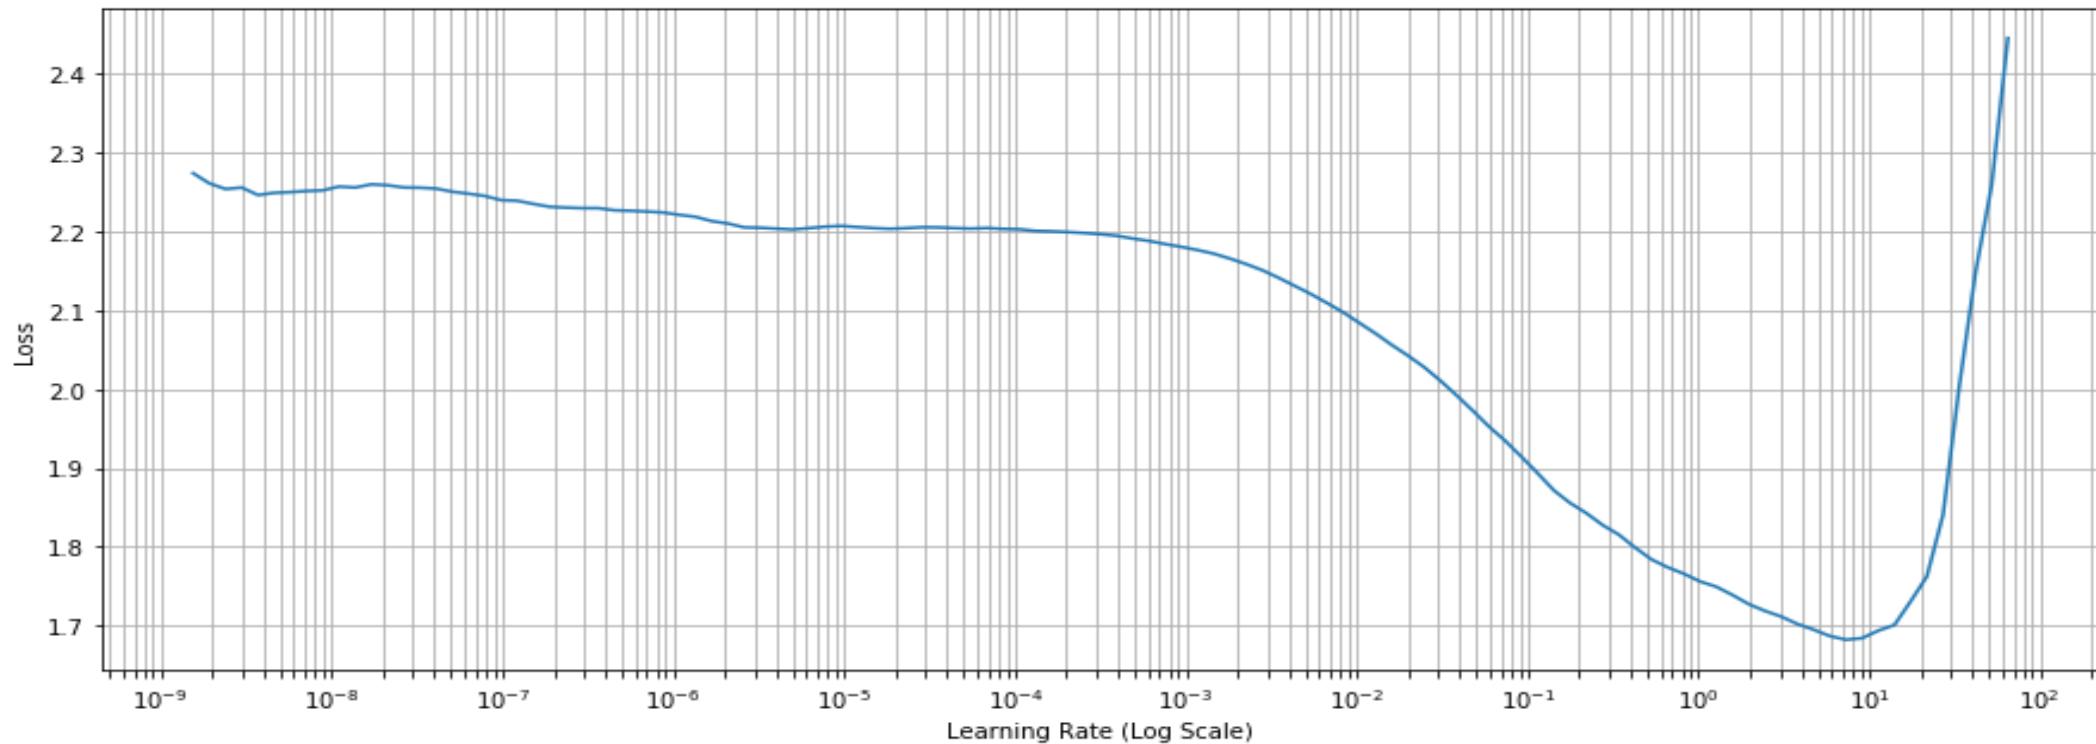


Finding a Good Learning Rate is critical

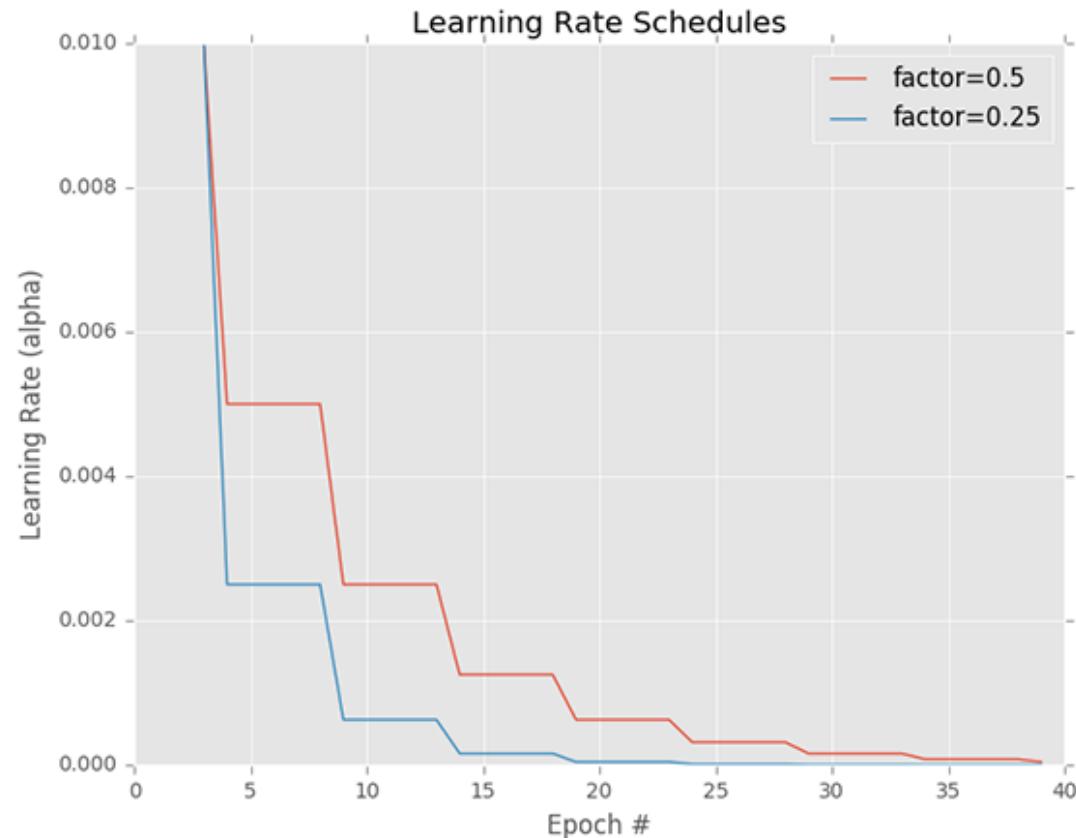


Learning Rate Finder

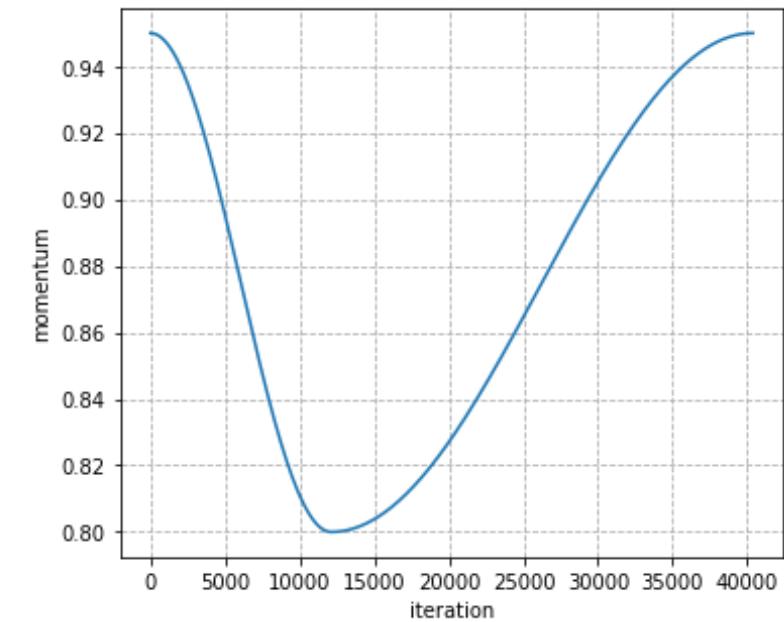
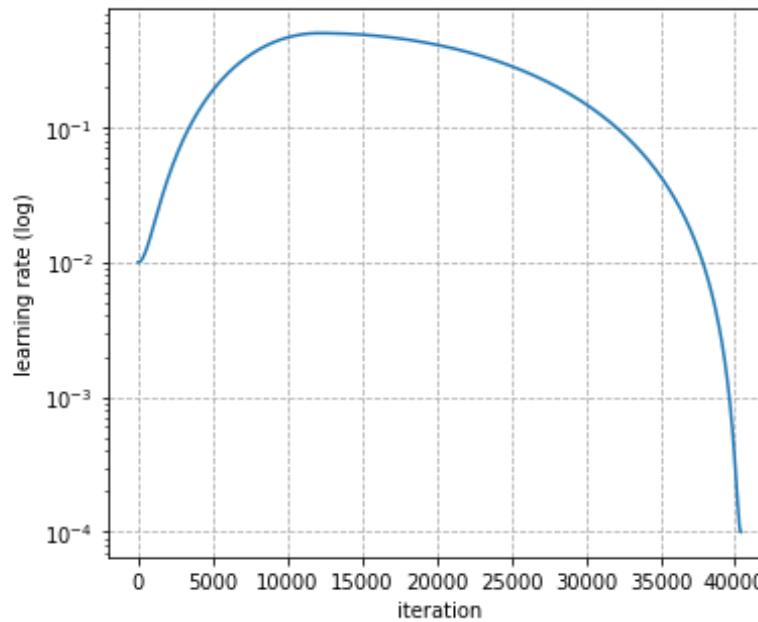
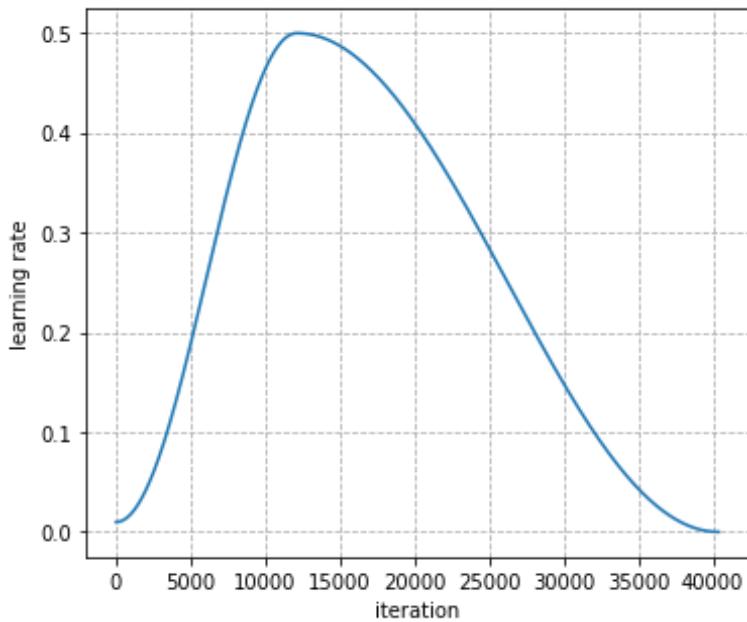
```
[ ]: lrf = LearningRateFinder(model)
lrf.find(train_ds, 1e-9, 1e+2, epochs=1, stepsPerEpoch=TRAIN_STEPS_PER_EPOCH)
lrf.plot_loss()
```



Training Strategy: Reduce on Plateau (Keras Callback)

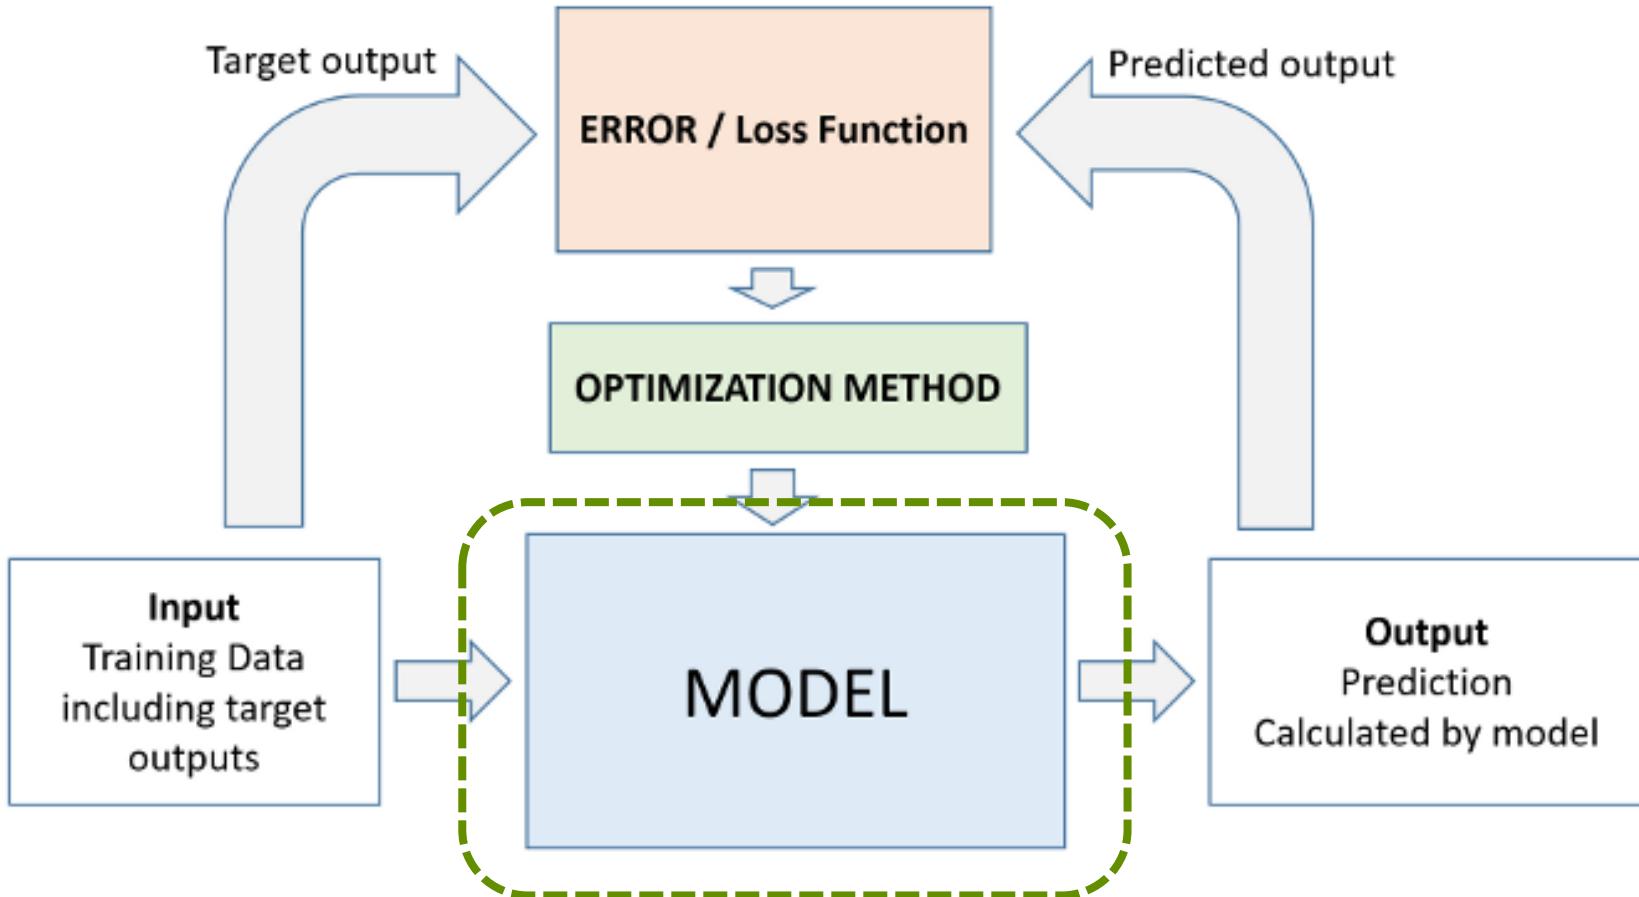


Training strategy: One Cycle Policy (Keras Callback)

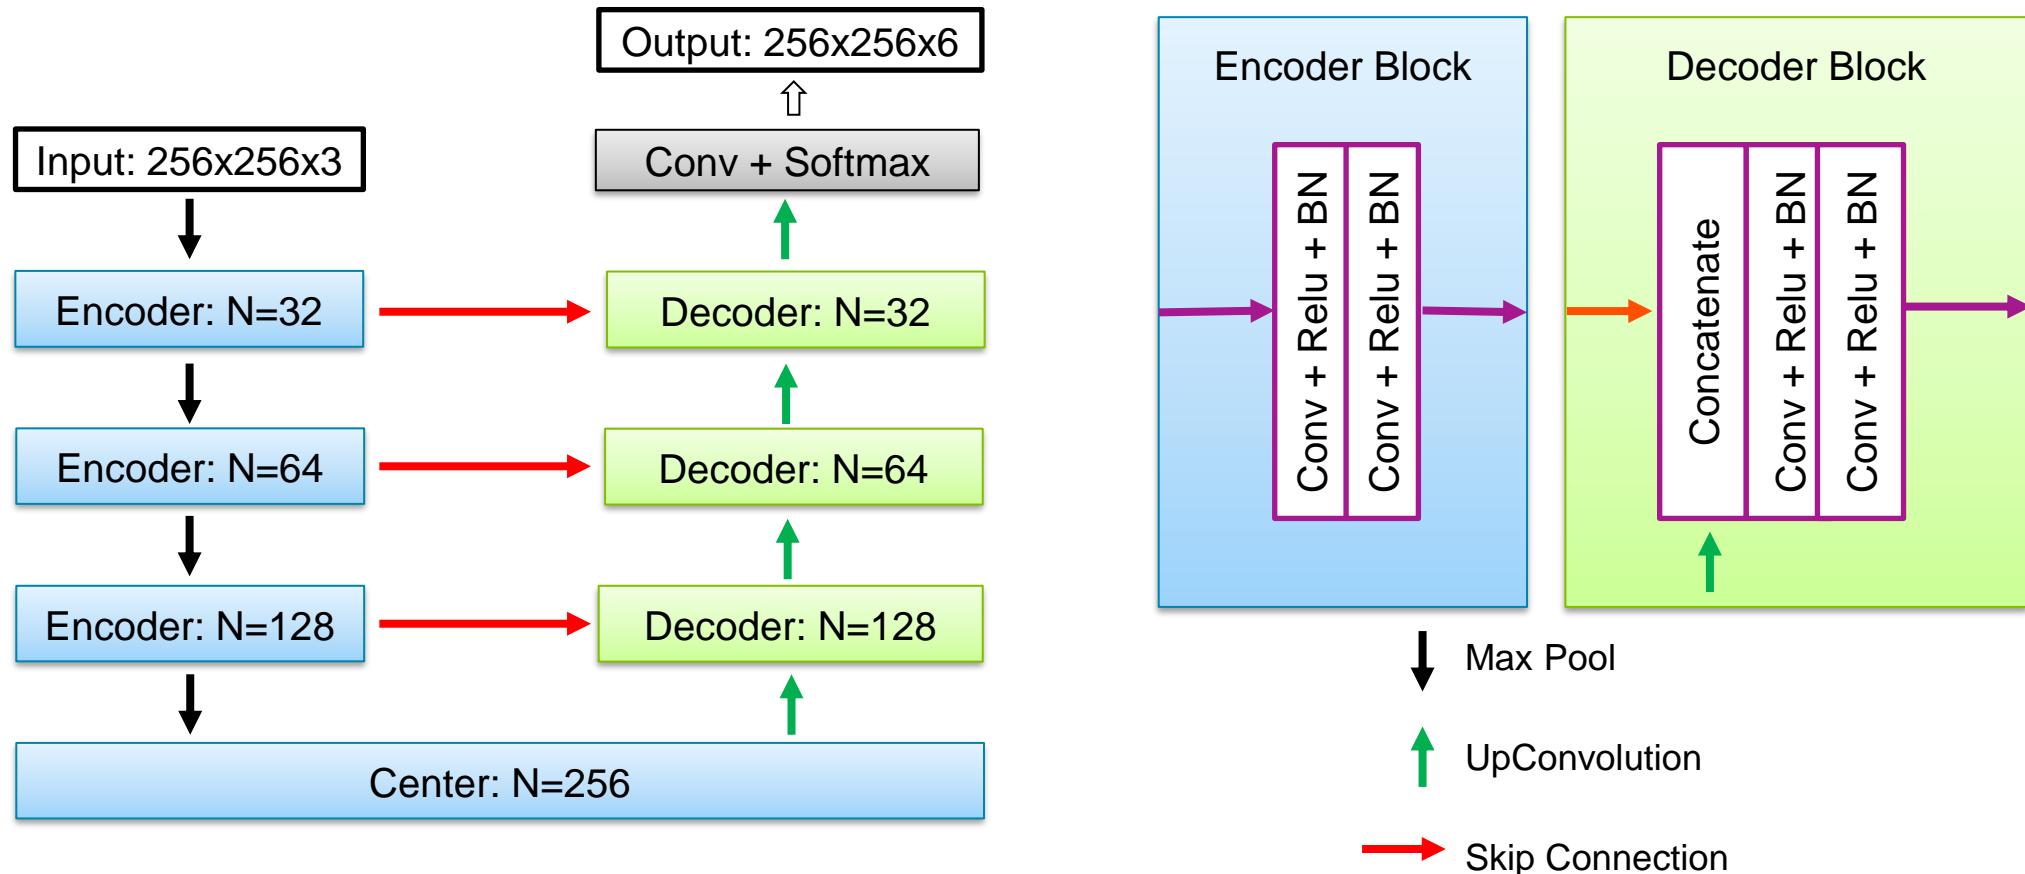


Model Architecture

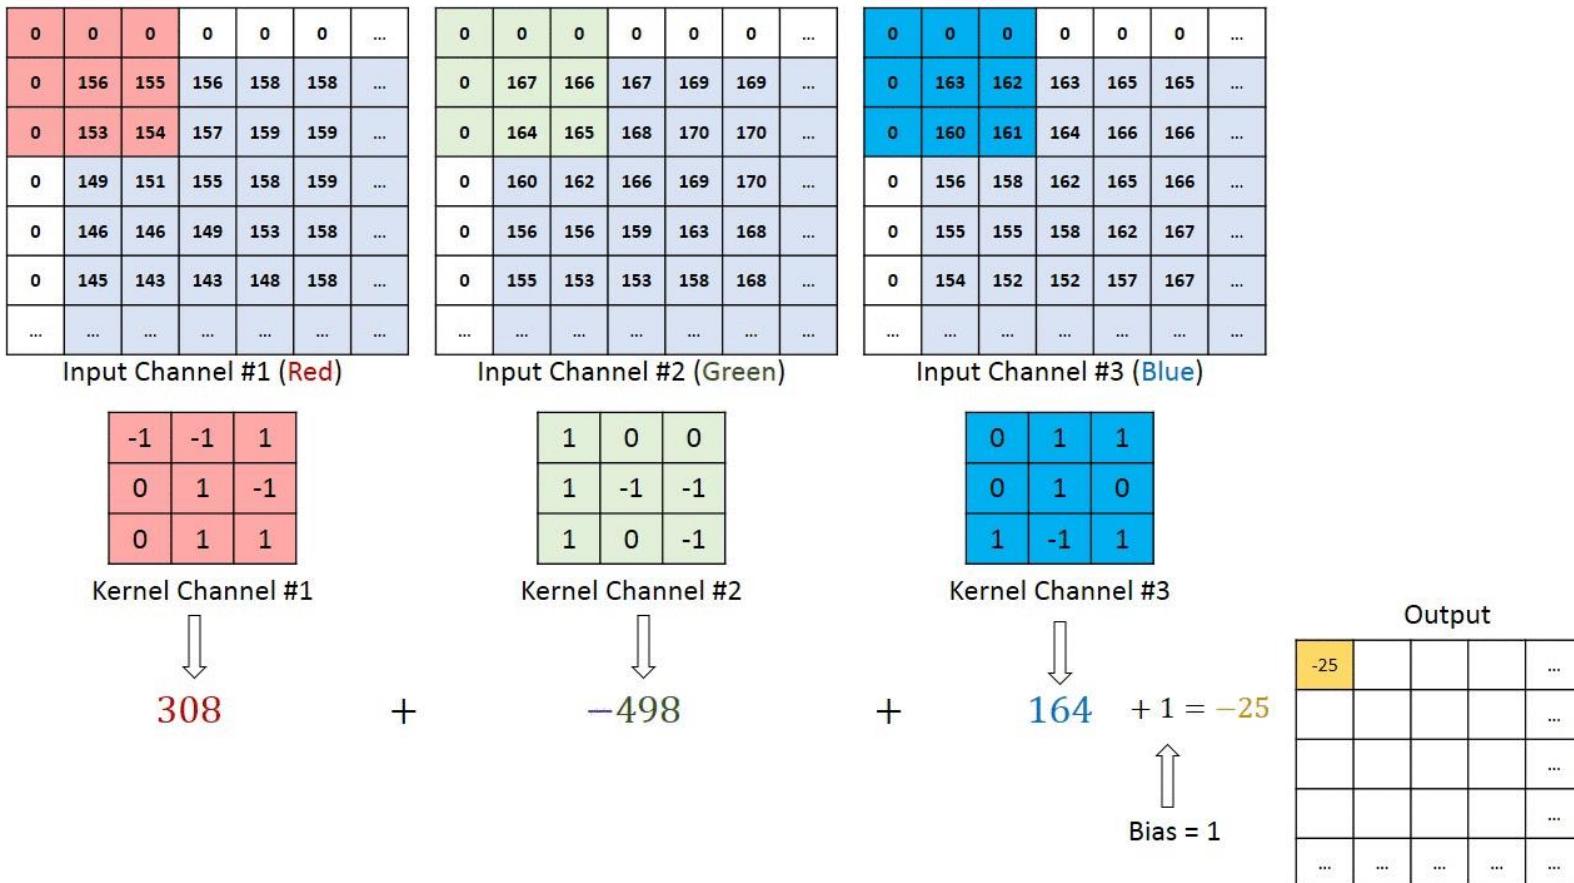
Model Architecture



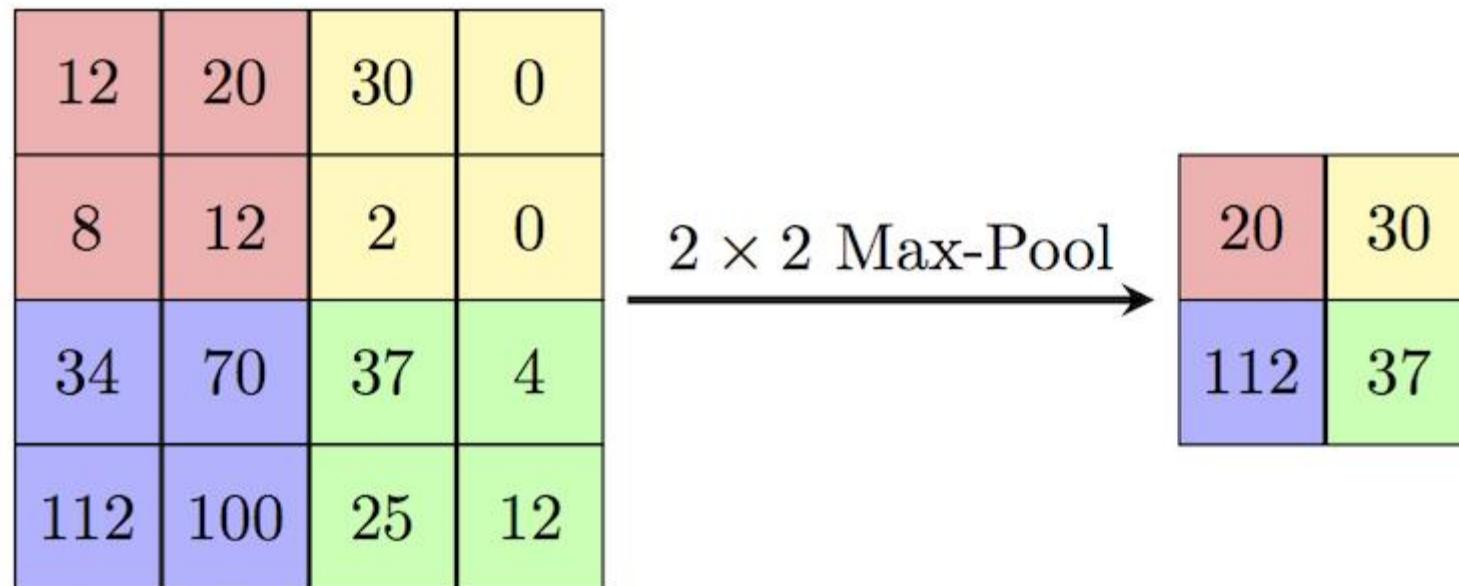
Unet: 3 Skips



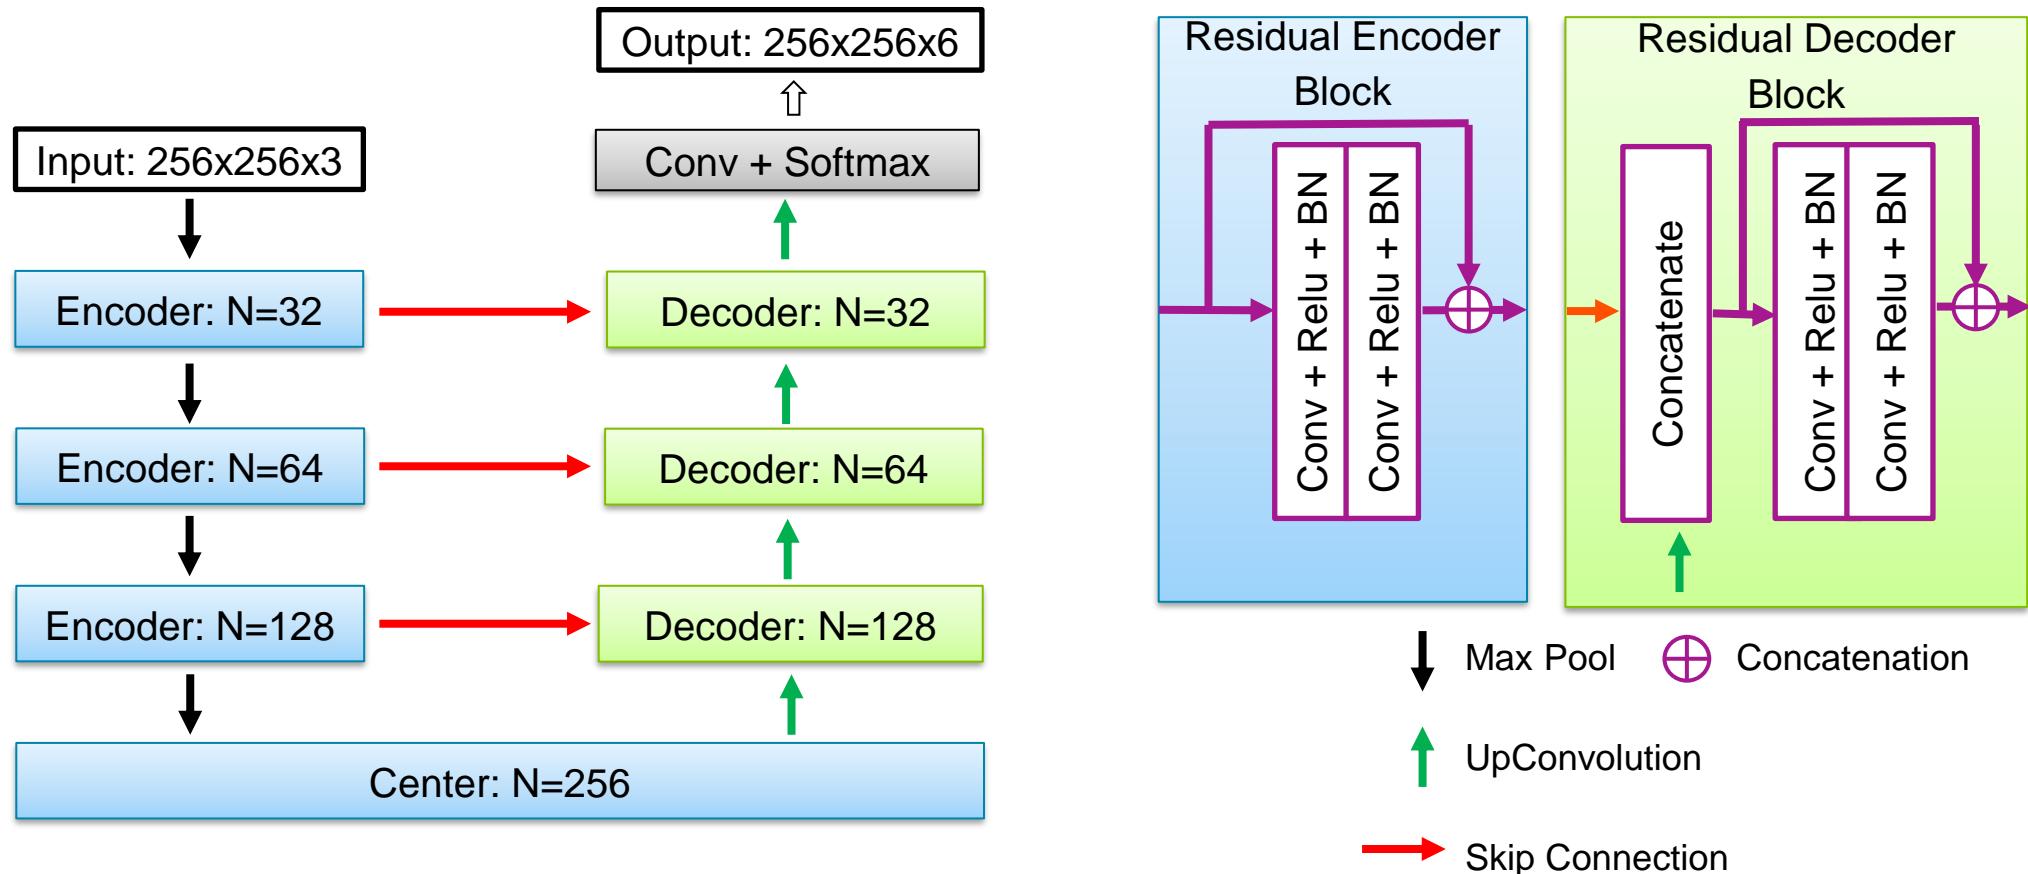
Convolution Layer in Neural Networks



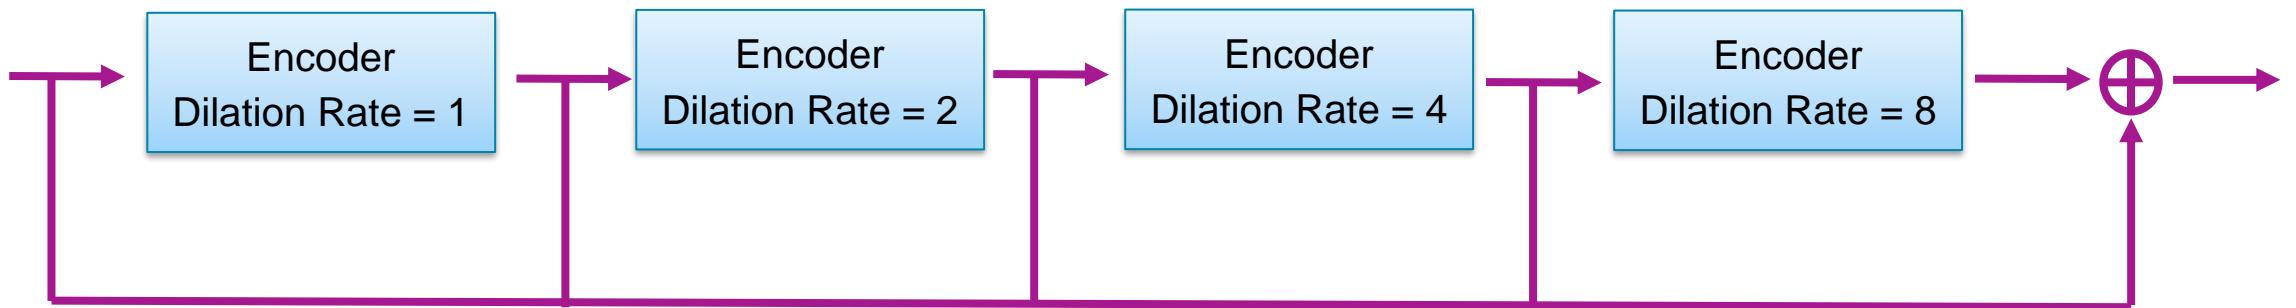
Max Pooling Layer in Neural Networks



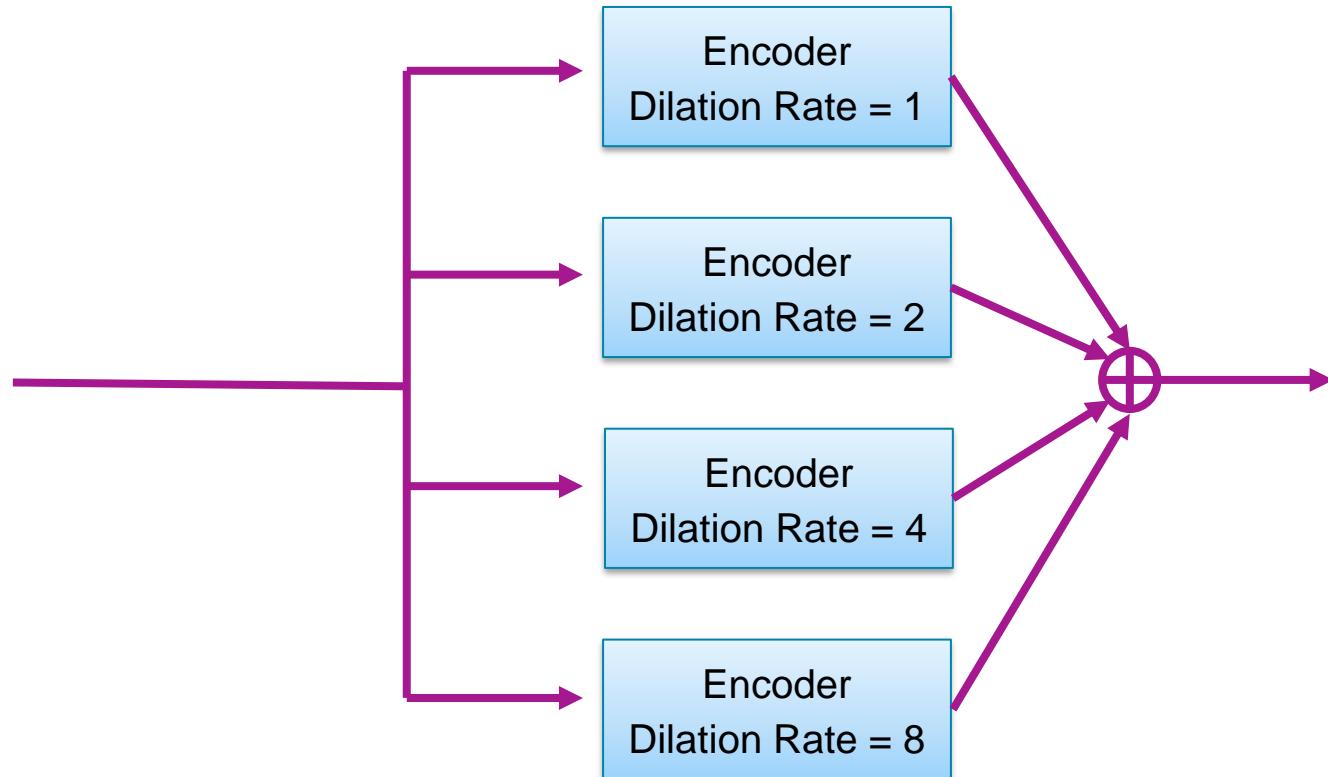
Unet: 3 Skips + Residual Blocks (like ResNet)



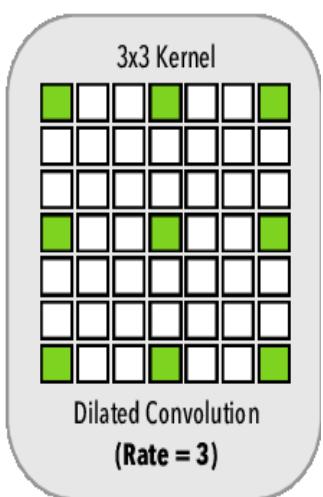
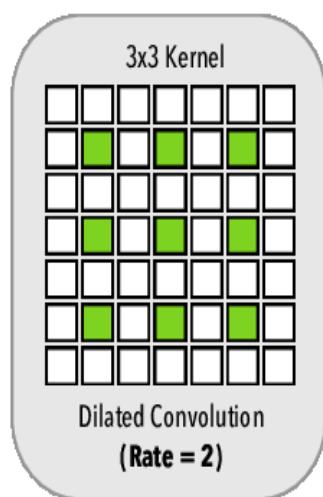
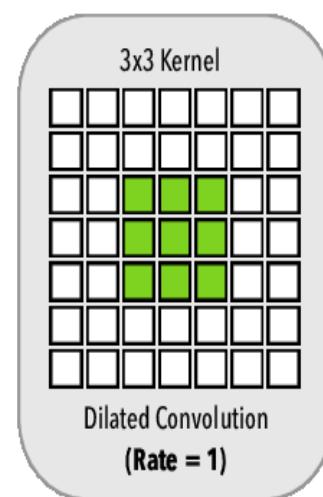
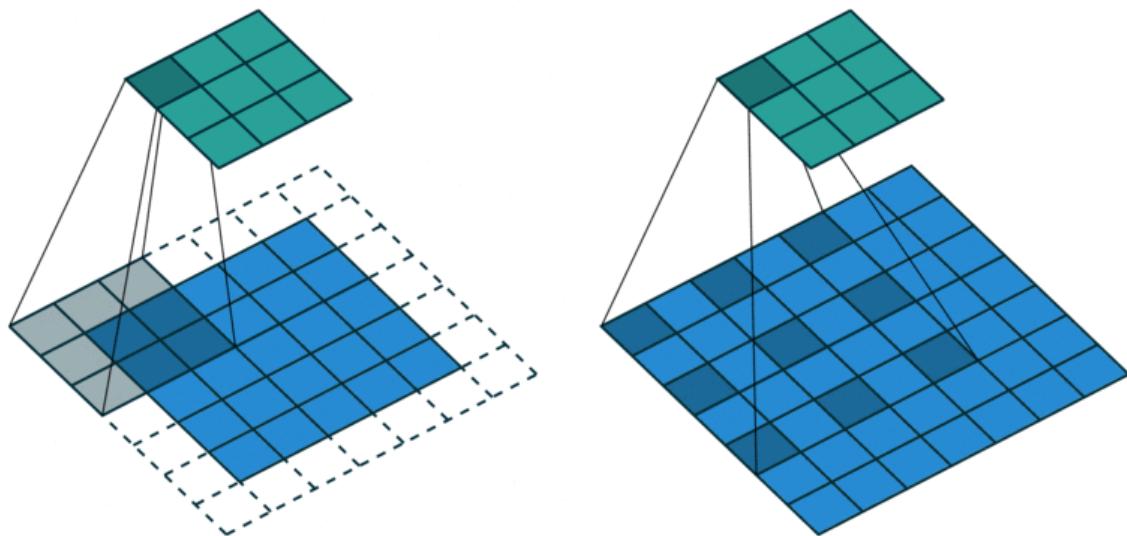
Unet: Serial Dilated Center (like DeepLab)



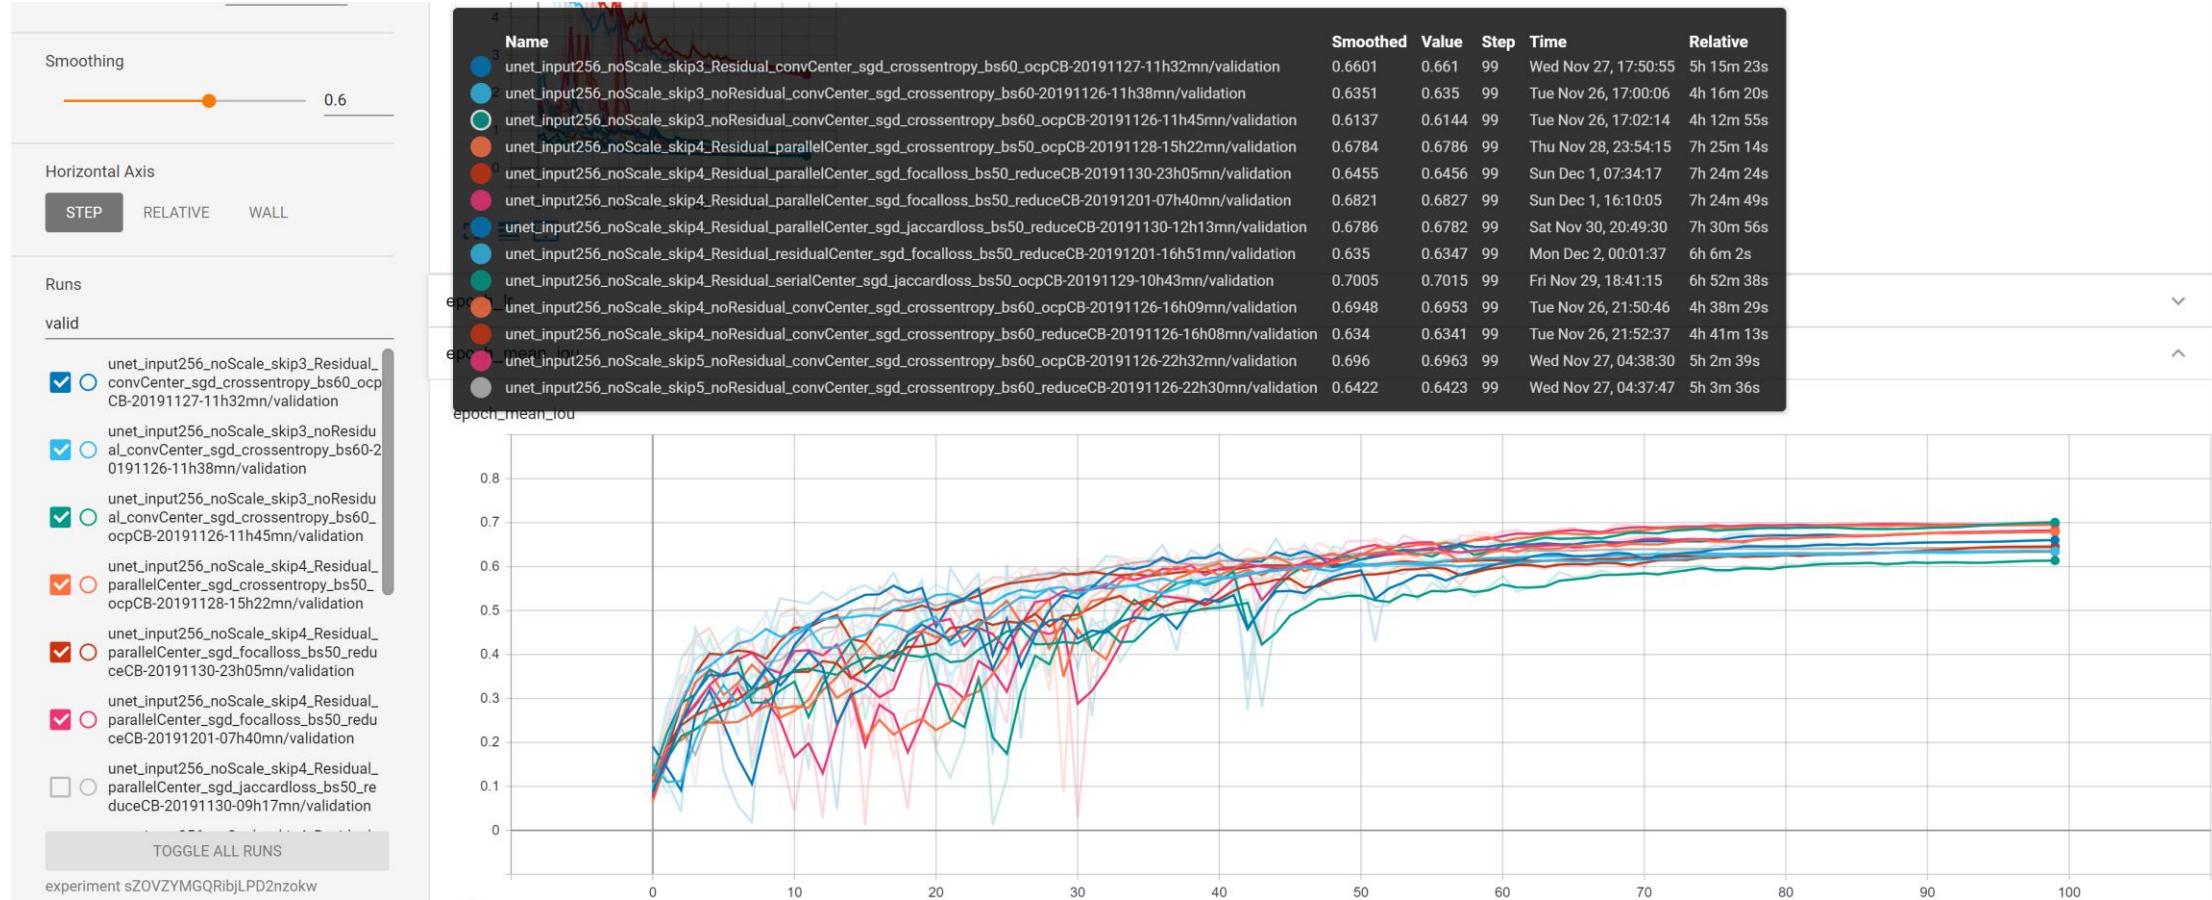
Unet: Parallel Dilated Center (like DeepLab)



Dilated Convolution Layer



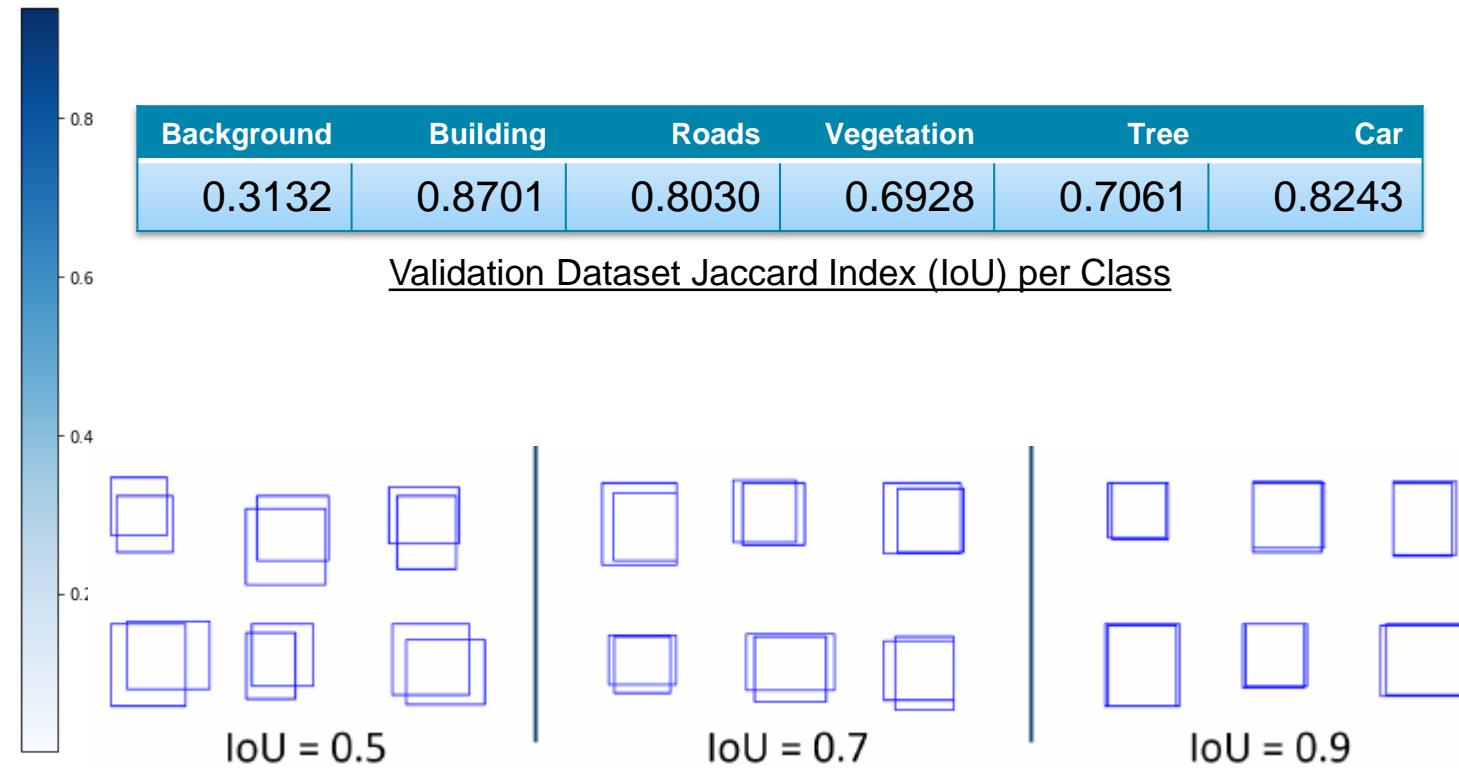
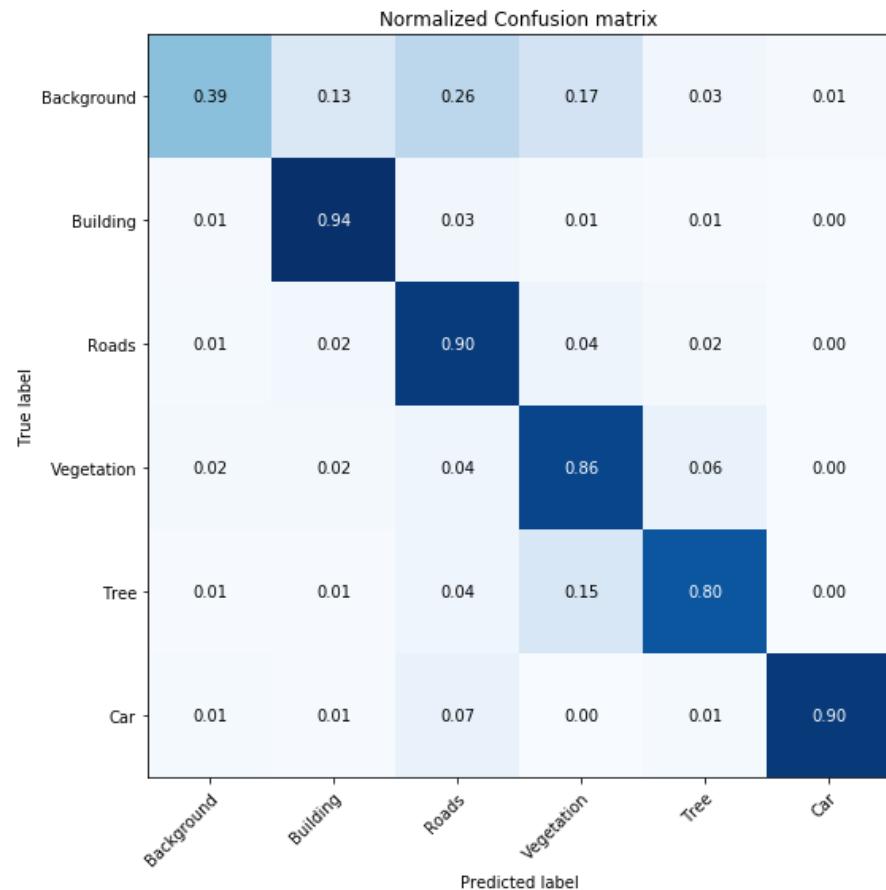
Tensorboard Experiment Results



Experiment Results

Name	Smoothed	Value	Step	Time	Relative
unet_input256_noScale_skip3_noResidual_convCenter_sgd_crossentropy_bs60-20191126-11h38mn/validation	0.6351	0.635	99	Tue Nov 26, 17:00:06	4h 16m 20s
unet_input256_noScale_skip4_Residual_convCenter_sgd_crossentropy_bs60_ocpCB-20191127-11h21mn/validation	0.689	0.689	99	Wed Nov 27, 18:24:11	5h 58m 11s
unet_input256_noScale_skip4_Residual_parallelCenter_sgd_jaccardloss_bs60_ocpCB-20191130-09h17mn/validation	0.6925	0.6929	99	Sat Nov 30, 17:53:51	7h 32m 7s
unet_input256_noScale_skip4_Residual_residualCenter_sgd_crossentropy_bs60_ocpCB-20191128-22h35mn/validation	0.6758	0.6756	99	Fri Nov 29, 04:47:44	5h 8m 49s
unet_input256_noScale_skip4_Residual_residualCenter_sgd_focalloss_bs60_ocpCB-20191201-16h50mn/validation	0.6751	0.6757	99	Mon Dec 2, 00:05:05	6h 10m 22s
unet_input256_noScale_skip4_Residual_residualCenter_sgd_jaccardloss_bs60_ocpCB-20191129-10h43mn/validation	0.6909	0.6918	99	Fri Nov 29, 18:06:55	6h 18m 45s
unet_input256_noScale_skip4_Residual_serialCenter_sgd_crossentropy_bs50_ocpCB-20191128-15h15mn/validation	0.6831	0.6845	99	Thu Nov 28, 23:07:05	6h 47m 12s
unet_input256_noScale_skip4_Residual_serialCenter_sgd_focalloss_bs60_ocpCB-20191130-22h50mn/validation	0.646	0.6467	99	Sun Dec 1, 06:41:48	6h 47m 16s
unet_input256_noScale_skip4_Residual_serialCenter_sgd_focalloss_bs60_ocpCB-20191201-07h40mn/validation	0.6879	0.6885	99	Sun Dec 1, 15:31:12	6h 46m 57s
epoch_mean_iou					
unet_input256_noScale_skip4_noResidual_convCenter_sgd_crossentropy_bs60_reduceCB-20191126-16h08mn/validation	0.634	0.6341	99	Tue Nov 26, 21:52:37	4h 41m 13s
unet_input256_noScale_skip5_noResidual_convCenter_sgd_crossentropy_bs60_reduceCB-20191126-22h30mn/validation	0.6422	0.6423	99	Wed Nov 27, 04:37:47	5h 3m 36s
epoch_lr					
Name	Smoothed	Value	Step	Time	Relative
unet_input256_noScale_skip3_Residual_convCenter_sgd_crossentropy_bs60_ocpCB-20191127-11h32mn/validation	0.6601	0.661	99	Wed Nov 27, 17:50:55	5h 15m 23s
unet_input256_noScale_skip3_noResidual_convCenter_sgd_crossentropy_bs60-20191126-11h38mn/validation	0.6351	0.635	99	Tue Nov 26, 17:00:06	4h 16m 20s
unet_input256_noScale_skip3_noResidual_convCenter_sgd_crossentropy_bs60_ocpCB-20191126-11h45mn/validation	0.6137	0.6144	99	Tue Nov 26, 17:02:14	4h 12m 55s
unet_input256_noScale_skip4_Residual_parallelCenter_sgd_crossentropy_bs50_ocpCB-20191128-15h22mn/validation	0.6784	0.6786	99	Thu Nov 28, 23:54:15	7h 25m 14s
unet_input256_noScale_skip4_Residual_parallelCenter_sgd_focalloss_bs50_reduceCB-20191130-23h05mn/validation	0.6455	0.6456	99	Sun Dec 1, 07:34:17	7h 24m 24s
unet_input256_noScale_skip4_Residual_parallelCenter_sgd_focalloss_bs50_reduceCB-20191201-07h40mn/validation	0.6821	0.6827	99	Sun Dec 1, 16:10:05	7h 24m 49s
unet_input256_noScale_skip4_Residual_parallelCenter_sgd_jaccardloss_bs50_reduceCB-20191130-12h13mn/validation	0.6786	0.6782	99	Sat Nov 30, 20:49:30	7h 30m 56s
unet_input256_noScale_skip4_Residual_residualCenter_sgd_focalloss_bs50_reduceCB-20191201-16h51mn/validation	0.635	0.6347	99	Mon Dec 2, 00:01:37	6h 6m 2s
unet_input256_noScale_skip4_Residual_serialCenter_sgd_jaccardloss_bs50_ocpCB-20191129-10h43mn/validation	0.7005	0.7015	99	Fri Nov 29, 18:41:15	6h 52m 38s
unet_input256_noScale_skip4_noResidual_convCenter_sgd_crossentropy_bs60_ocpCB-20191126-16h09mn/validation	0.6948	0.6953	99	Tue Nov 26, 21:50:46	4h 38m 29s
unet_input256_noScale_skip4_noResidual_convCenter_sgd_crossentropy_bs60_reduceCB-20191126-16h08mn/validation	0.634	0.6341	99	Tue Nov 26, 21:52:37	4h 41m 13s
epoch_mean_iou					
unet_input256_noScale_skip5_noResidual_convCenter_sgd_crossentropy_bs60_ocpCB-20191126-22h32mn/validation	0.696	0.6963	99	Wed Nov 27, 04:38:30	5h 2m 39s
unet_input256_noScale_skip5_noResidual_convCenter_sgd_crossentropy_bs60_reduceCB-20191126-22h30mn/validation	0.6422	0.6423	99	Wed Nov 27, 04:37:47	5h 3m 36s

Best Model Detailed Results Computed on Validation Set

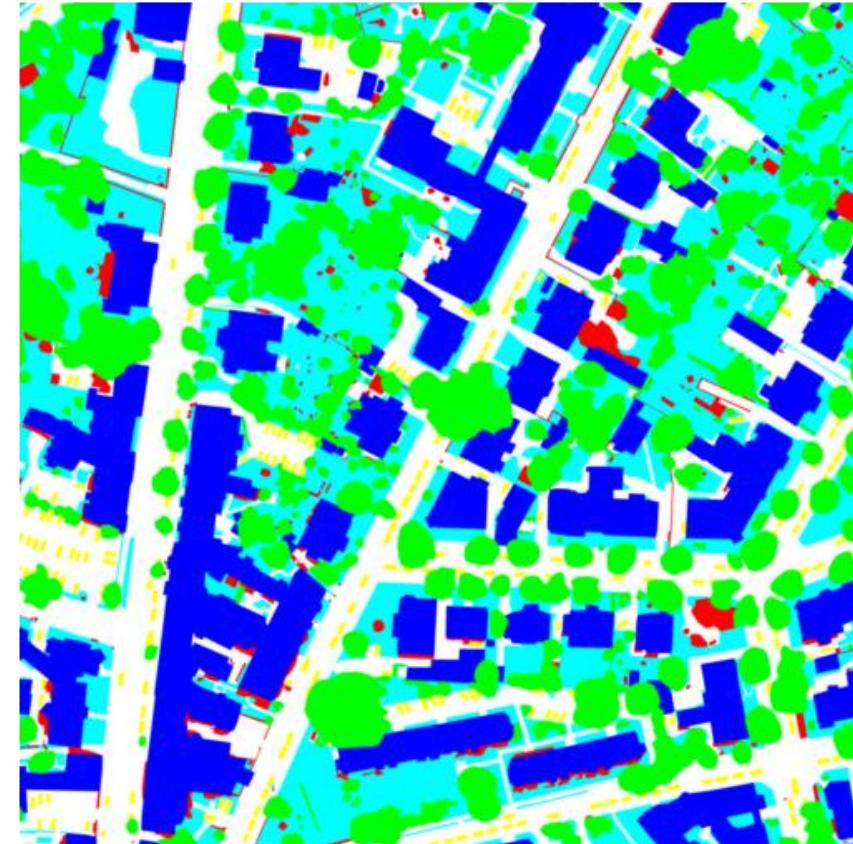


Predicting on a new Tile (from Validation set)

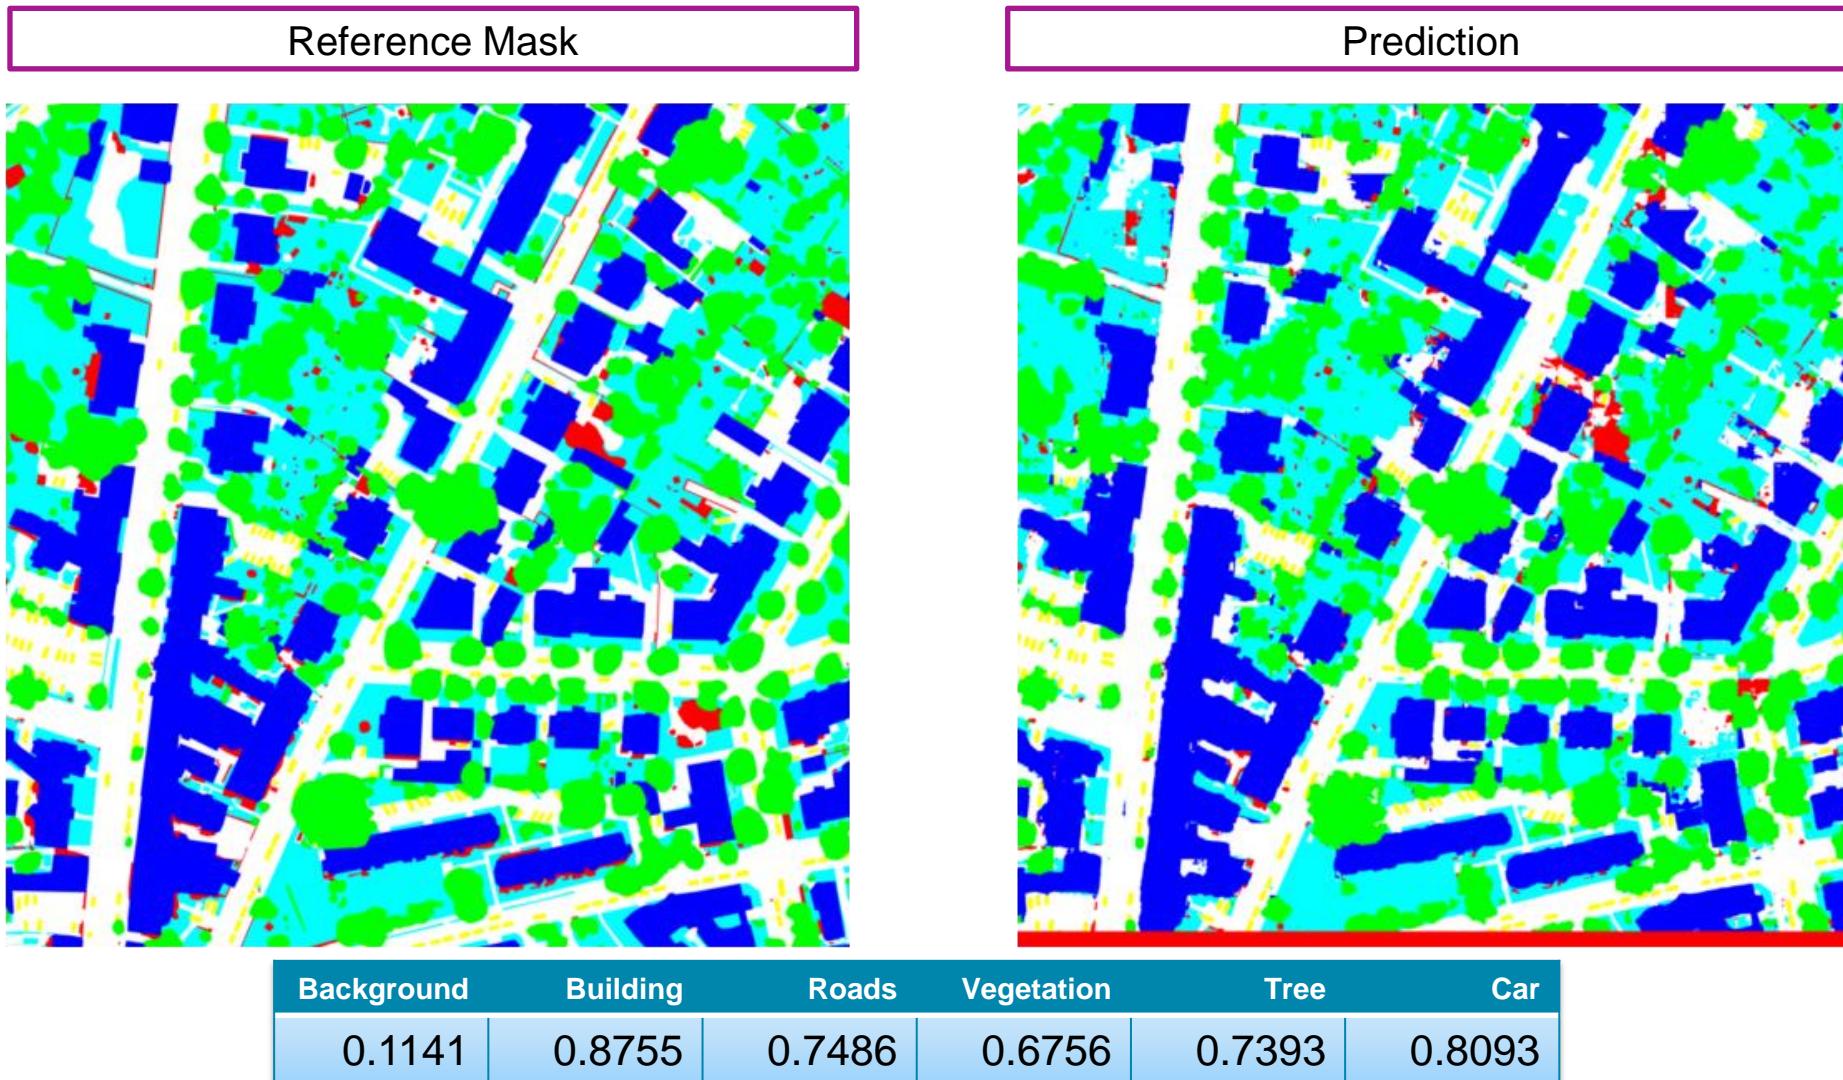
Aerial Tile



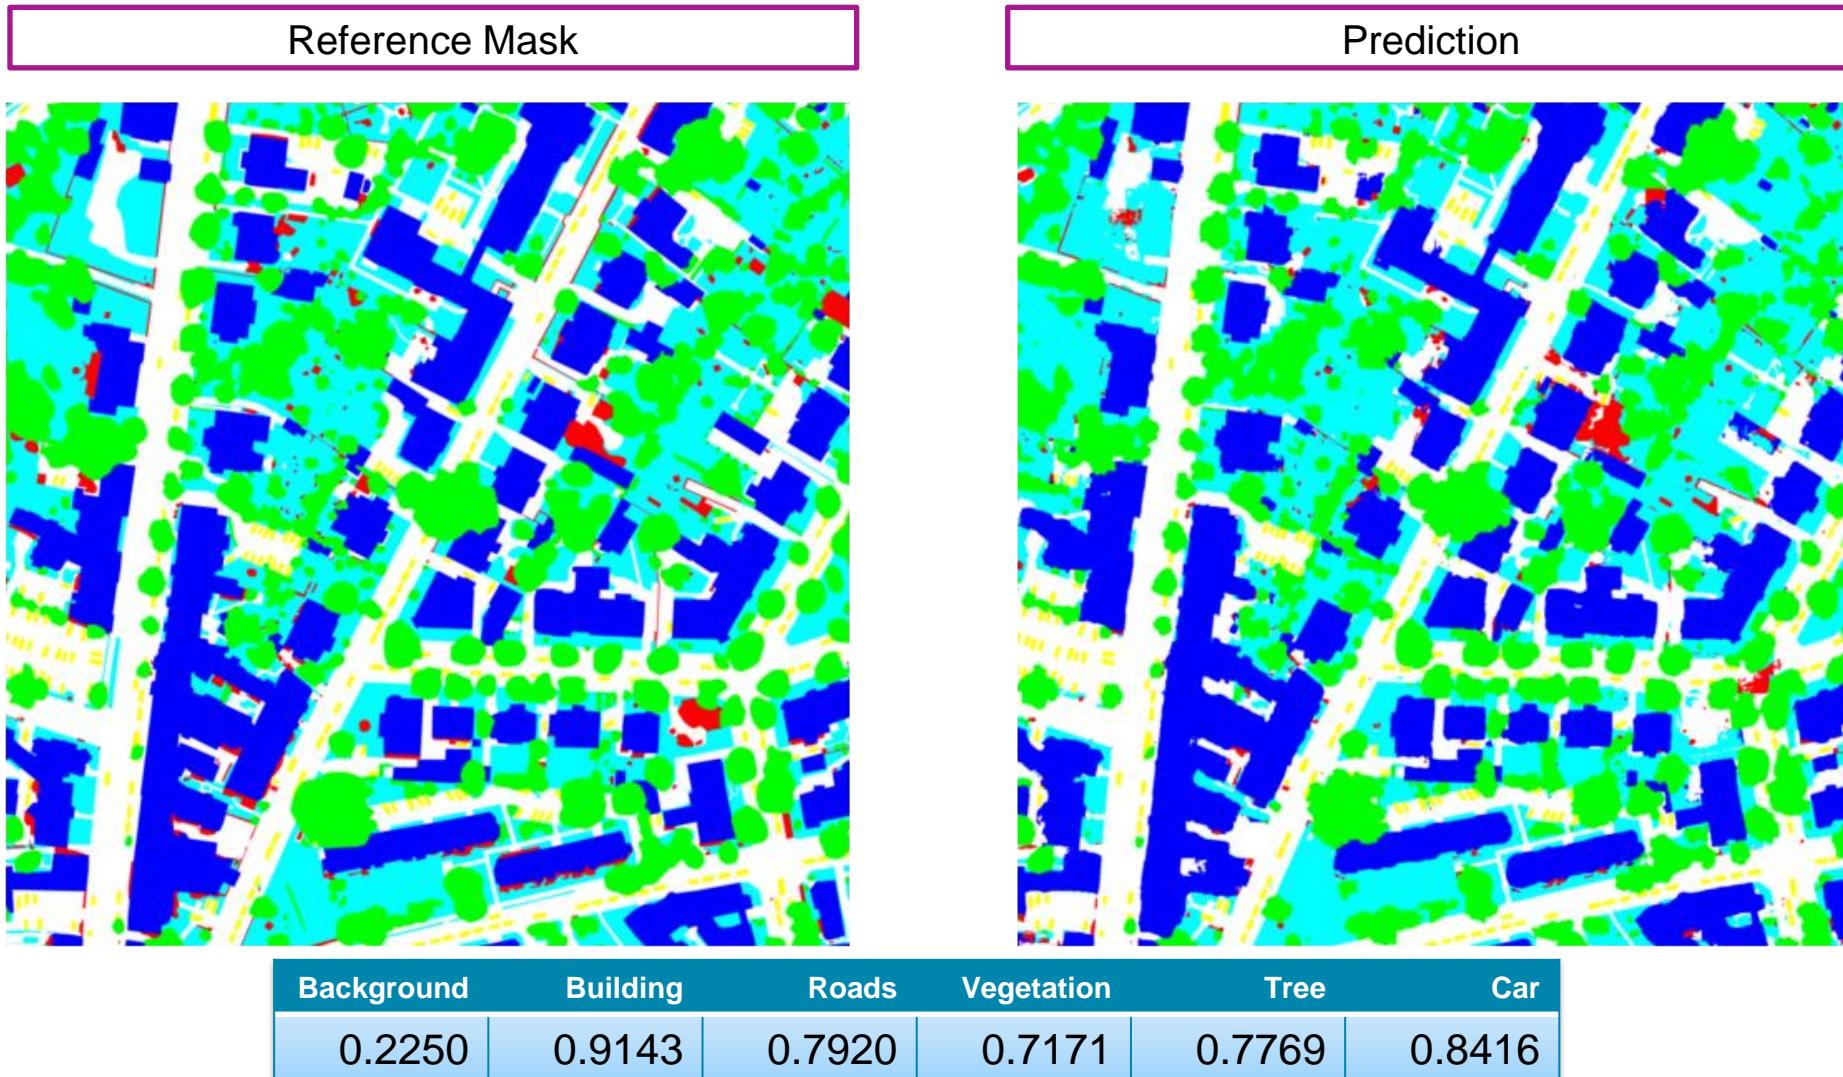
Reference Mask



Predicting on a new Tile (from Validation set)



Predicting on a new Tile with Test Time Augmentation

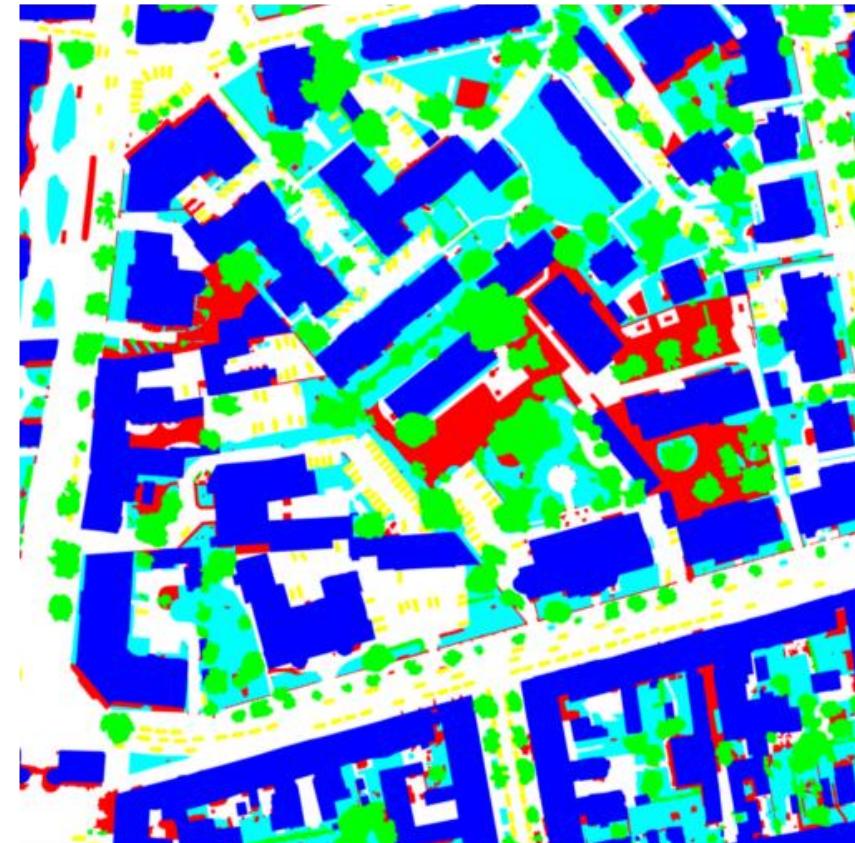


Predicting on a new Tile (Example 2)

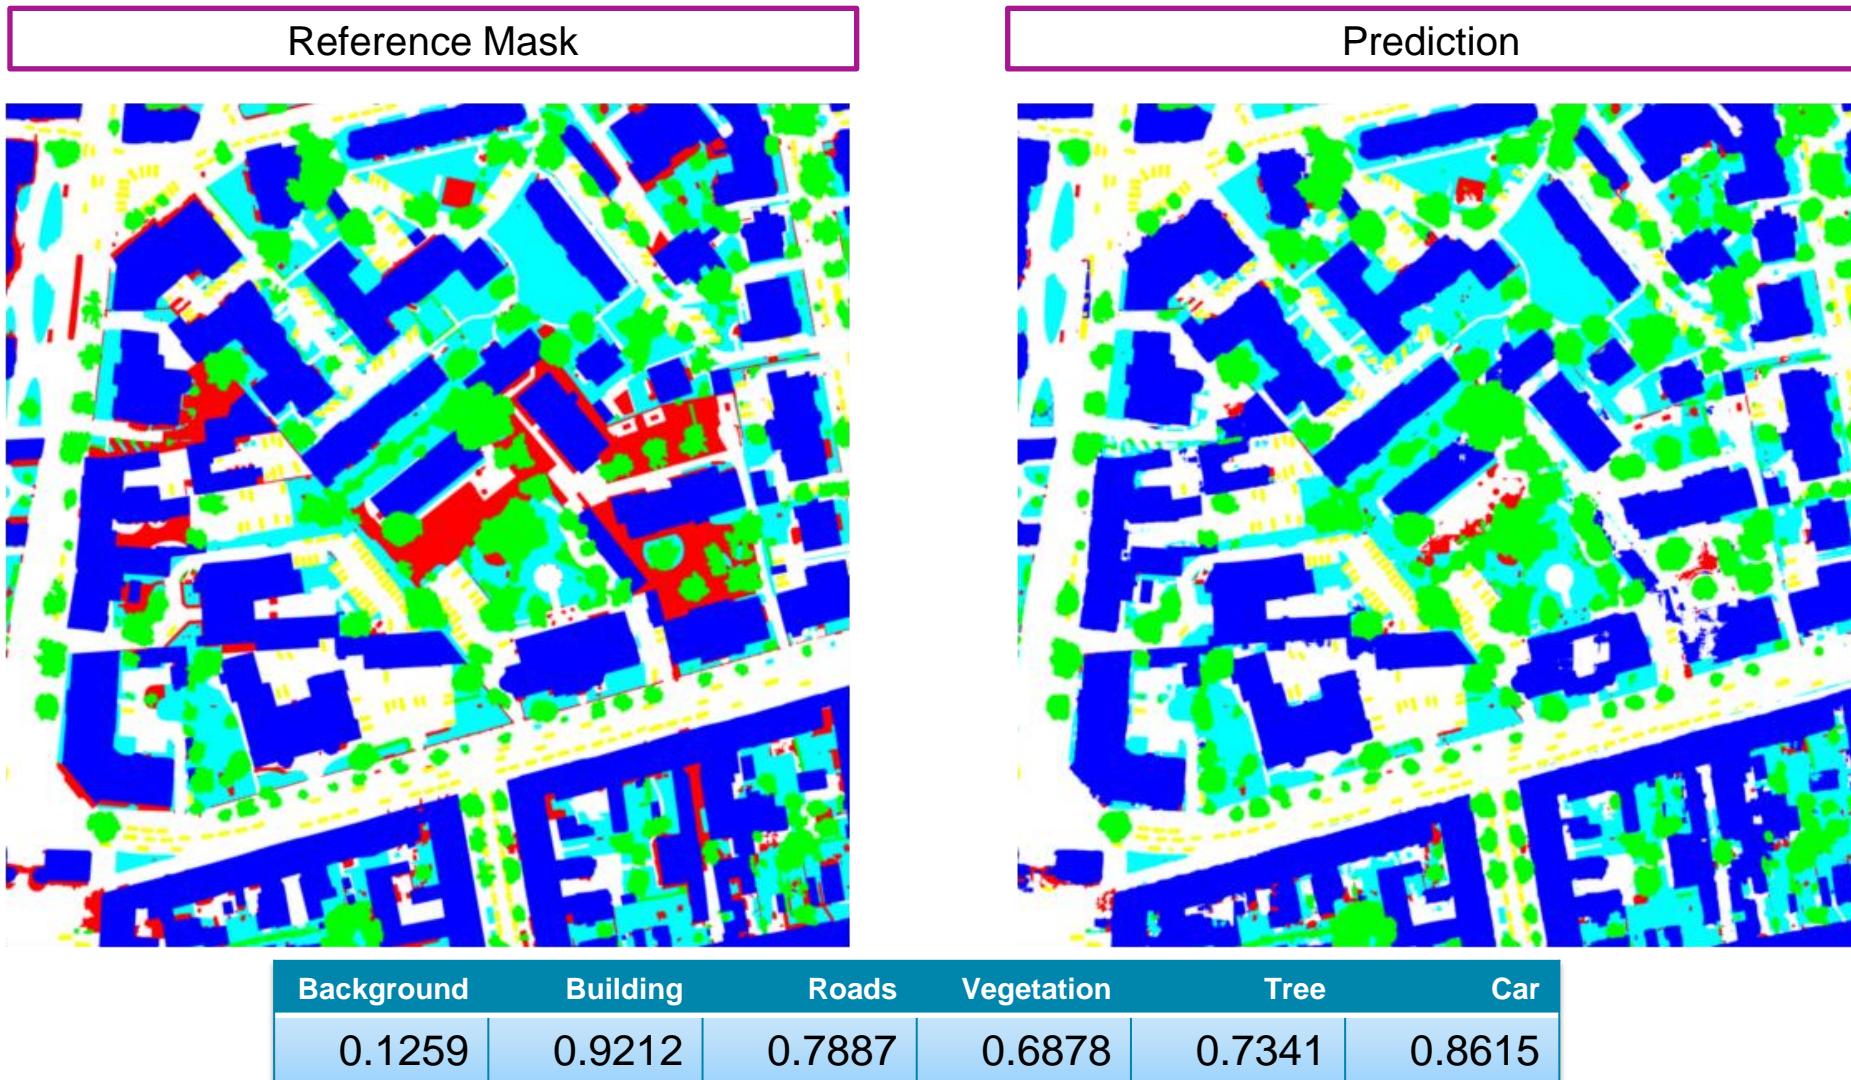
Aerial Tile



Reference Mask

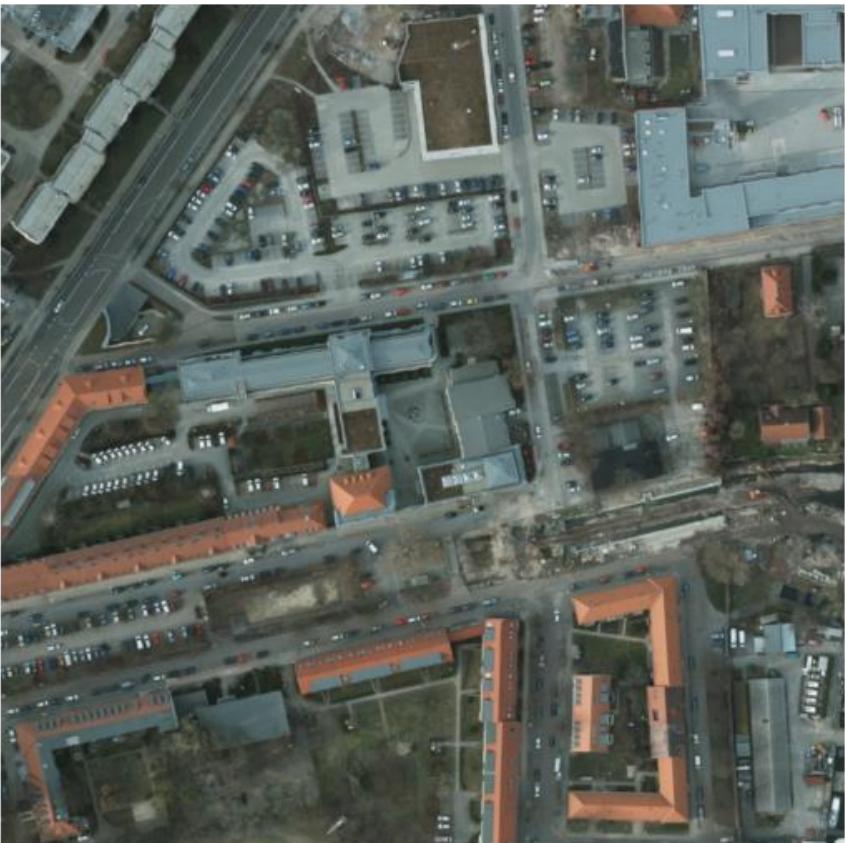


Predicting on a new Tile (Example 2)

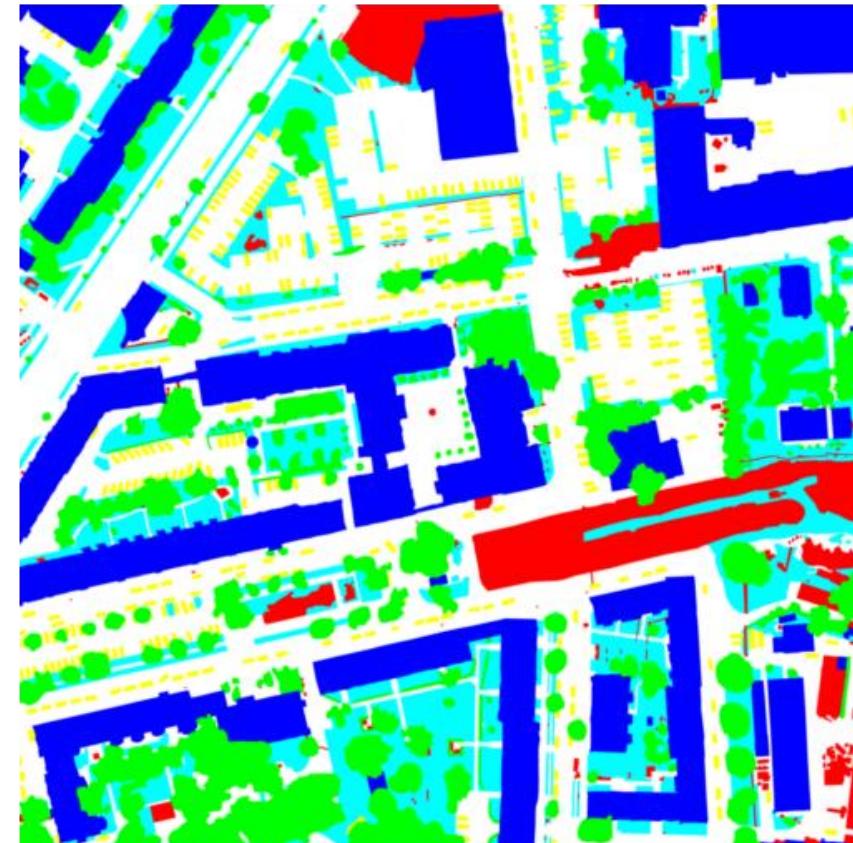


Predicting on a new Tile (Example 3)

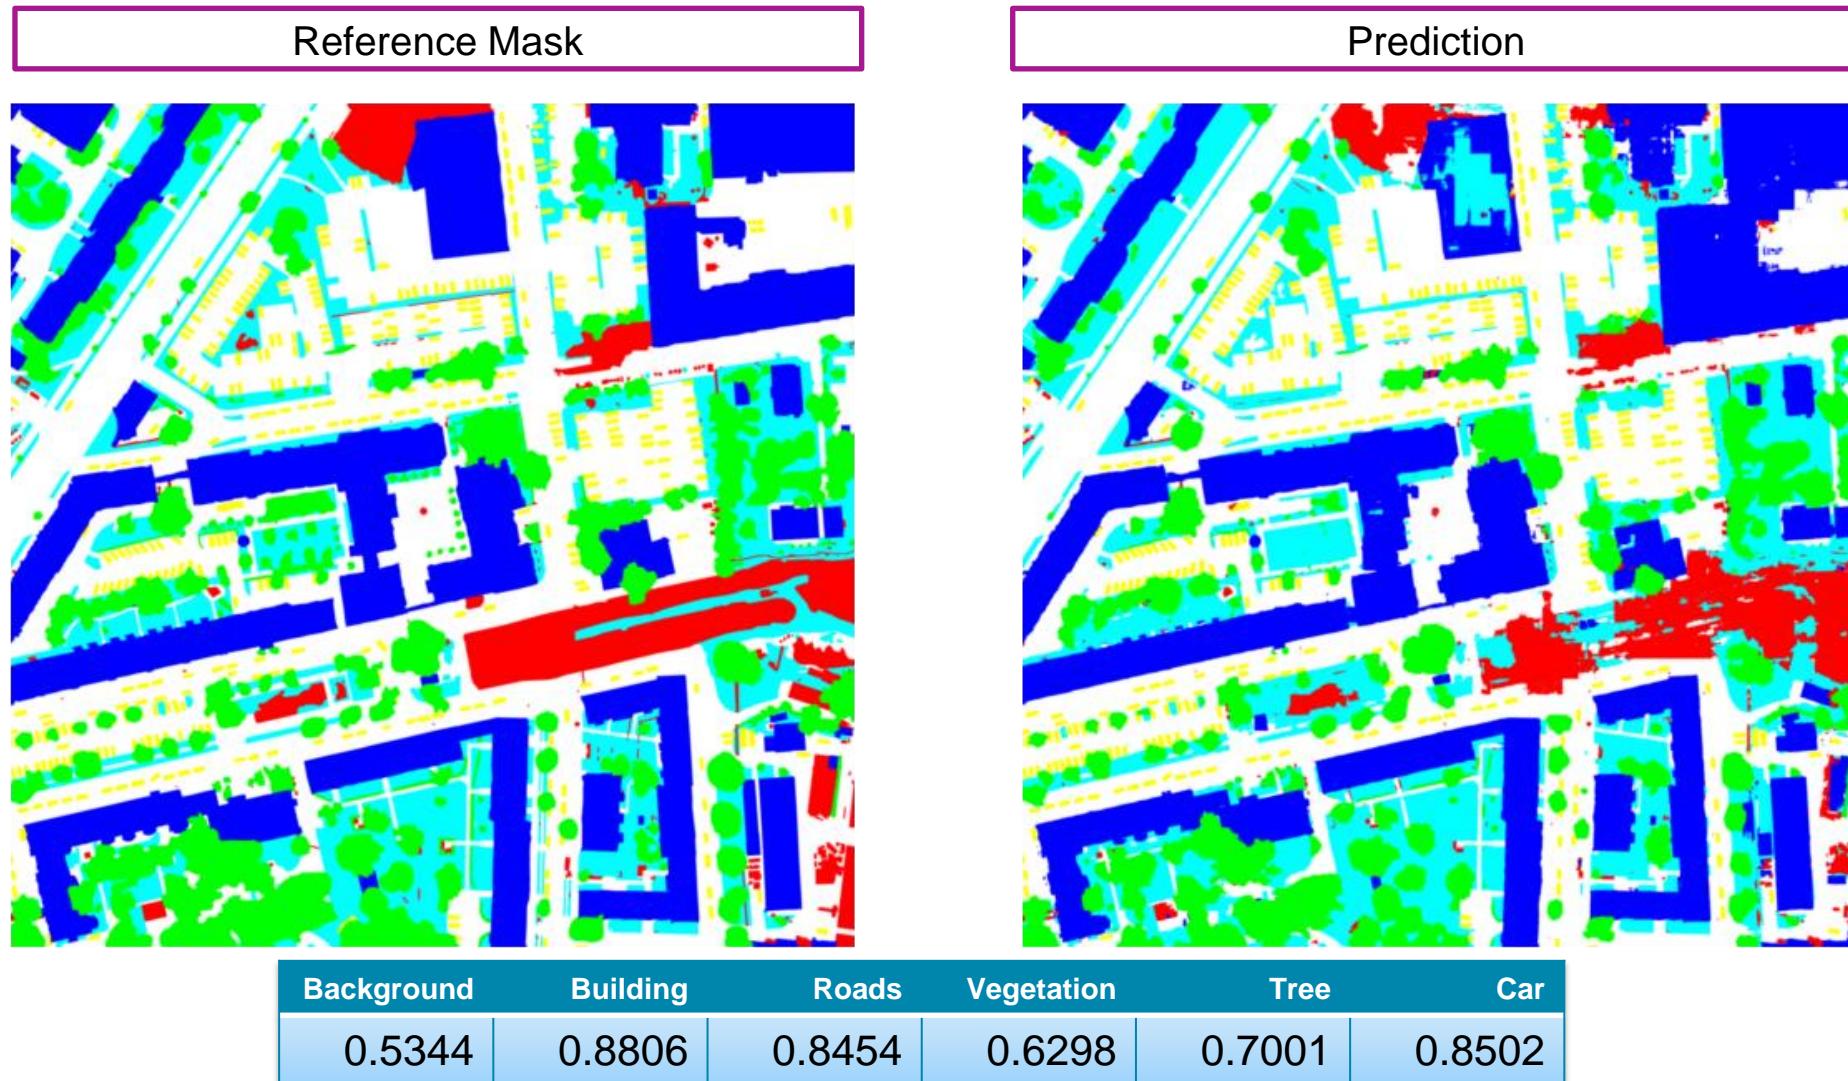
Aerial Tile



Reference Mask



Predicting on a new Tile (Example 3)





Model into the wild

Divider subtitle

Predicting on Marseille

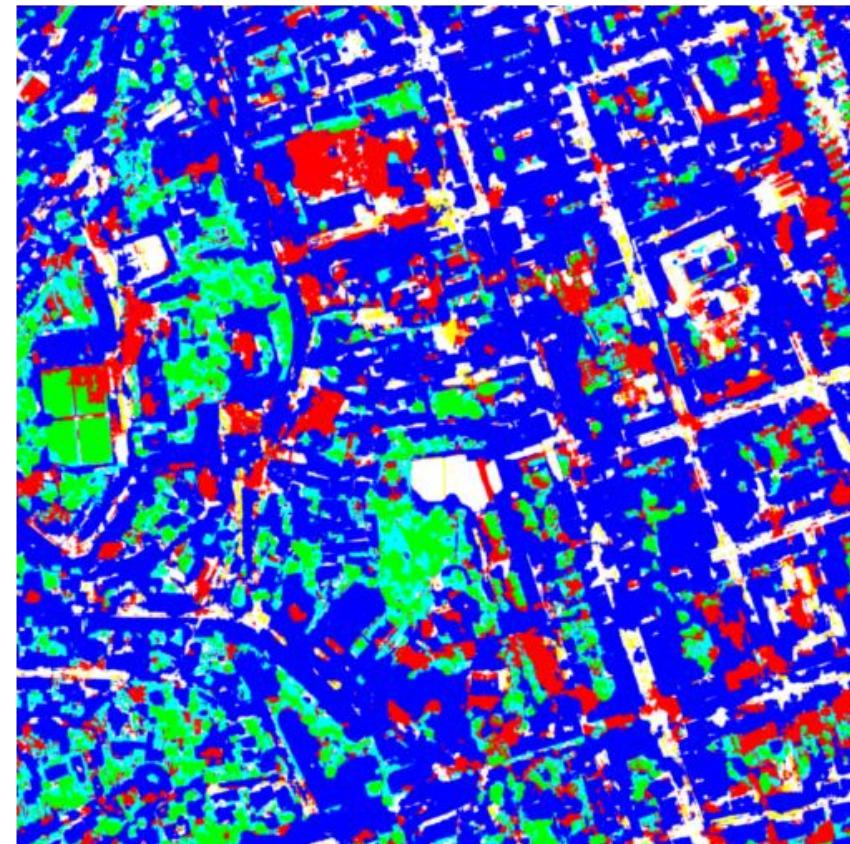


Predicting on Marseille

Aerial Tile

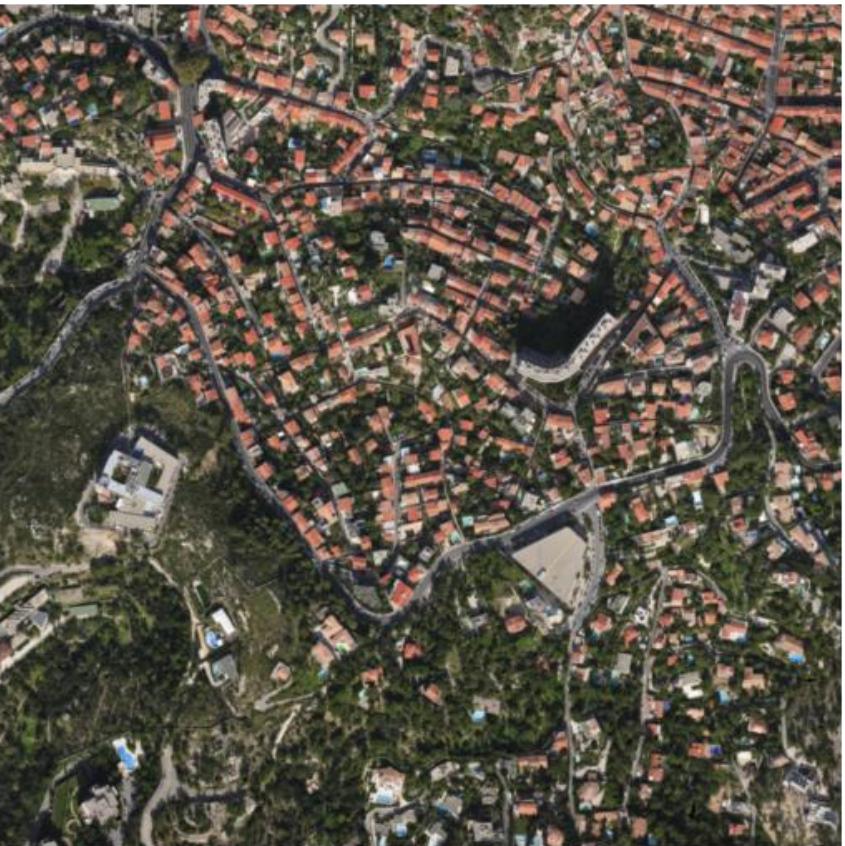


Prediction

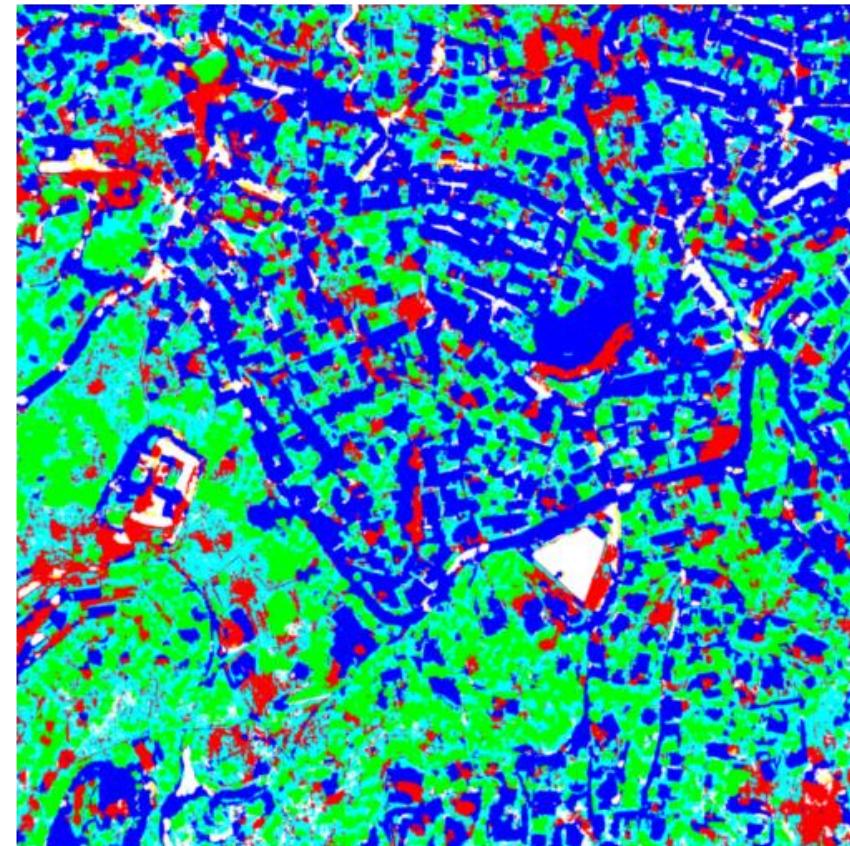


Predicting on Marseille

Aerial Tile



Prediction

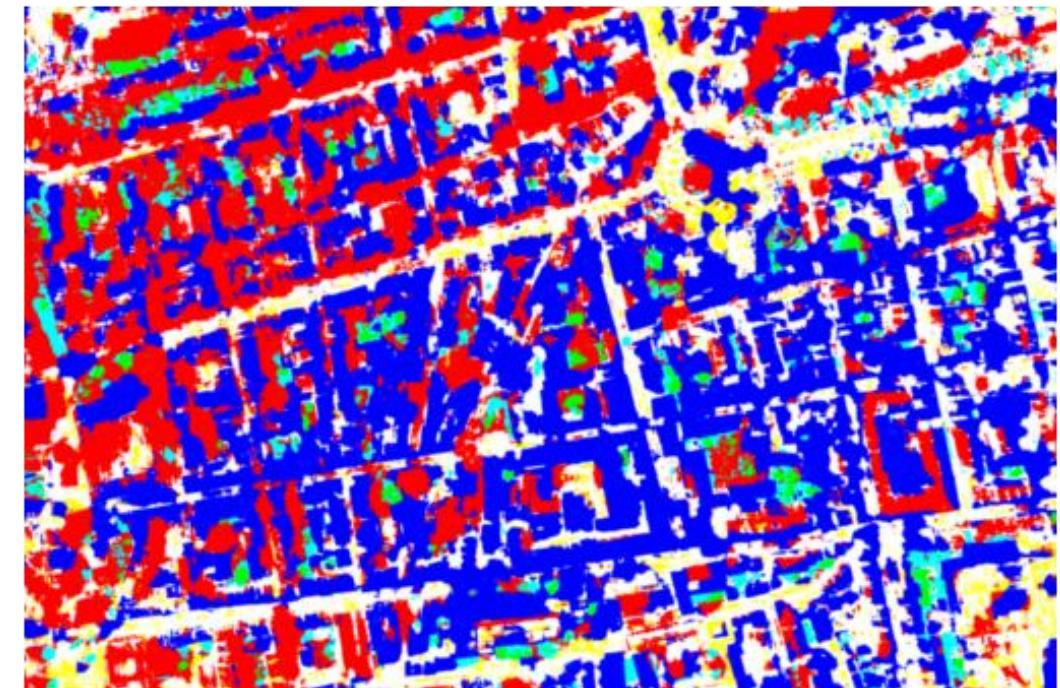


Predicting on Toulouse (NADIR)

Aerial Tile



Prediction



What Happened ?

- Our Model did not generalize well.
- Possible root causes:
 - Training set is too small: limited to one city, tile resolution
 - Implemented Data Augmentation is limited: need more
 - Inference tile have challenging lightning aspects
 - Off-Nadir images does not provide good results

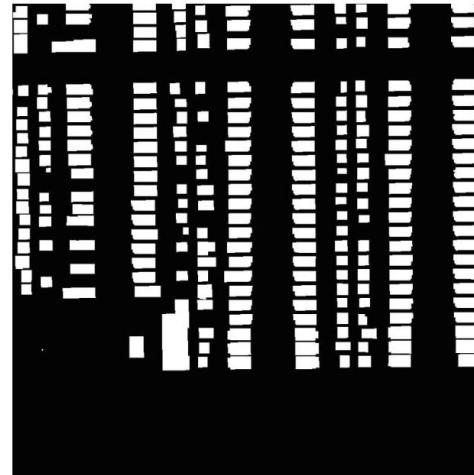
INRIA Aerial Building Dataset

INRIA Building Dataset

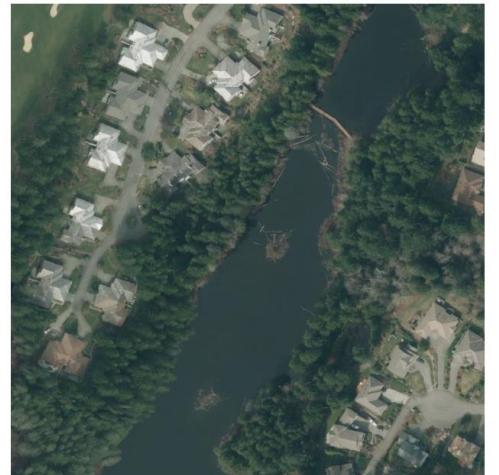
- The training set contains 180 color image tiles of size 5000×5000, covering a surface of 1500 m × 1500 m each (at a 30 cm resolution).
- There are 36 tiles for each of the following regions:
 - Austin
 - Chicago
 - Kitsap County
 - Western Tyrol
 - Vienna



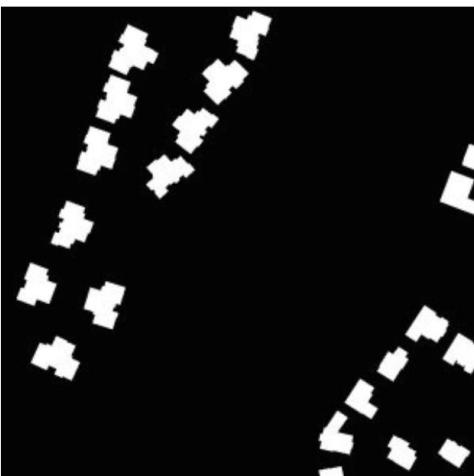
Chicago



Chicago – reference



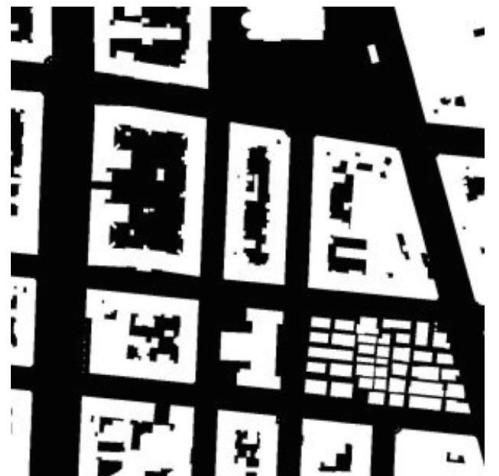
Kitsap County, WA



Kitsap County, WA – Reference



Vienna



Vienna – Reference

INRIA: Defining the Train and the Validation dataset objects

```
[20]: train_ds = get_dataset(img_dir=inria_train_img_dir, mask_dir=inria_train_mask_dir, batch_size=BATCH_SIZE, shuffle=True, repeat=True, augment=True, cachefile='inria_train_cache')
val_ds   = get_dataset(img_dir=inria_val_img_dir, mask_dir=inria_val_mask_dir, batch_size=BATCH_SIZE, shuffle=False, repeat=True, augment=False, cachefile=None)
print("train patch number = {}, valid patch number = {}".format(TRAIN_PATCH_NUMBER, VALID_PATCH_NUMBER))
print("train steps = {}, valid steps = {}".format(TRAIN_STEPS_PER_EPOCH, VAL_STEPS_PER_EPOCH))
```

Removing old cache file: /tmp/inria_train_cache.index
Removing old cache file: /tmp/inria_train_cache.data-00000-of-00001
train patch number = 361, valid patch number = 361
train steps = 1120, valid steps = 181

```
[22]: timelpipeline(train_ds, batch_size=BATCH_SIZE, steps=4000)
```

Error displaying widget: model not found

batch size: 50
4000 batches: 414.9433560371399 s
9.63987 Batches/s
481.99350 Images/s

```
[28]: timelpipeline(val_ds, batch_size=BATCH_SIZE, steps=1000)
```

Error displaying widget: model not found

batch size: 50
1000 batches: 41.72989892959595 s
23.96363 Batches/s
1198.18167 Images/s

INRIA: Defining the Model

```
[47]: inria_model = get_unet_model(MODEL_SHAPE, filters=32, n_block=4, n_class=2, mode=None, residual=True)
inria_model.summary()
```

Total params: 10,730,984

Trainable params: 10,725,096

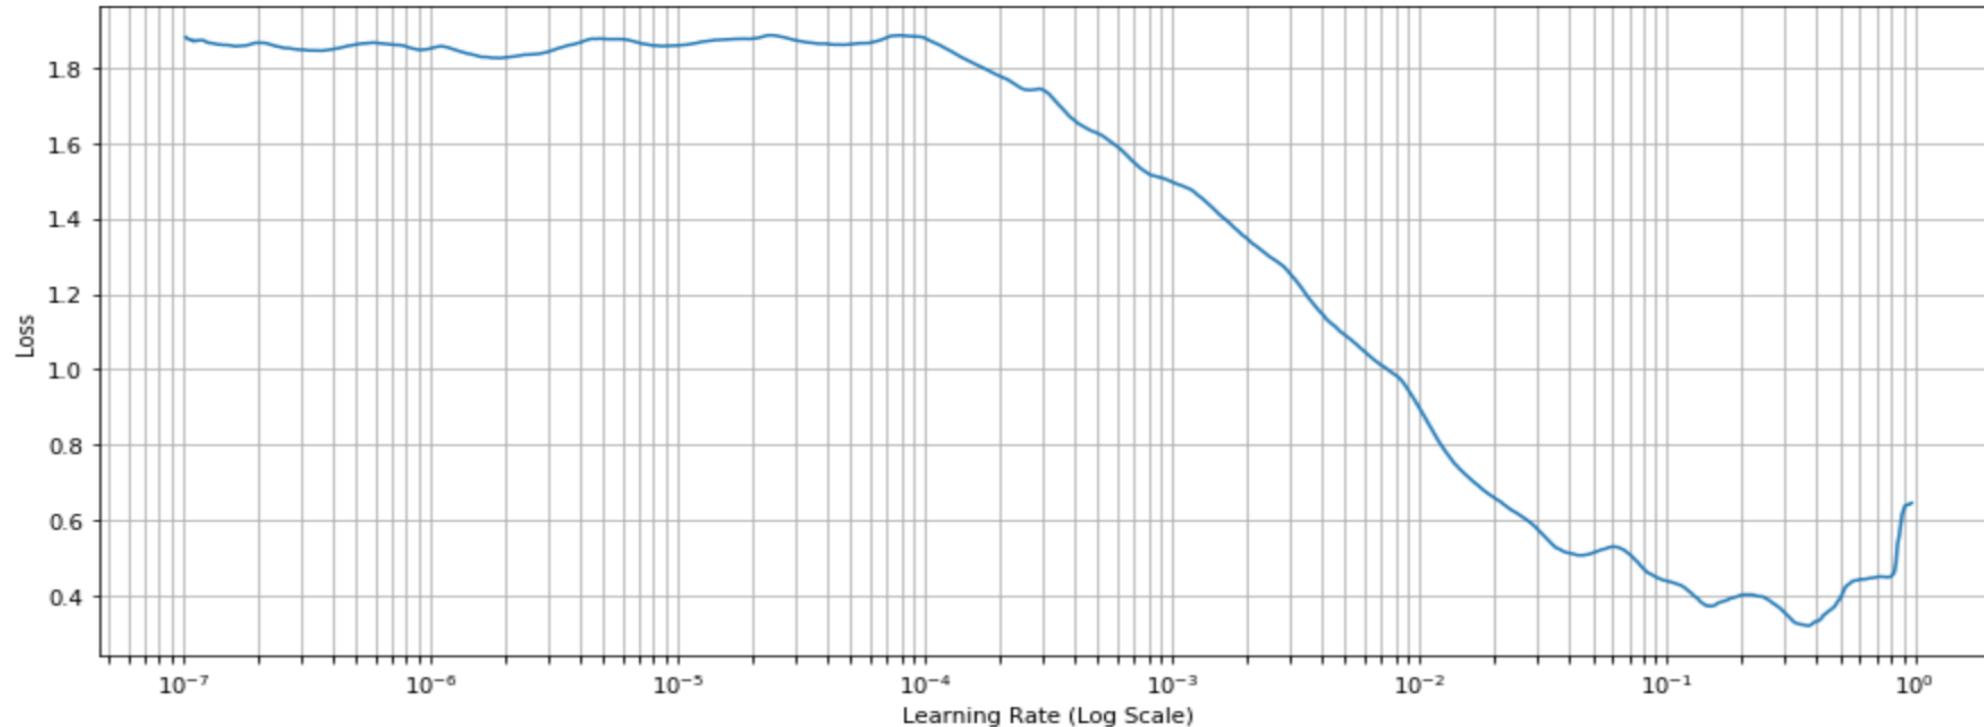
Non-trainable params: 5,888

```
[ ]: experiment_name = "inria_unet_input256_noScale_skip4_Residual_residualCenter_sgd_cross_entropy_bs50_ocpCB"
sgd_optim = tf.keras.optimizers.SGD(lr=0.1, nesterov=True)
acc_metric = tf.keras.metrics.SparseCategoricalAccuracy(name='Accuracy')
iou_metric = tf.keras.metrics.MeanIoU(num_classes=NUM_CLASSES)
fixed_iou = CategoricalMeanIoU(num_classes=NUM_CLASSES)
inria_model.compile(optimizer=sgd_optim, loss='sparse_categorical_crossentropy', metrics=[acc_metric,fixed_iou])
```

INRIA: Finding the Learning Rate

```
[41]: lrf = LearningRateFinder(inria_model)
lrf.find(train_ds, 1e-7, 1, epochs=1, stepsPerEpoch=TRAIN_STEPS_PER_EPOCH)
lrf.plot_loss()

Train for 1120 steps
1120/1120 [=====] - 794s 709ms/step - loss: 1.3063 - Accuracy: 0.6399 - mean_iou: 0.3948
```



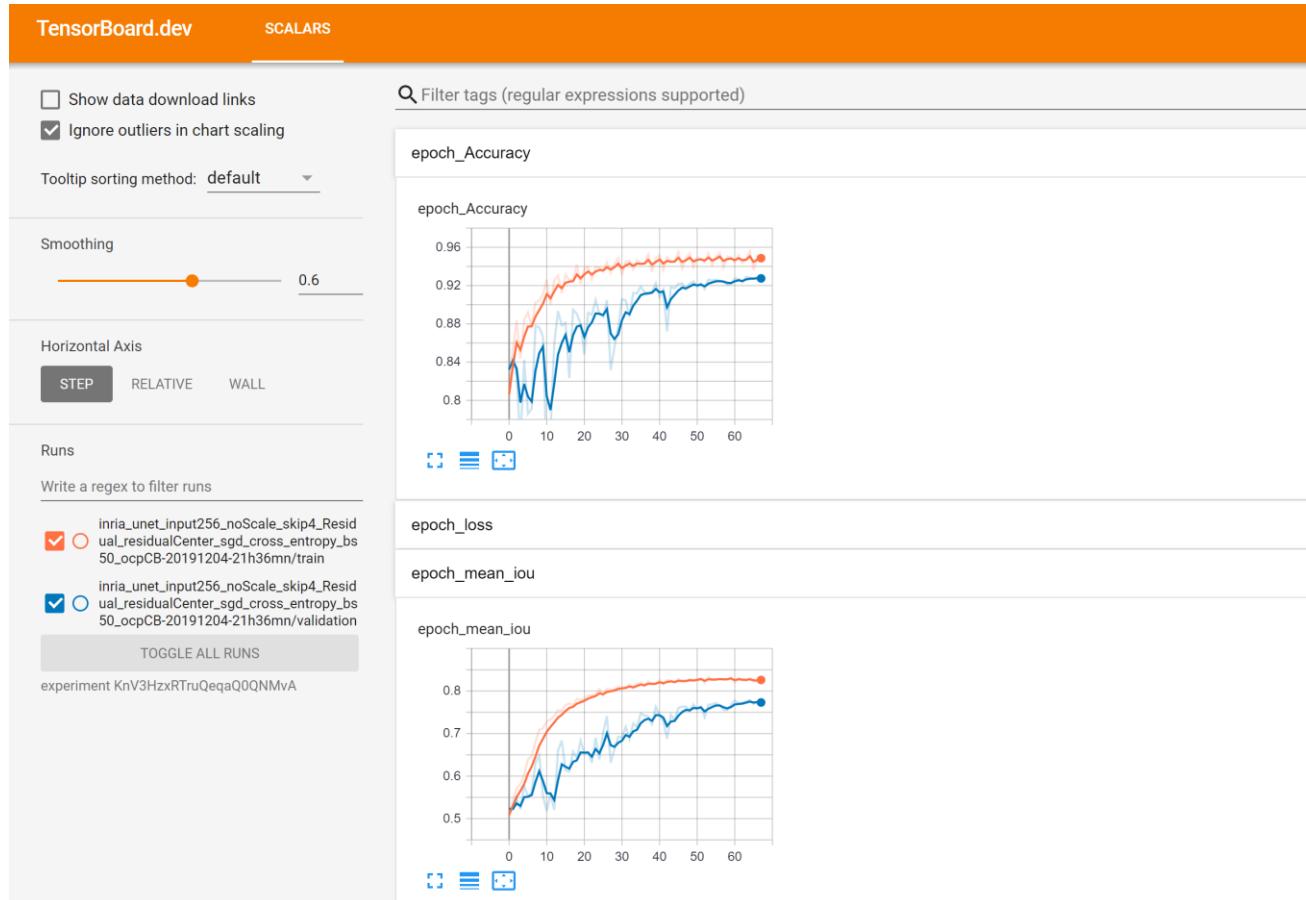
INRIA: Training Phase

```
[ ]: # One cycle policy training
checkpoint_path = "../inria_checkpoint_dir/" + experiment_name + datetime.now().strftime("%Y%m%d-%Hh%Mmn") + "/best-cp-" + experiment_name
tb_log_dir = "../inria_log_dir/" + experiment_name + datetime.now().strftime("%Y%m%d-%Hh%Mmn")
# callbacks
ocp_cb = OneCycleScheduler(max_lr=1e-1)
checkpoint_cb = tf.keras.callbacks.ModelCheckpoint(filepath=checkpoint_path, monitor='val_mean_iou', mode='max', save_weights_only=True, save_best_only=True, verbose=1)
tensorboard_cb = tf.keras.callbacks.TensorBoard(log_dir=tb_log_dir)

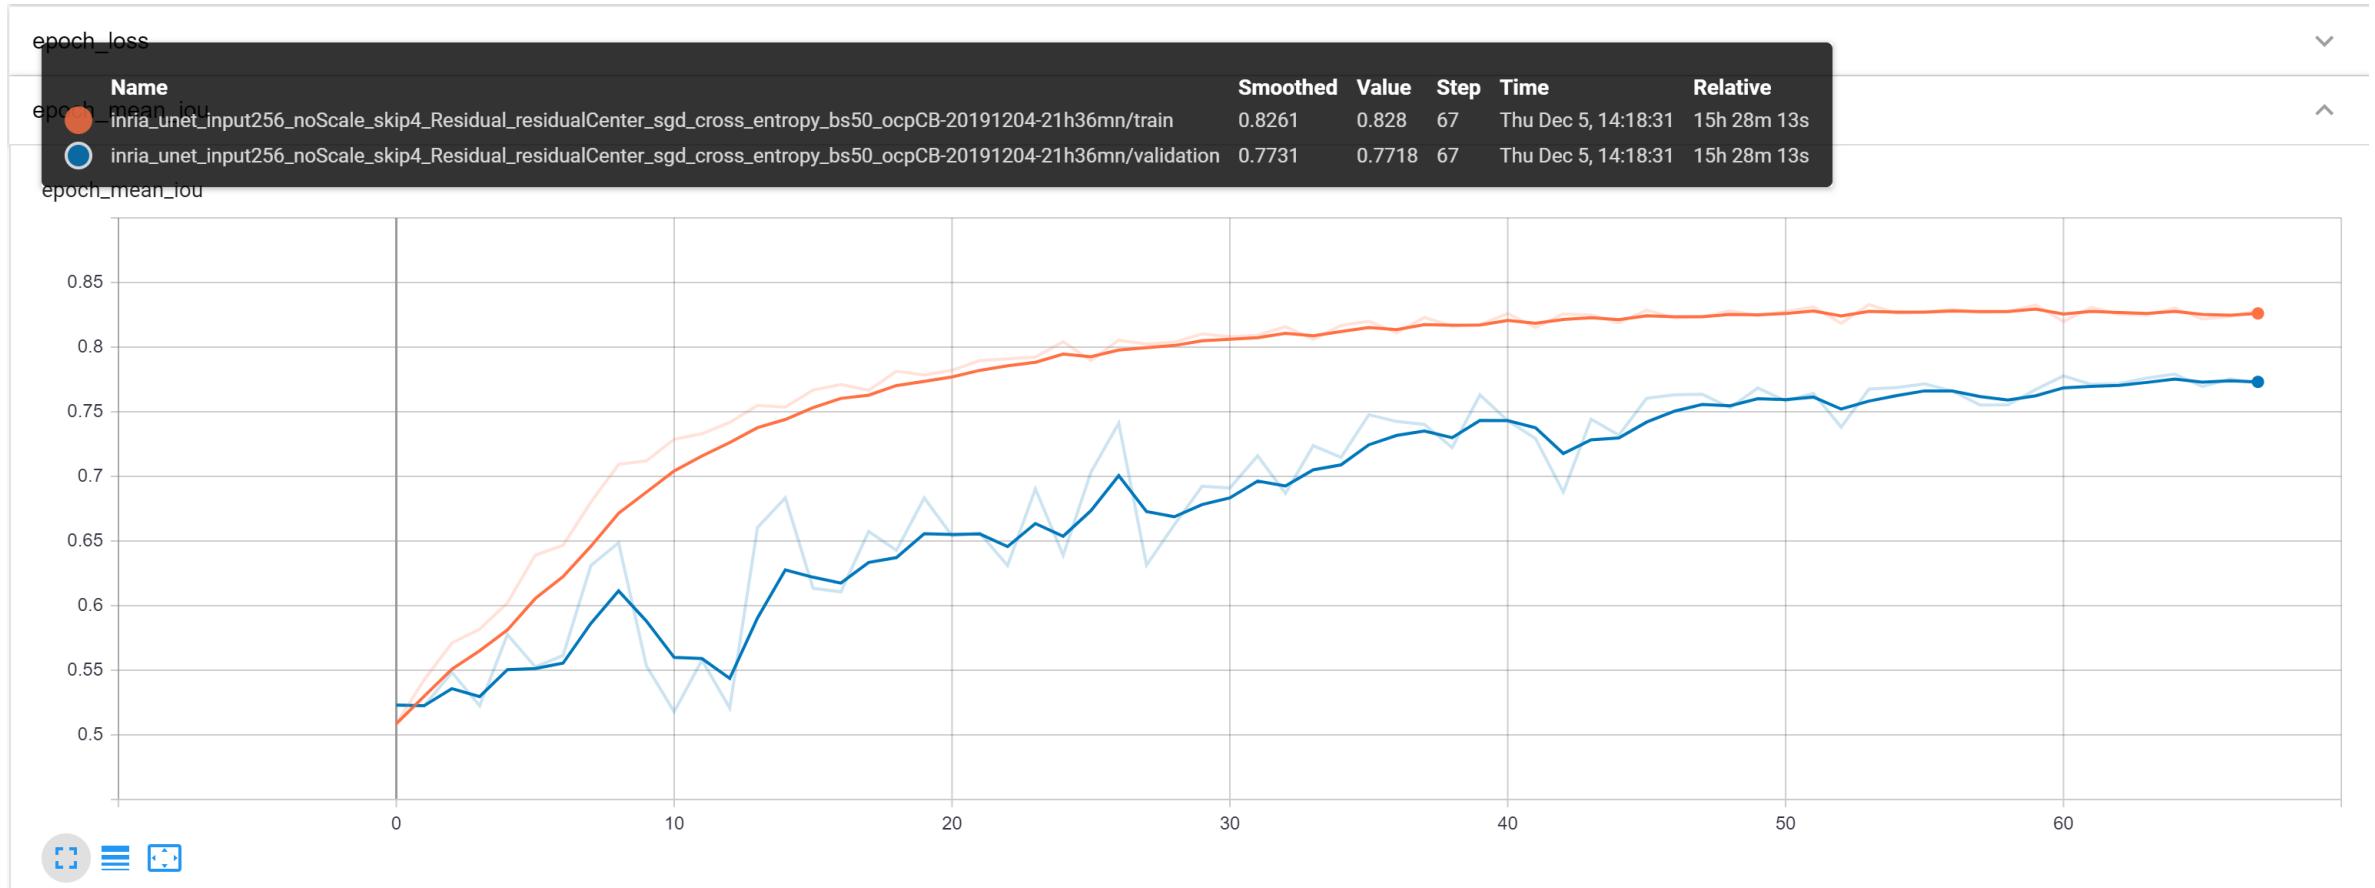
keras_history = inria_model.fit(train_ds,
                                steps_per_epoch=TRAIN_STEPS_PER_EPOCH,
                                epochs=70,
                                validation_data=val_ds,
                                validation_steps=VAL_STEPS_PER_EPOCH,
                                callbacks=[ocp_cb, checkpoint_cb, tensorboard_cb],
                                verbose=1)

Train for 1120 steps, validate for 181 steps
Epoch 1/70
1119/1120 [=====>.] - ETA: 0s - loss: 0.5033 - Accuracy: 0.8061 - mean_iou: 0.5086 - OneCycleScheduler, lr: 0.0045361, momentum: 0.9491623
Epoch 00001: val_mean_iou improved from -inf to 0.52302, saving model to ../inria_checkpoint_dir/inria_unet_input256_noScale_skip4_Residual_residualCenter_sgd_cross_entropy_bs50_ocpCB-2019
1120/1120 [=====] - 833s 744ms/step - loss: 0.5031 - Accuracy: 0.8062 - mean_iou: 0.5086 - val_loss: 0.4252 - val_Accuracy: 0.8324 - val_mean_iou: 0.5230
Epoch 2/70
1119/1120 [=====>.] - ETA: 0s - loss: 0.3631 - Accuracy: 0.8538 - mean_iou: 0.5426 - OneCycleScheduler, lr: 0.0061325, momentum: 0.9466680
Epoch 00002: val_mean_iou did not improve from 0.52302
1120/1120 [=====] - 830s 741ms/step - loss: 0.3631 - Accuracy: 0.8538 - mean_iou: 0.5426 - val_loss: 0.3866 - val_Accuracy: 0.8473 - val_mean_iou: 0.5223
Epoch 3/70
296/1120 [=====>.....] - ETA: 9:31 - loss: 0.3196 - Accuracy: 0.8729 - mean_iou: 0.5586
```

INRIA: Monitoring the Training Phase



INRIA: Monitoring the Training Phase



INRIA: Evaluation

```
[ ]: inria_model.load_weights('..../inria_checkpoint_dir/.../best-cp-inria_unet_input256_noScale_skip4_Residual_residualCenter_sgd_cross_entropy_bs50_ocpCB')
```

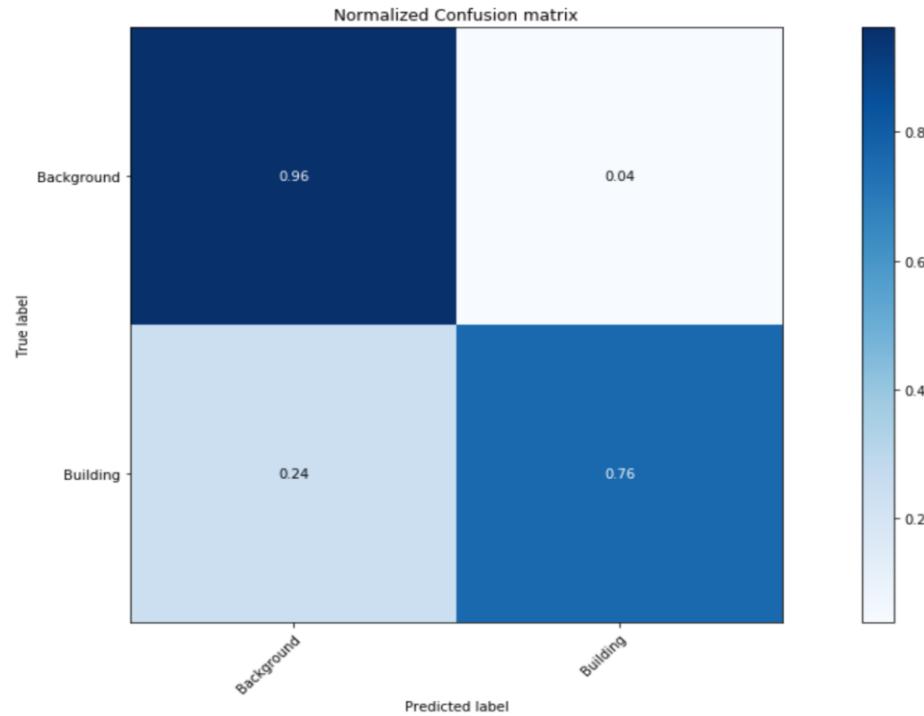
```
[46]: inria_model.evaluate(val_ds, steps=VAL_STEPS_PER_EPOCH)
```

```
181/181 [=====] - 53s 293ms/step - loss: 0.1810 - Accuracy: 0.9283 - mean_iou: 0.7791
```

```
[46]: [0.1810000821824561, 0.92829496, 0.77909243]
```

```
[50]: compute_dataset_iou(inria_model, val_ds, VAL_STEPS_PER_EPOCH)
```

Error displaying widget: model not found

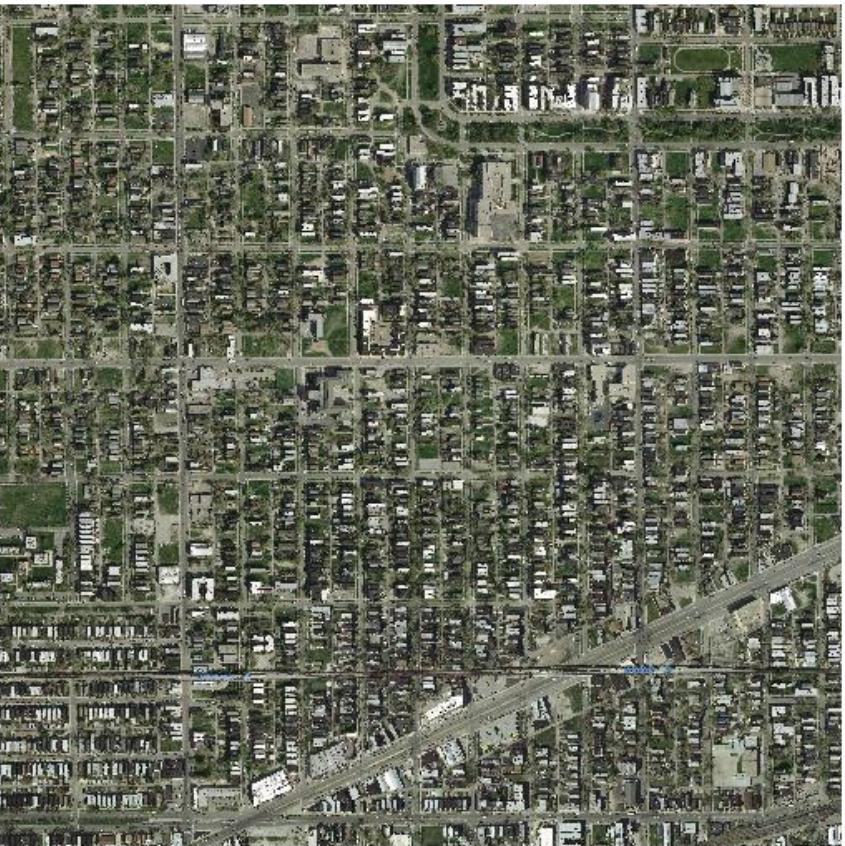


```
[50]: Background Building
```

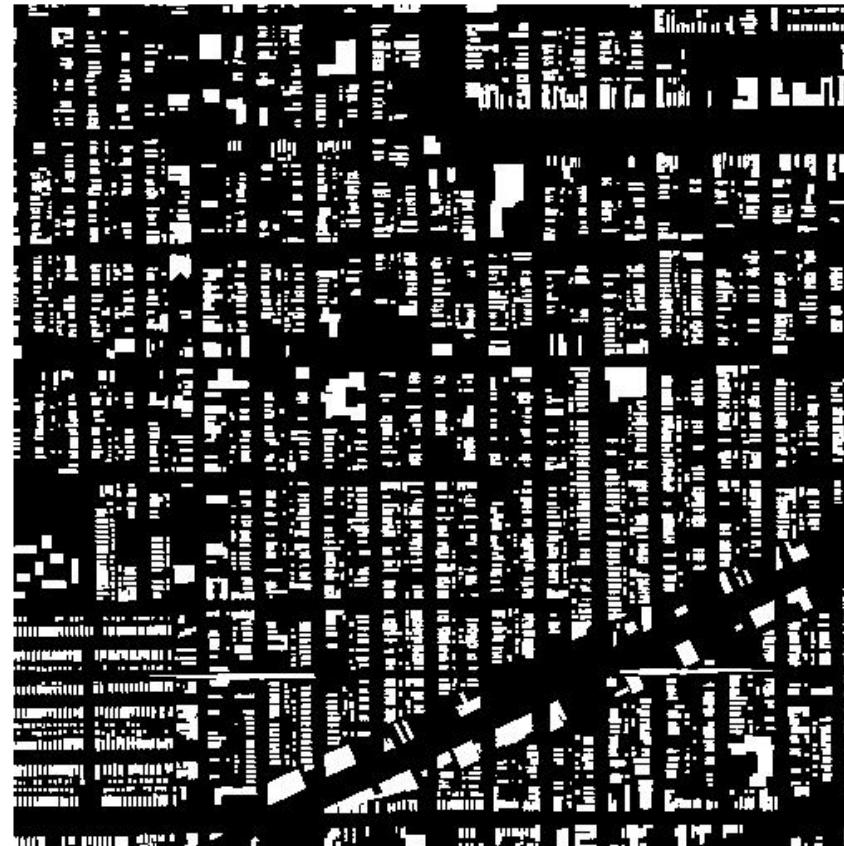
```
0 0.9178 0.6404
```

INRIA: predicting on a Validation tile (chicago)

Aerial Tile

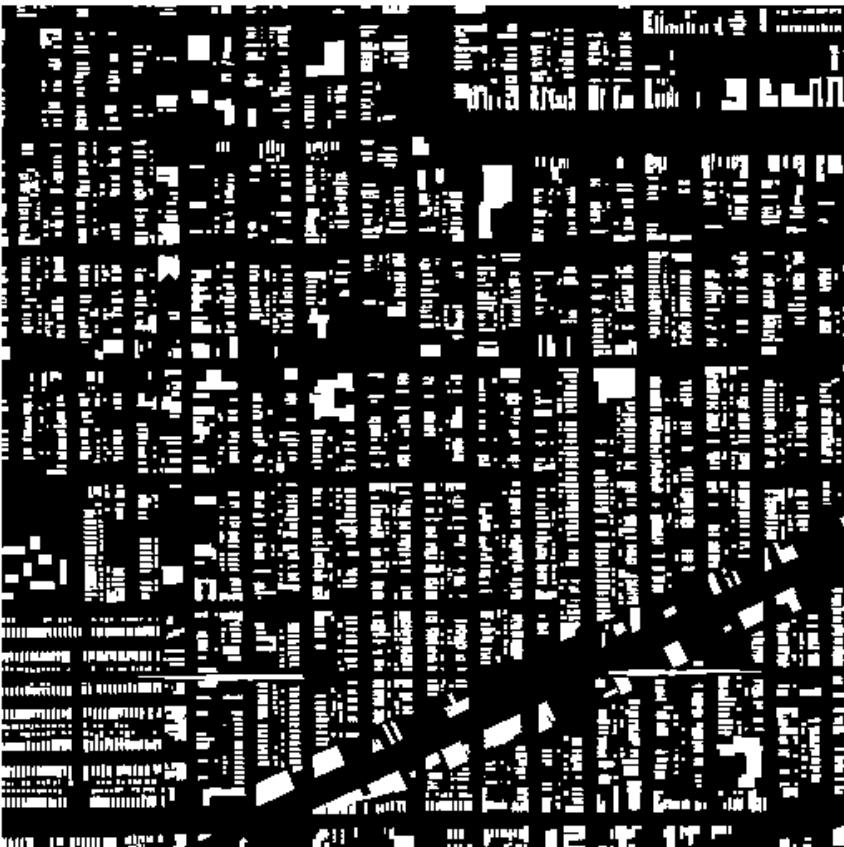


Reference Mask

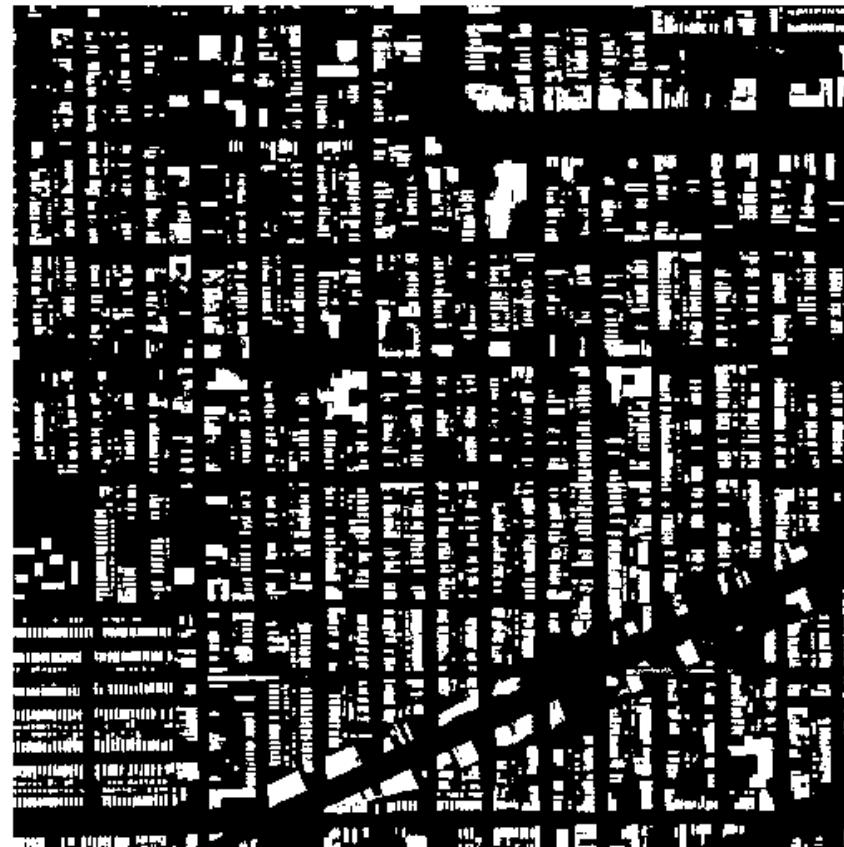


INRIA: predicting on a Validation tile (chicago)

Reference Mask



Prediction



```
[159]: get_tile_iou(msk, y_pred2)
```

```
[159]: Background Building
```

0	0.8908	0.6419
---	--------	--------

INRIA: Predicting on Marseille



INRIA: Predicting on Marseille

Aerial Tile



Prediction



INRIA: Predicting on Marseille (oblique image)

Aerial Tile



Prediction



Conclusion

- Produced a fast, flexible and robust input pipeline that make optimal usage of CPU & GPU resources
- Trained Deep learning Neural Network for segmentation with transparent performance metrics and reproducible flow
- Tried cutting-edge research technique to train neural networks (learning rate finder, One Cycle Policy)
- Tried New Loss function (Focal Loss, Jaccard Index Loss)
- Reached state of the Art results on the ISPRS dataset
- Learned (the hard way) that Model Performance is tightly coupled with its training data: Variety and Quantity are paramount

Where to go from here

- Synthetic datasets: Generative model for data augmentation
- Combine Geometric Computer vision with Deep Learning
- Prediction post processing

Bibliography

- UNet: Convolutional Networks for Biomedical Image Segmentation: <https://arxiv.org/abs/1505.04597>
- Cyclical Learning Rates for Training Neural Networks: <https://arxiv.org/abs/1506.01186>
- A disciplined approach to neural network hyper-parameters: Part 1 learning rate, batch size, momentum, and weight decay: <https://arxiv.org/abs/1803.09820>
- Super-Convergence: Very Fast Training of Neural Networks Using Large Learning Rates: <https://arxiv.org/abs/1708.07120>
- Focal Loss for Dense Object Detection: <https://arxiv.org/abs/1708.02002>
- Deep Residual Learning for Image Recognition: <https://arxiv.org/abs/1512.03385>
- Densely Connected Convolutional Networks: <https://arxiv.org/abs/1608.06993>
- Semantic Labeling in Very High Resolution Images via a Self-Cascaded Convolutional Neural Network:
https://www.researchgate.net/publication/321970661_Semantic_Labeling_in_Very_High_Resolution_Images_via_a_Self-Cascaded_Convolutional_Neural_Network_ISPRS_JPRS_IF6942

Thank you